

April , 2021

# RabbitMQ

## BuildingBlocks and messaging patterns



# RabbitMQ

## The Building Blocks

**Routing**

Fanout Exchange

Direct Exchange

Topic Exchange

Header Exchange

Consistent Hashing Exchange

Random Exchange

Sharding

**Queues and Consumers**

Queues

Competing Consumers

Non-Competing Consumers

Ephemeral Queues

Lazy Queues

Priority Queues

**More**

Deadletter Exchange

Alternate Exchange

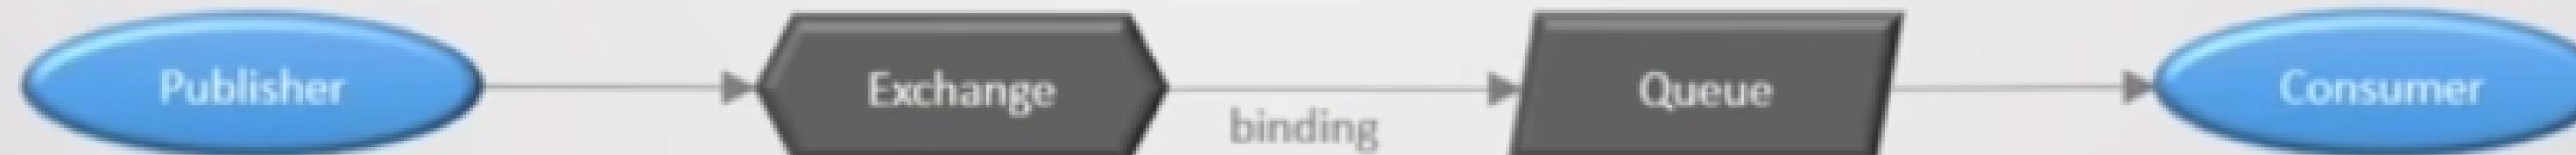
Virtual Hosts

# RabbitMQ Building Blocks

Traditional queues on  
steroids

## Exchanges, Queues and Bindings

- Publishers send messages to Exchanges
- Exchanges route messages to Queues and even other Exchanges.
- Consumers read from Queues
- Bindings link Exchanges to Queues and even Exchanges to Exchanges

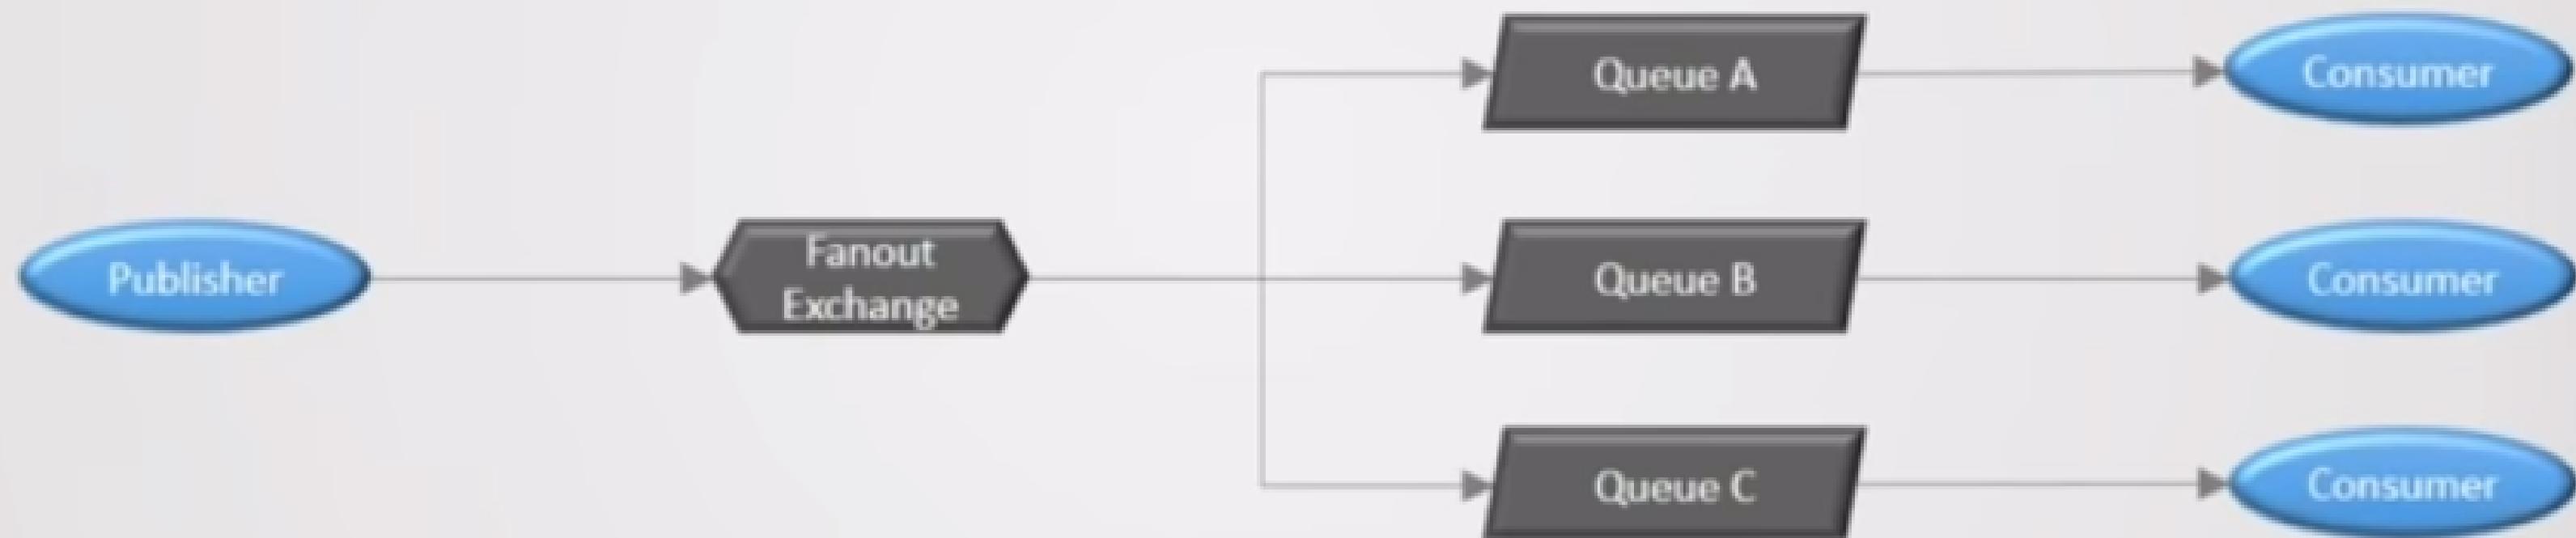


# Publish - Subscribe

## Fanout

Simplest Publish-Subscribe pattern.

Each consumer receives every message

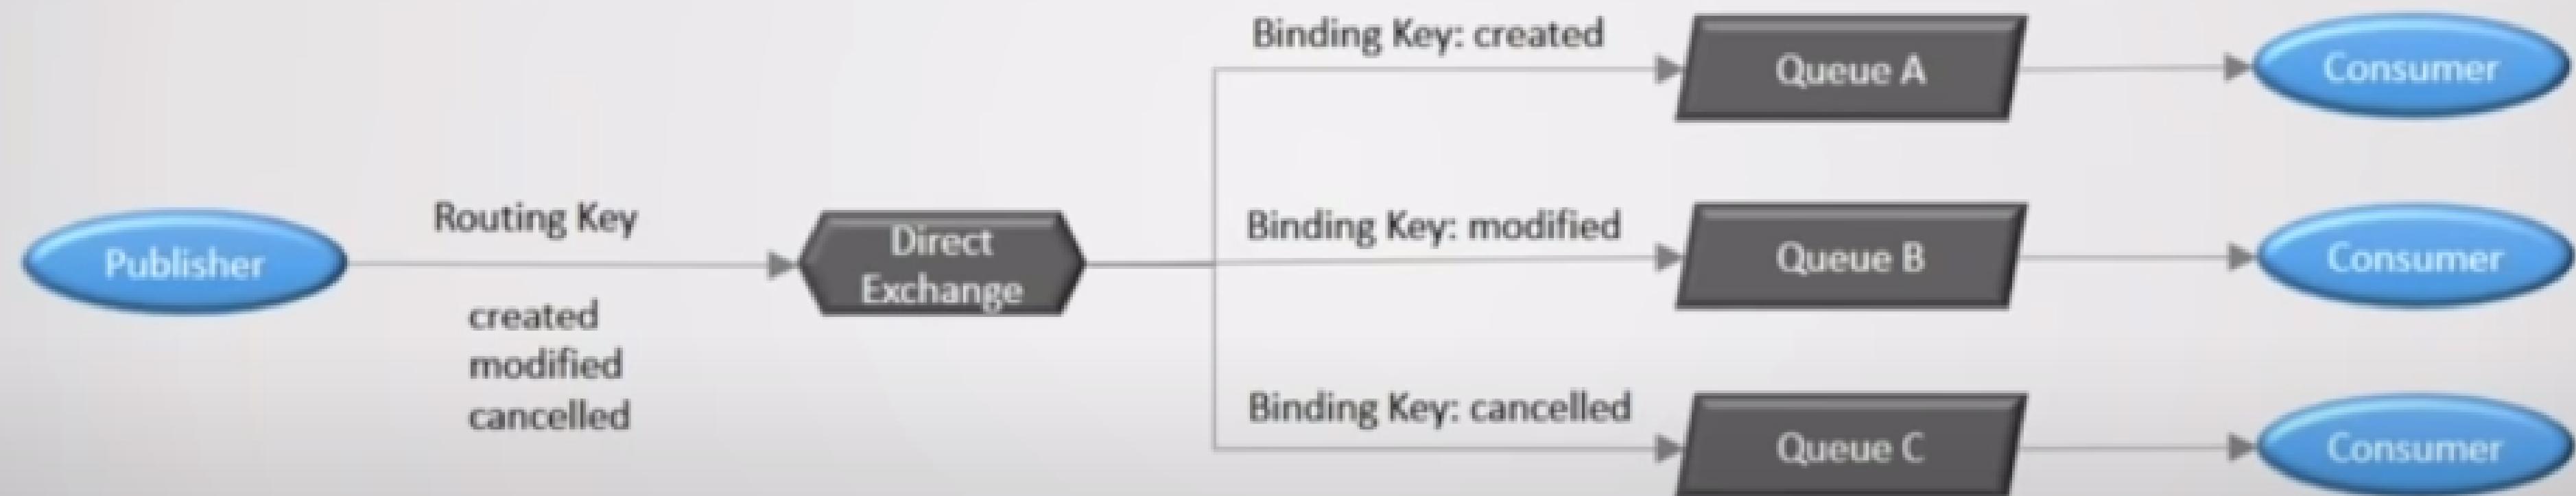


# Publish - Subscribe with Routing

## Direct Exchange

### Routing by Routing Key

Exact match routing key = binding key



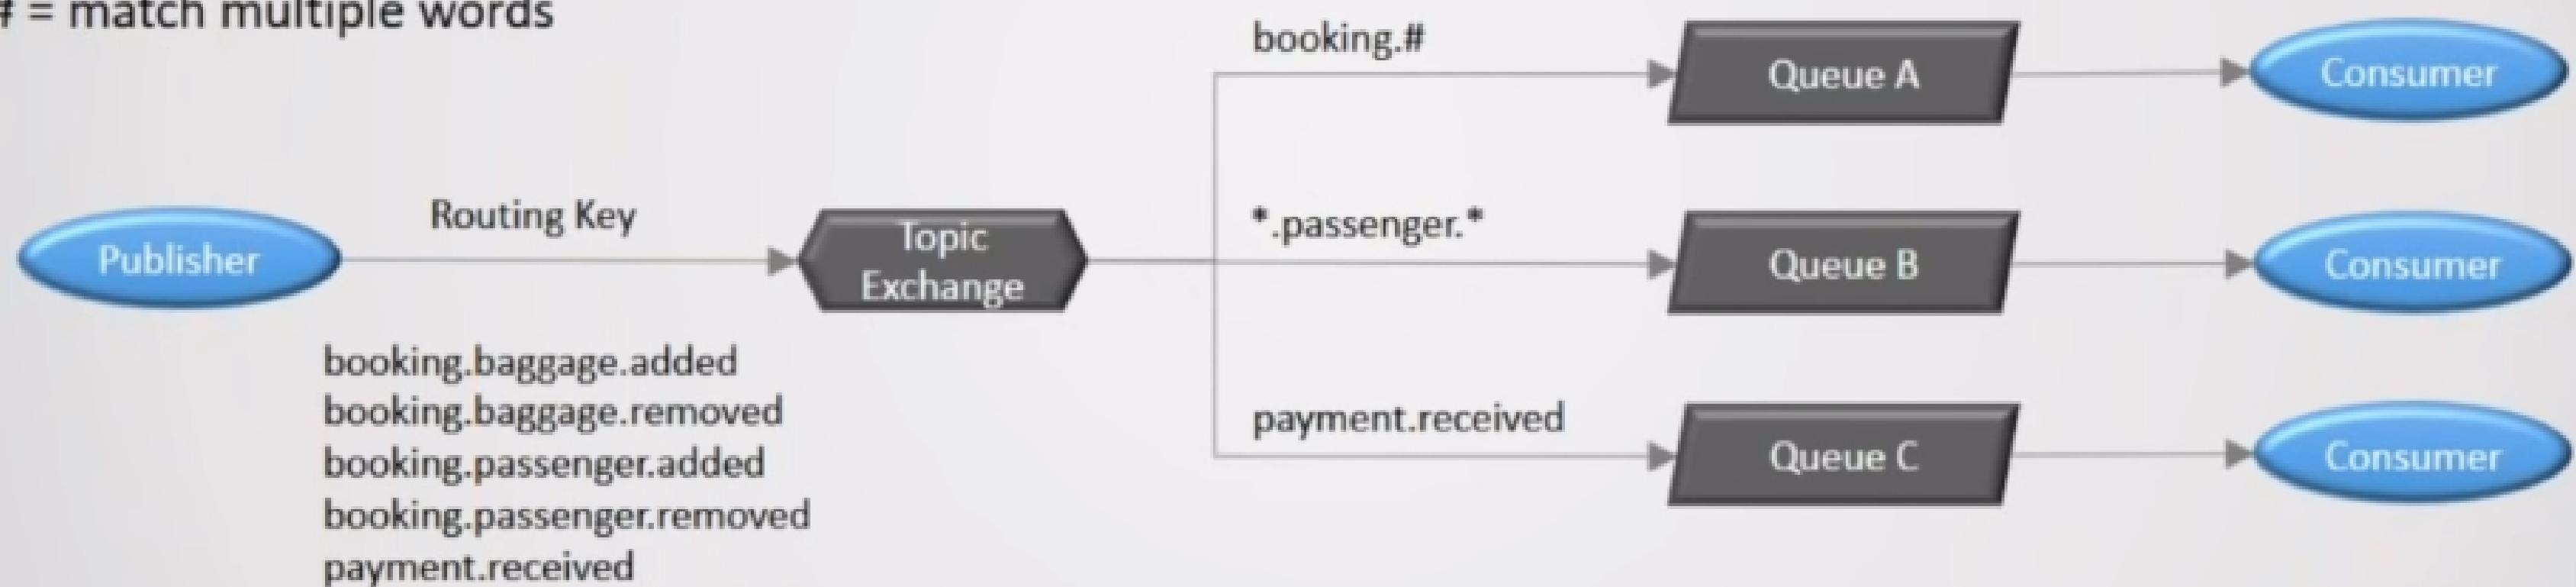
# Publish - Subscribe with Routing

## Topic Exchange

### Wildcard routing by Routing Key

\* = match 1 word

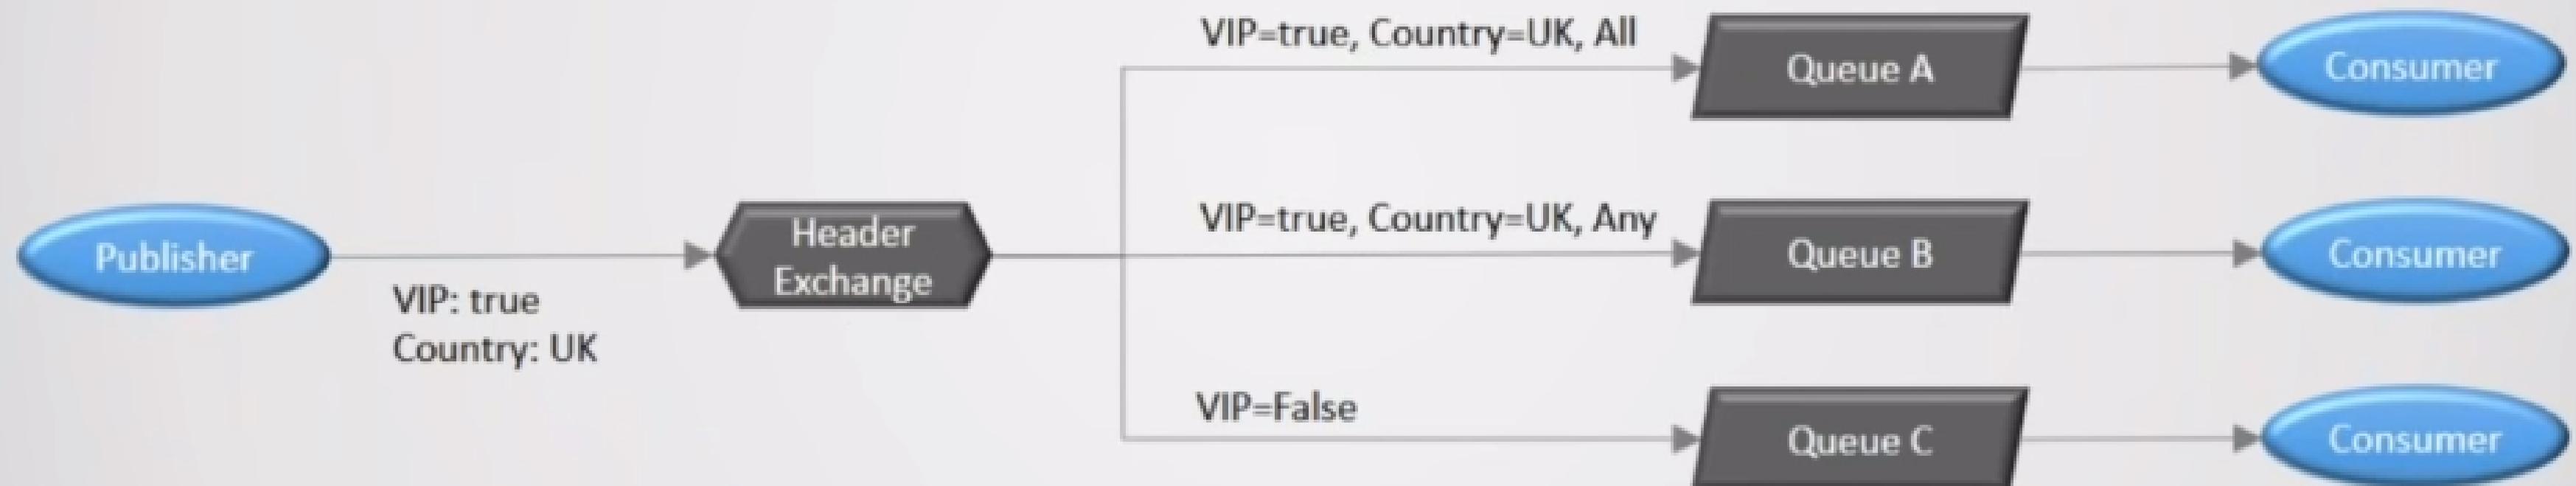
# = match multiple words



# Publish - Subscribe with Routing

## Header Exchange

Routing by message headers



# Point-to-Point Messaging

## Default Exchange

Direct exchange type

All queues have implicit binding to the default exchange.

Routes messages by Routing Key to Queue Name

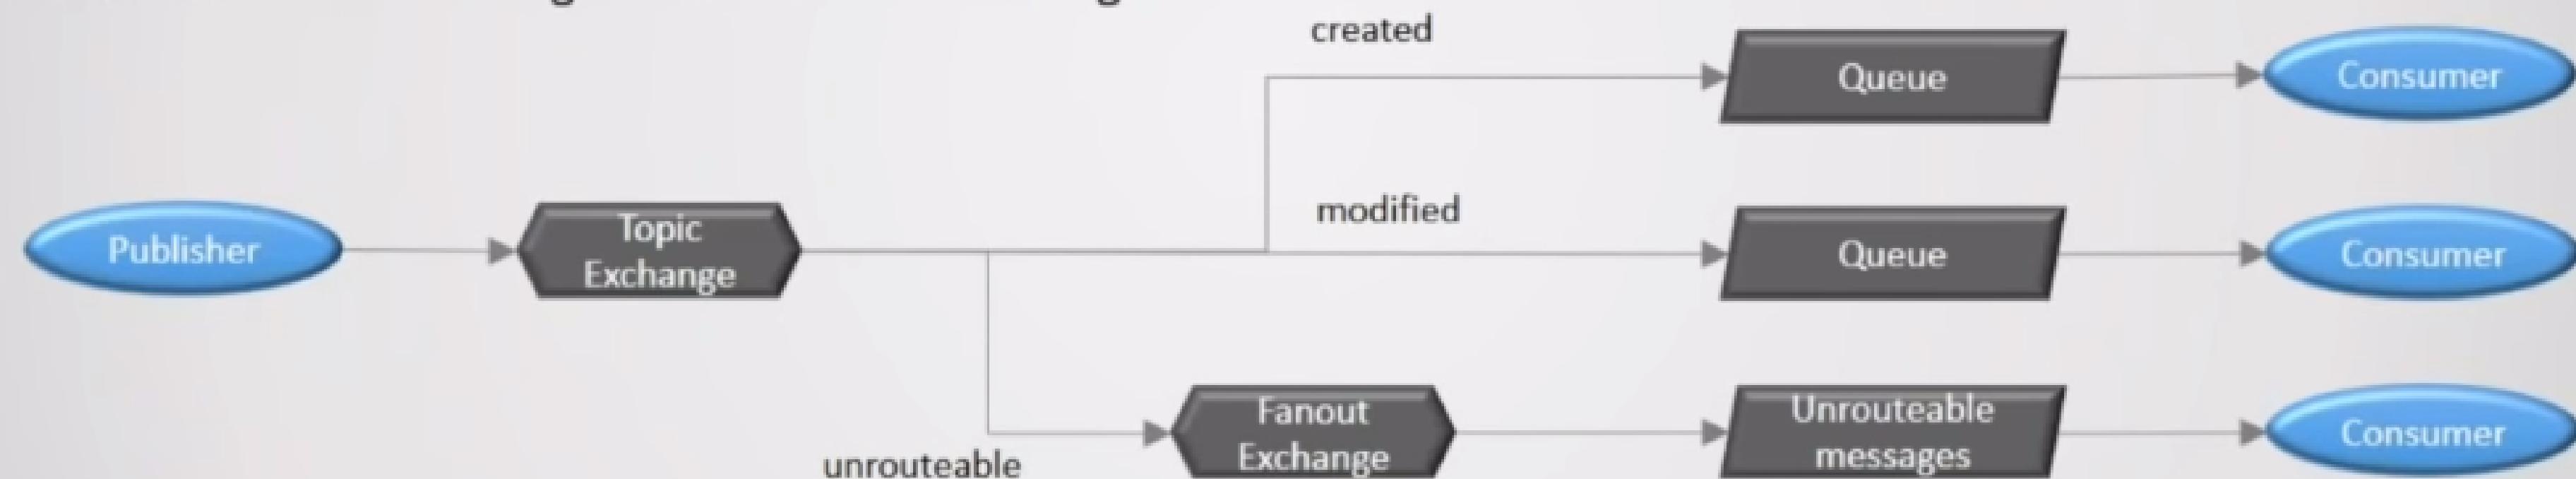


# Protecting Against Routing Failures

## Alternate Exchange

Not an exchange type but a configured relationship between two exchanges.

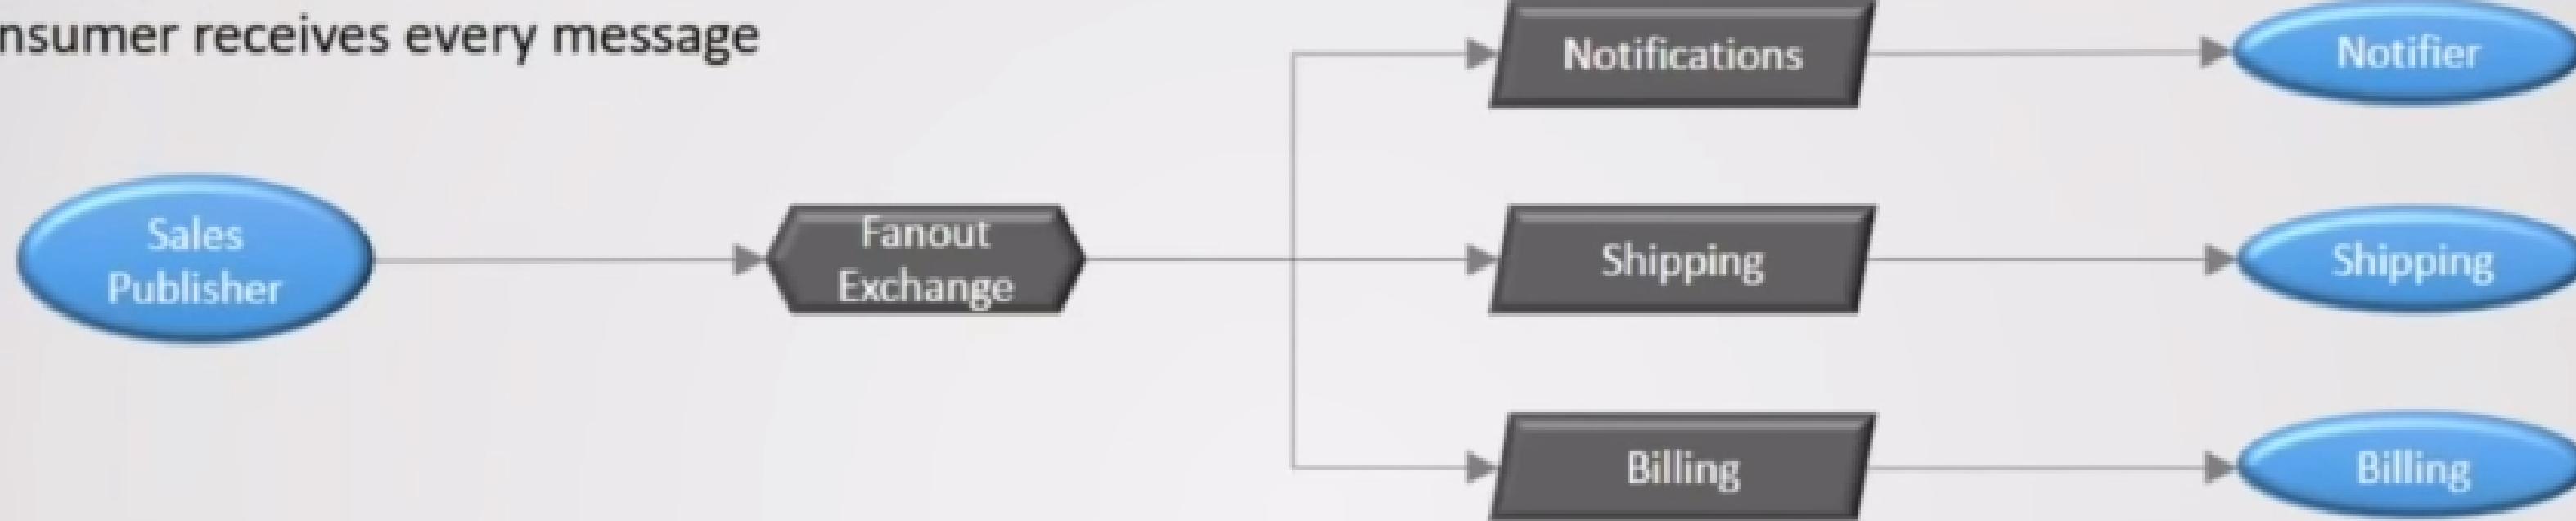
Route unrouteable messages to an alternate exchange



# Scaling Out Consumers

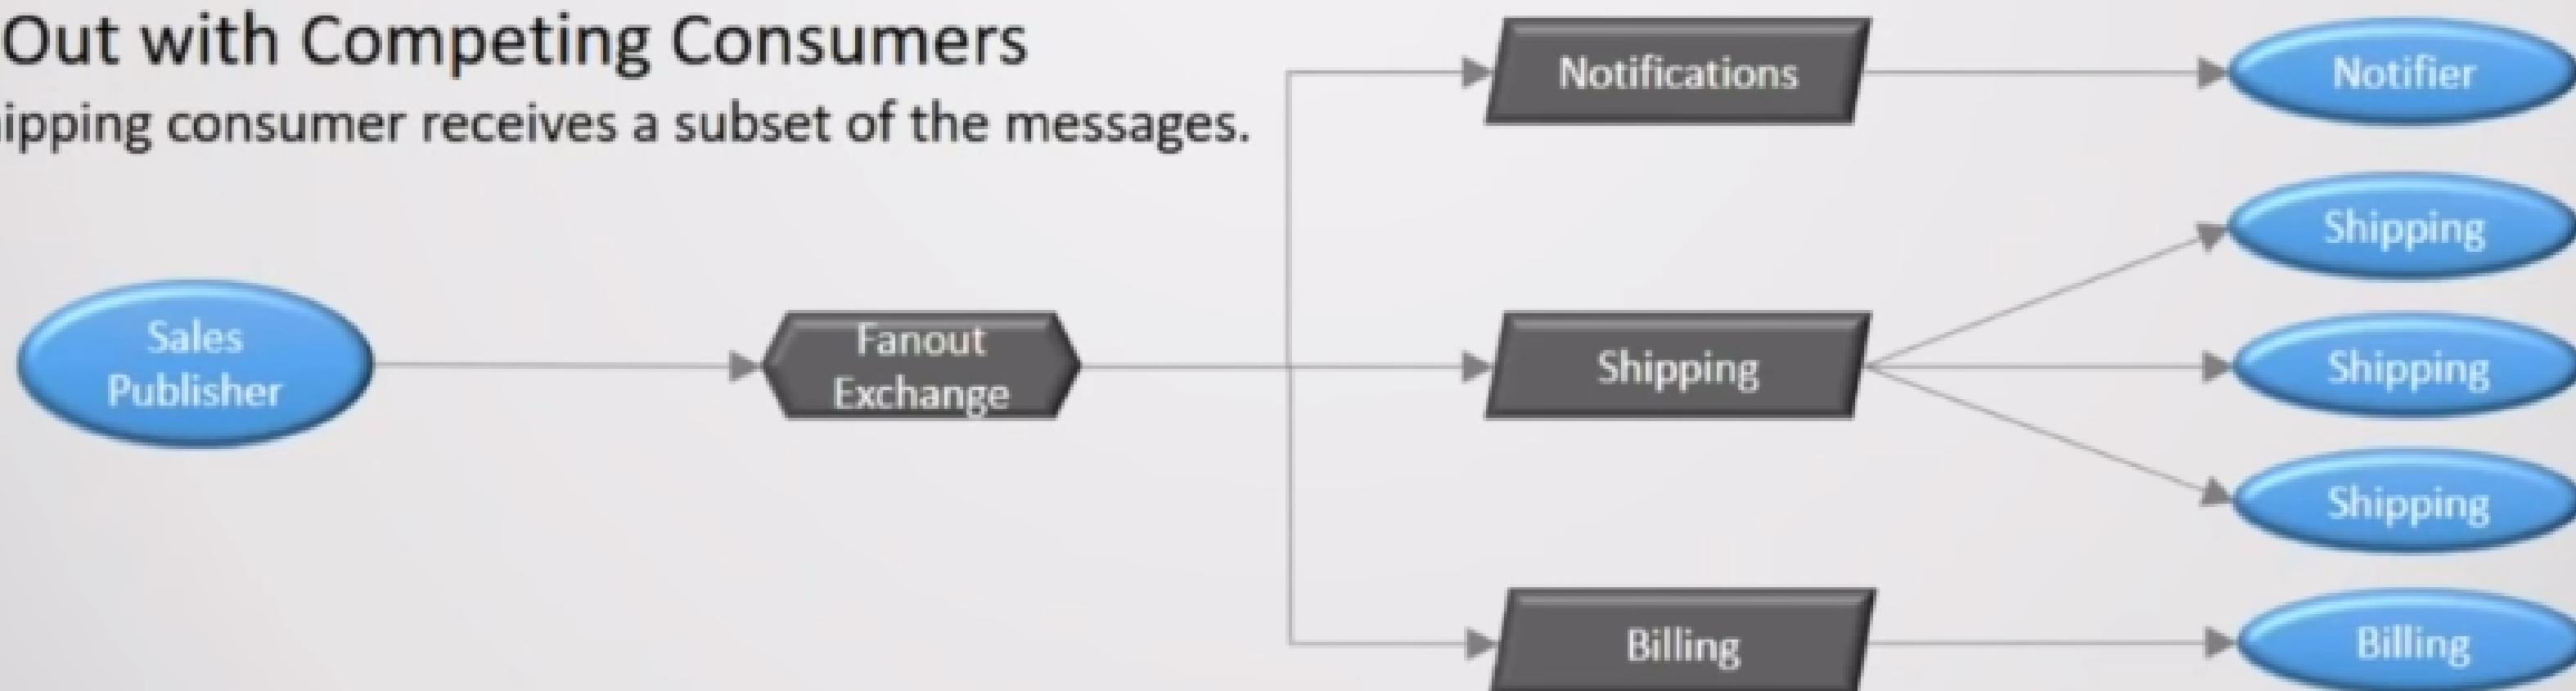
## Non-Competing Consumers

Each consumer receives every message



## Scale Out with Competing Consumers

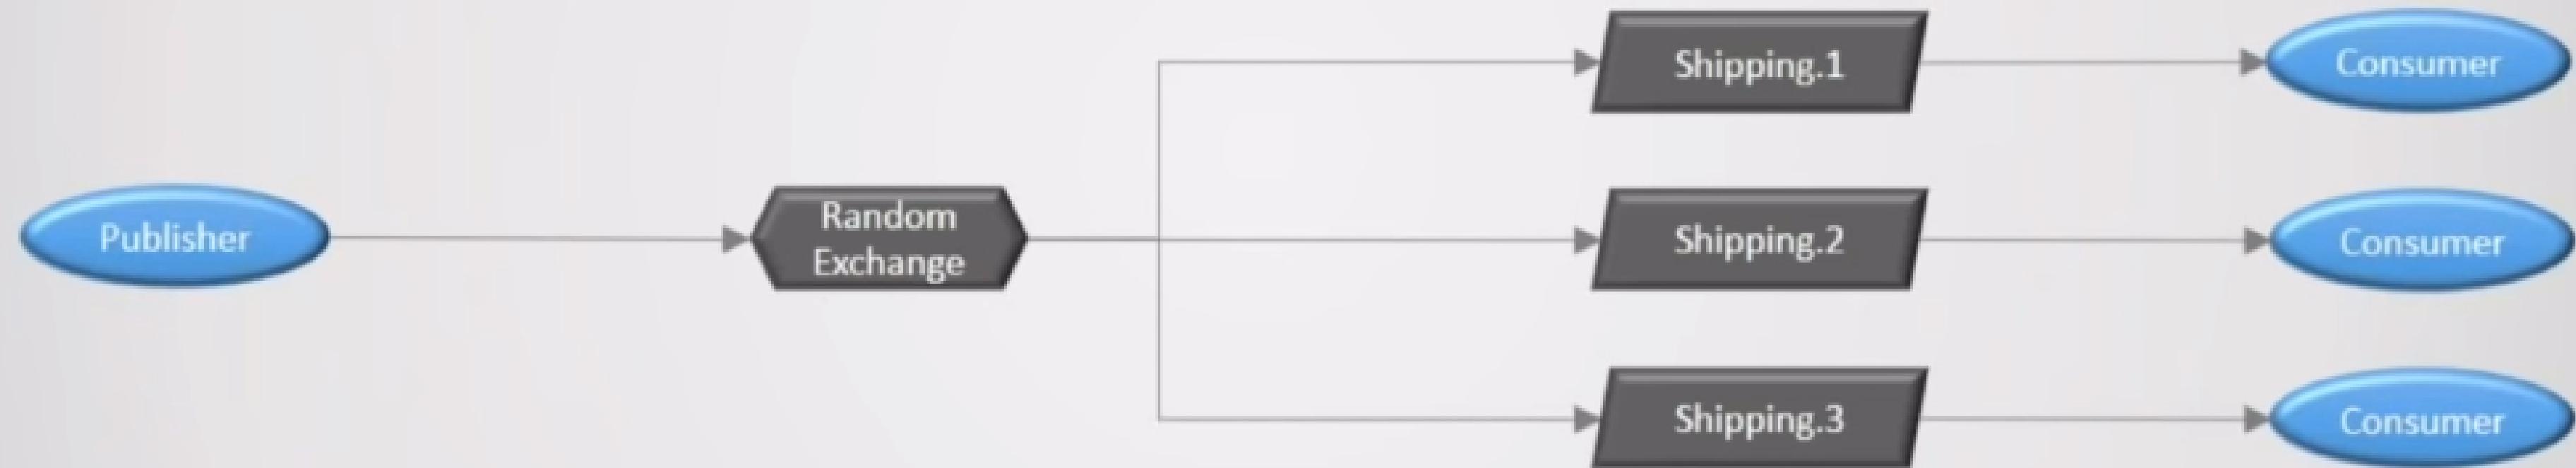
Each Shipping consumer receives a subset of the messages.



# Scaling Out Queues

## Random Exchange

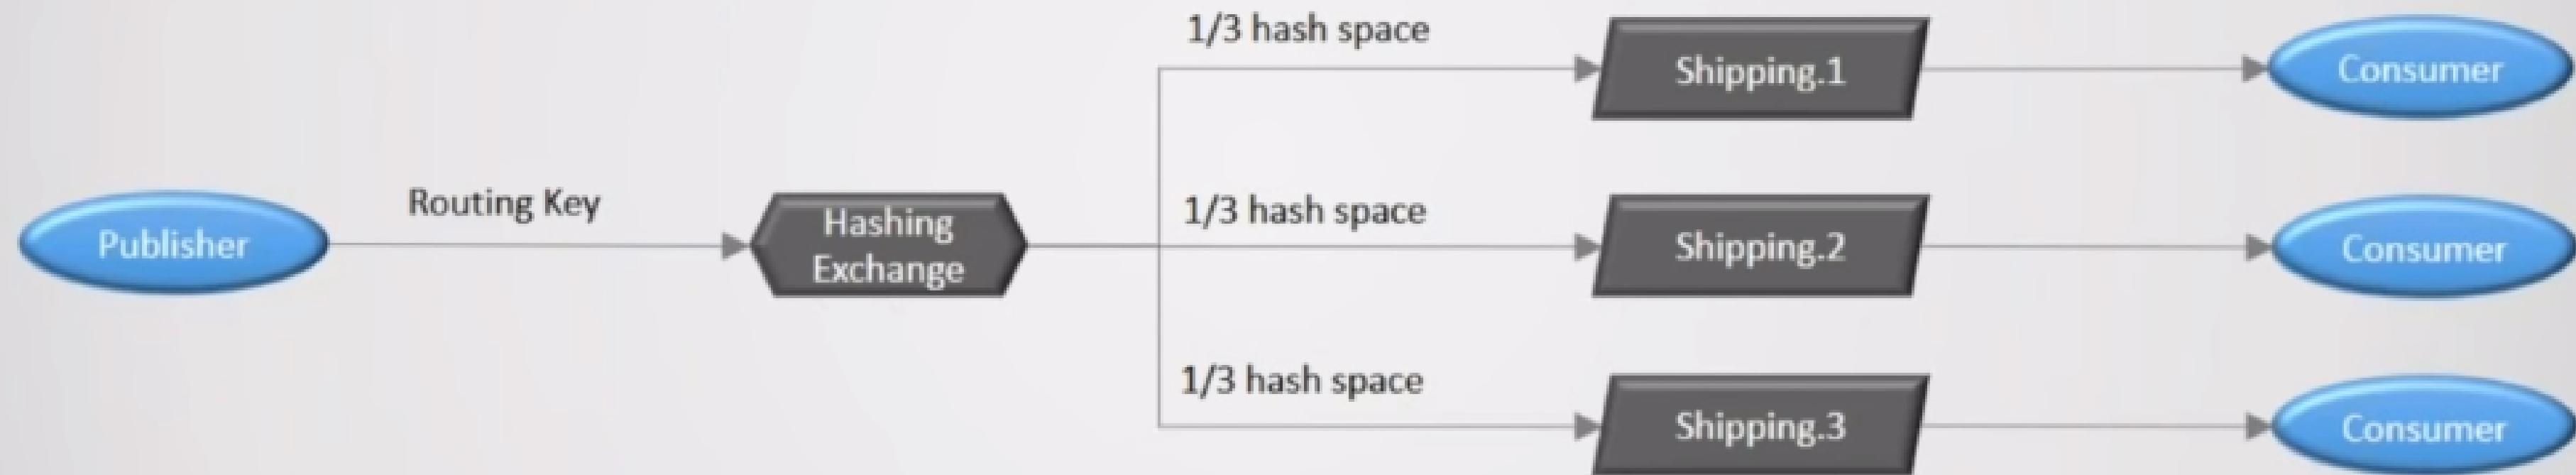
Load balance messages across queues randomly



# Scaling Out Queues With Data Locality

## Consistent Hashing Exchange

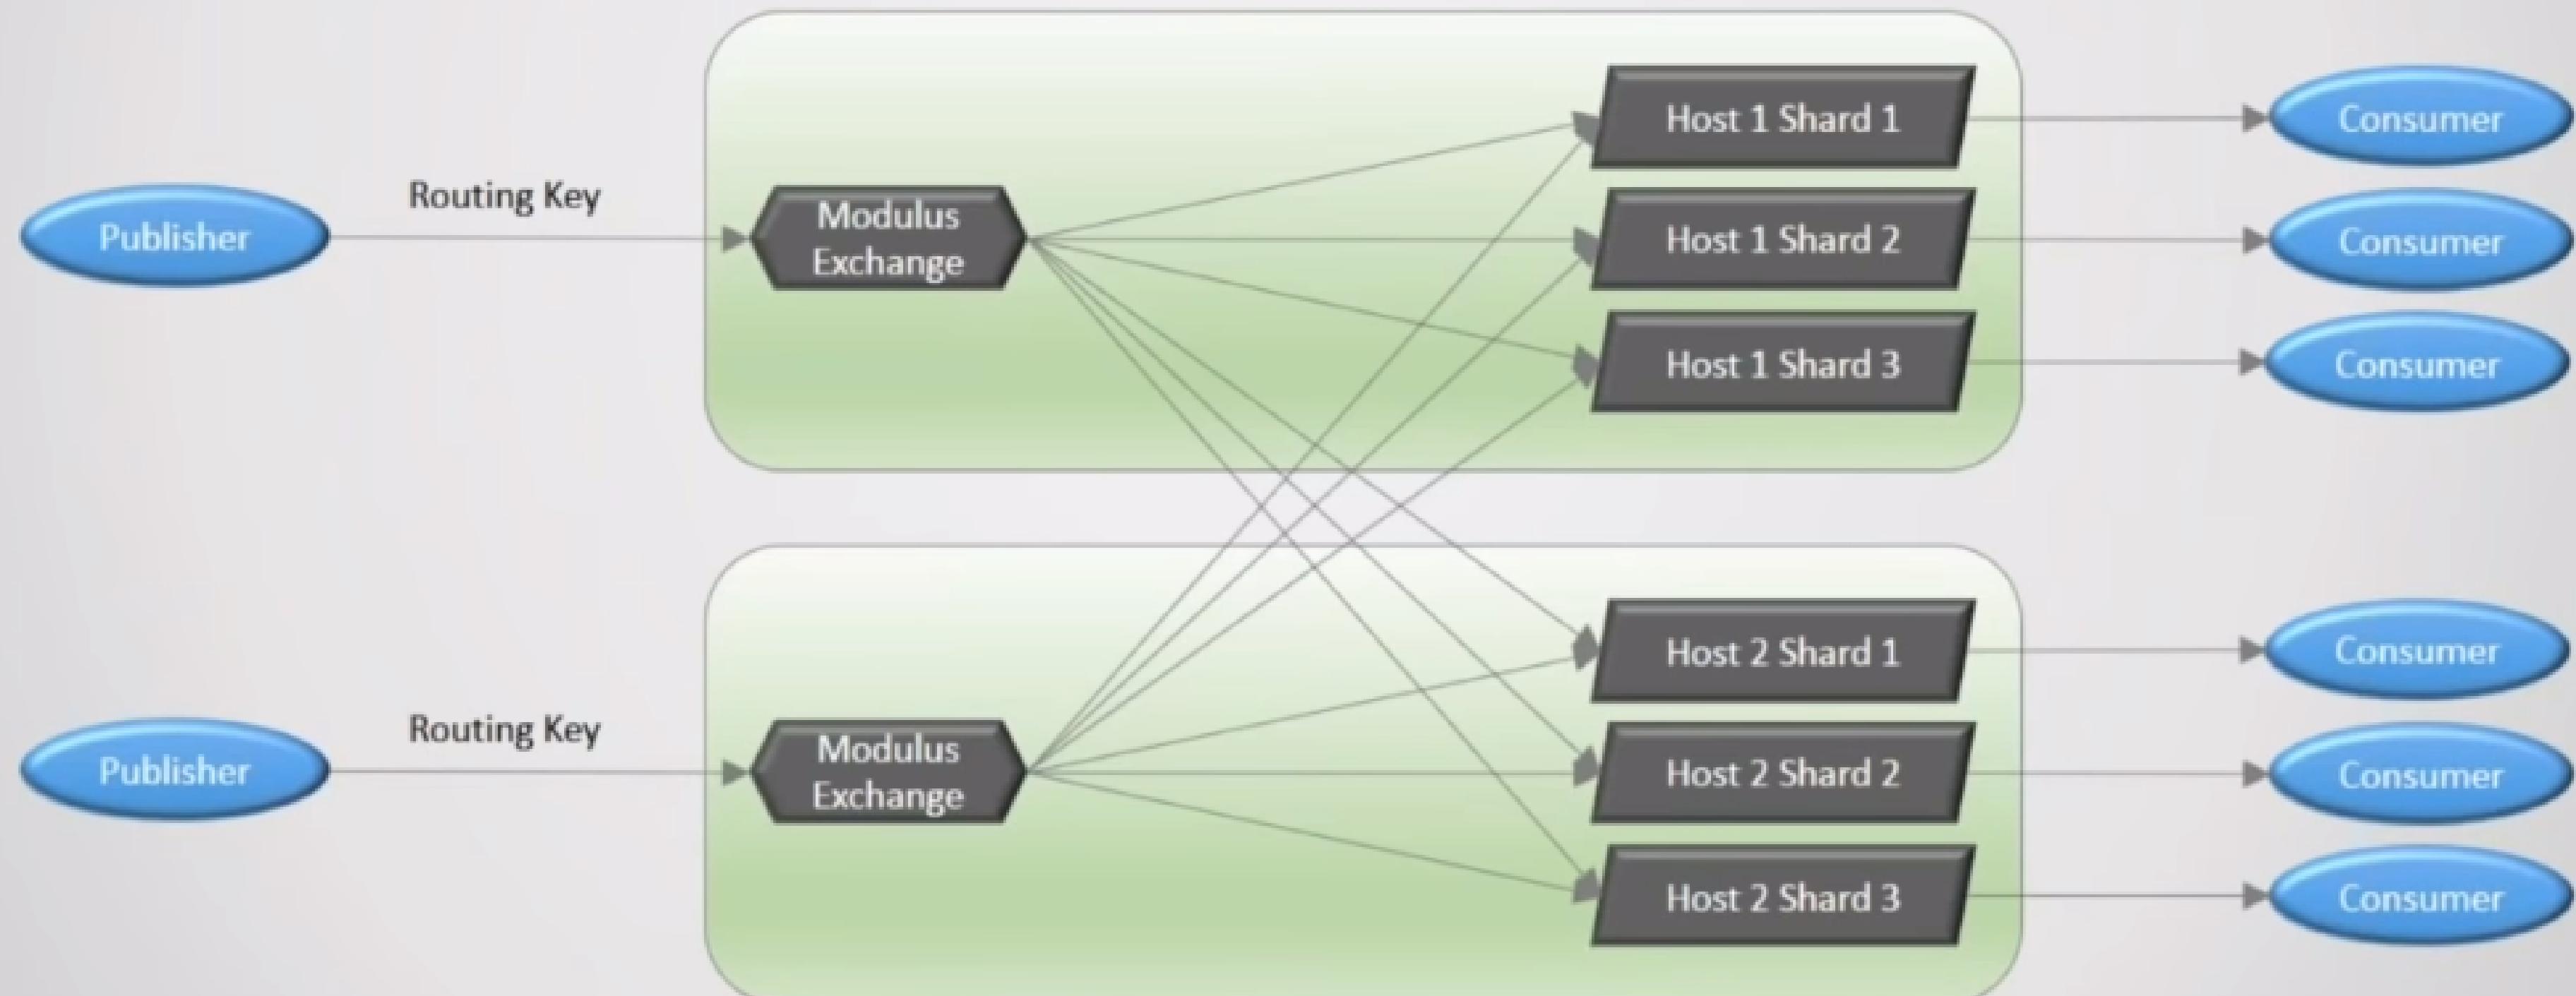
Partition messages across queues via hashing function over routing key, message header or message property



# Scaling Out Queues With Data Locality

## Sharding Plugin and Modulus Exchange

Partition messages across queues over multiple hosts via hashing function on the routing key.



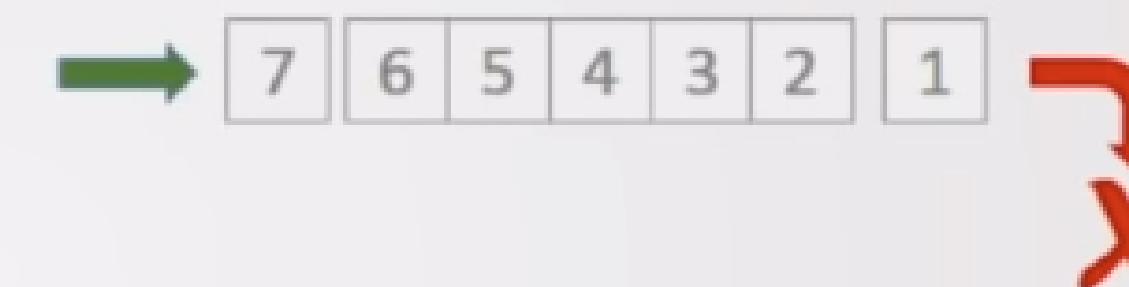
# Queue Limits and Deadletter Exchanges

Apply limits to queues

- Length
- Size
- Time limits (TTL)

Eject messages from a queue when:

- A queue size/length limit reached.



- A message has spent longer than the TTL time limit in the queue

Route to a deadletter exchange to avoid message loss.



# Ephemeral Exchanges and Queues

## Lazy Queues

## Priority Queues

### Ephemeral

- Auto-Delete Queue
- Queue Expiration (TTL)
- Exclusive Queues
- Auto-Delete Exchanges

### Lazy Queues

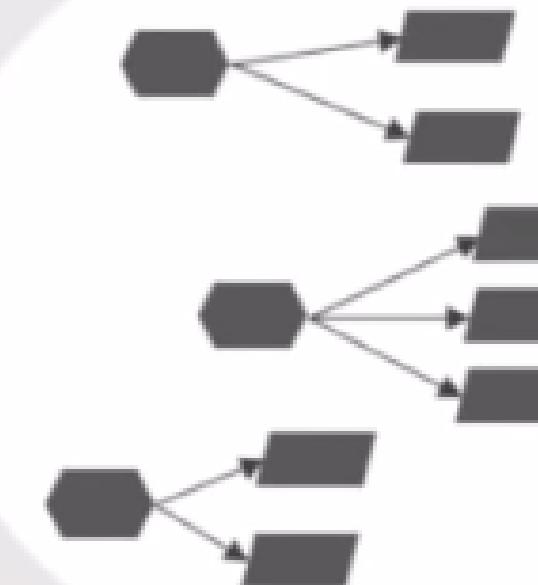
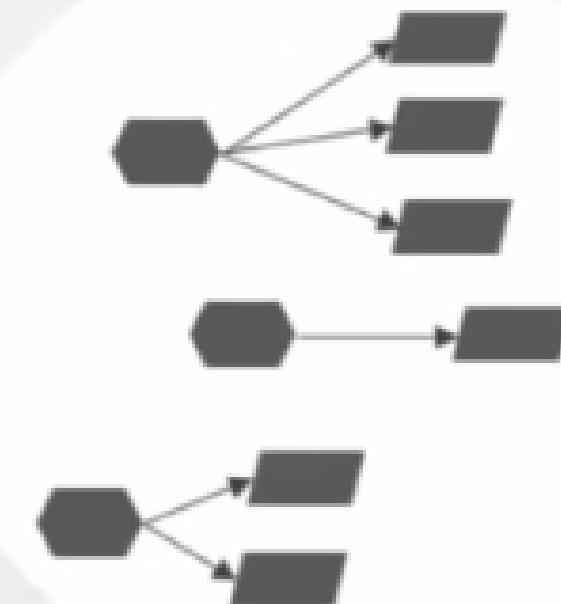
- Memory optimized queues.

### Priority Queues

- No longer FIFO.
- Publisher sets priority on messages
- Priority Queue moves higher priority messages ahead of lower
- Drawbacks – blocked low priority messages, low priority can eject high priority when queue is full

# Virtual Hosts

- Allows Multi-Tenant architecture
- Isolation of groups of exchanges, queues and users
- No routing between two virtual hosts



# Decoupling of publishers from subscribers

## Endpoint Decoupling

**RabbitMQ:** The endpoint we publish to is decoupled from the endpoint we consume from.

**Kafka:** The endpoint we publish to is the same as we subscribe from.

## Temporal Decoupling

**RabbitMQ:** Consumers read a message once, most likely close to the time of publishing.

**Kafka:** Consumers are decoupled temporally from the publisher. Consumers can read messages multiple times, and go back in history to read old messages.

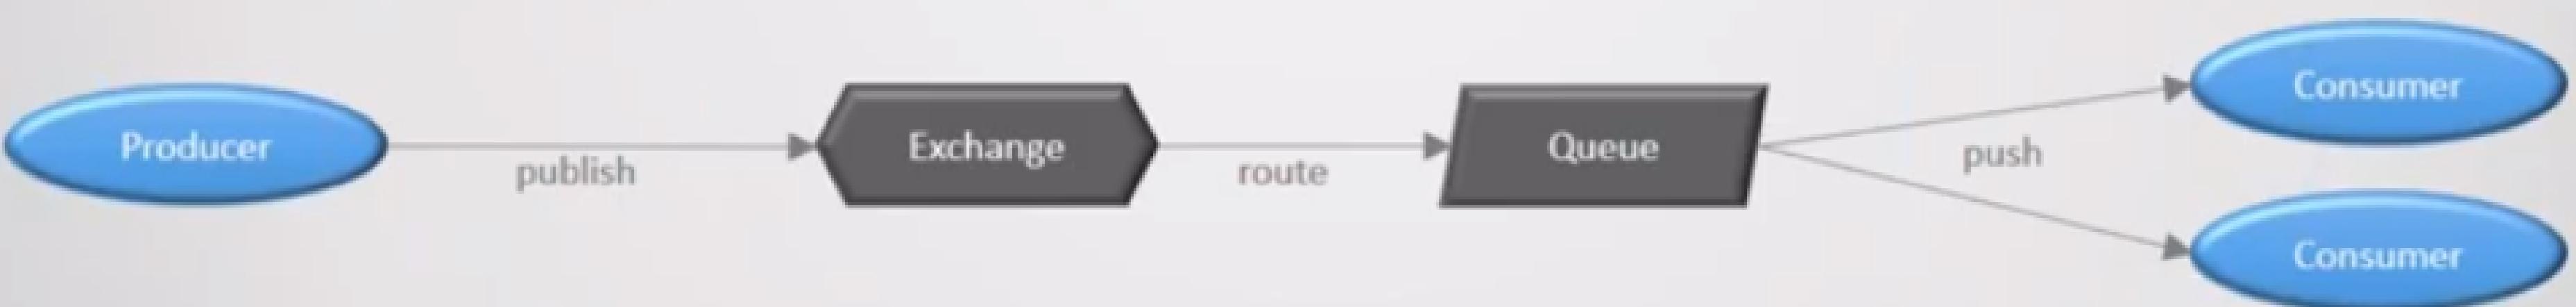
# RabbitMQ - Push Model

## Producer Publish

- Send messages one at a time

## Consumer Push

- Long-lived TCP connection
- Consumer registers interest in queues
- Broker pushes messages down connection in real-time



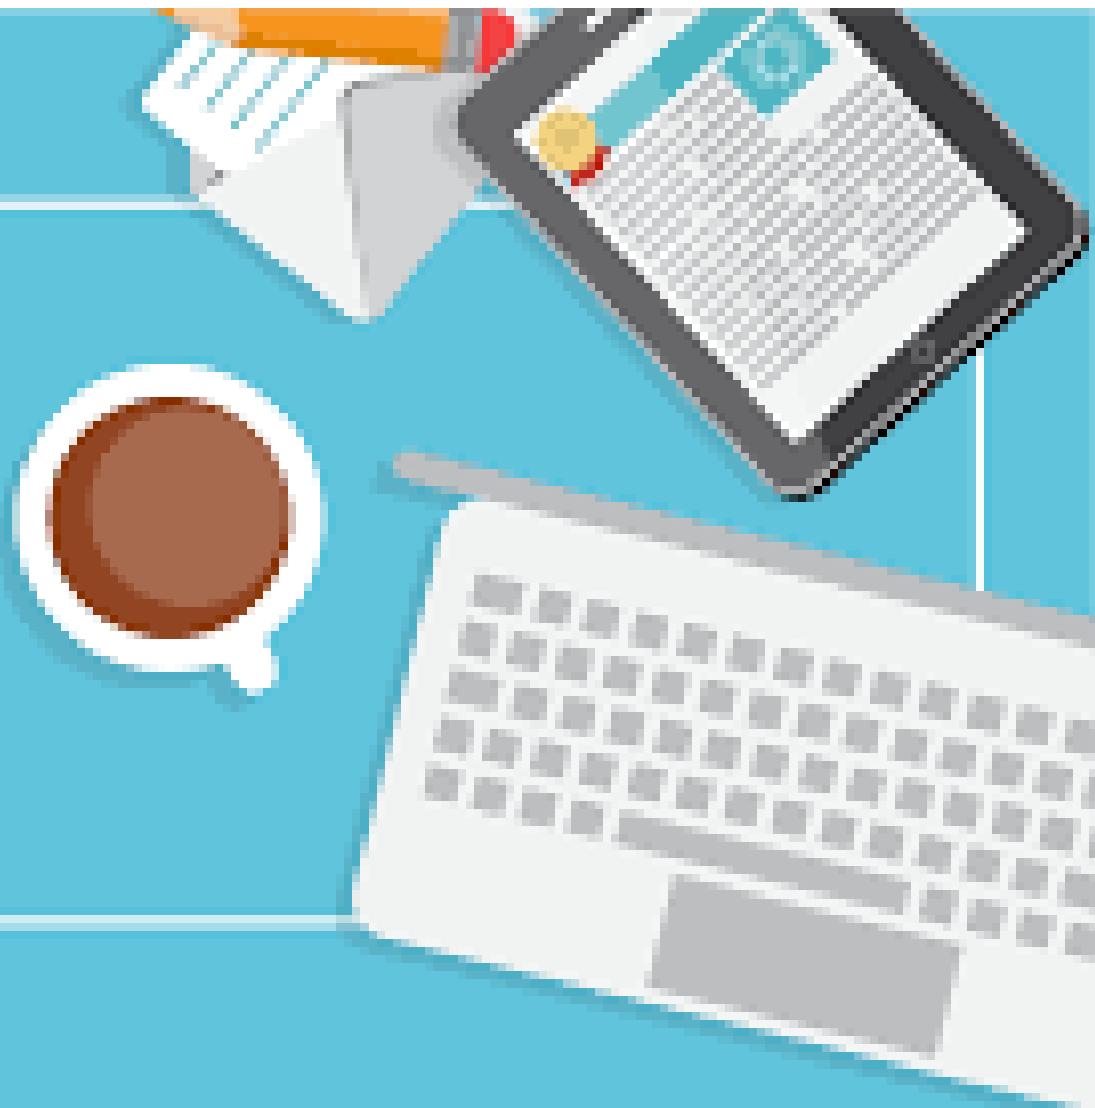
## RabbitMQ – Why Push?

The push model allows RabbitMQ to:

- Offer low latency messaging.
- Evenly distribute messages across competing consumers.
- Keep processing order closer to delivery order in the face of competing consumers.

A push model requires Back-Pressure: Consumer Prefetch.

# CASE STUDY

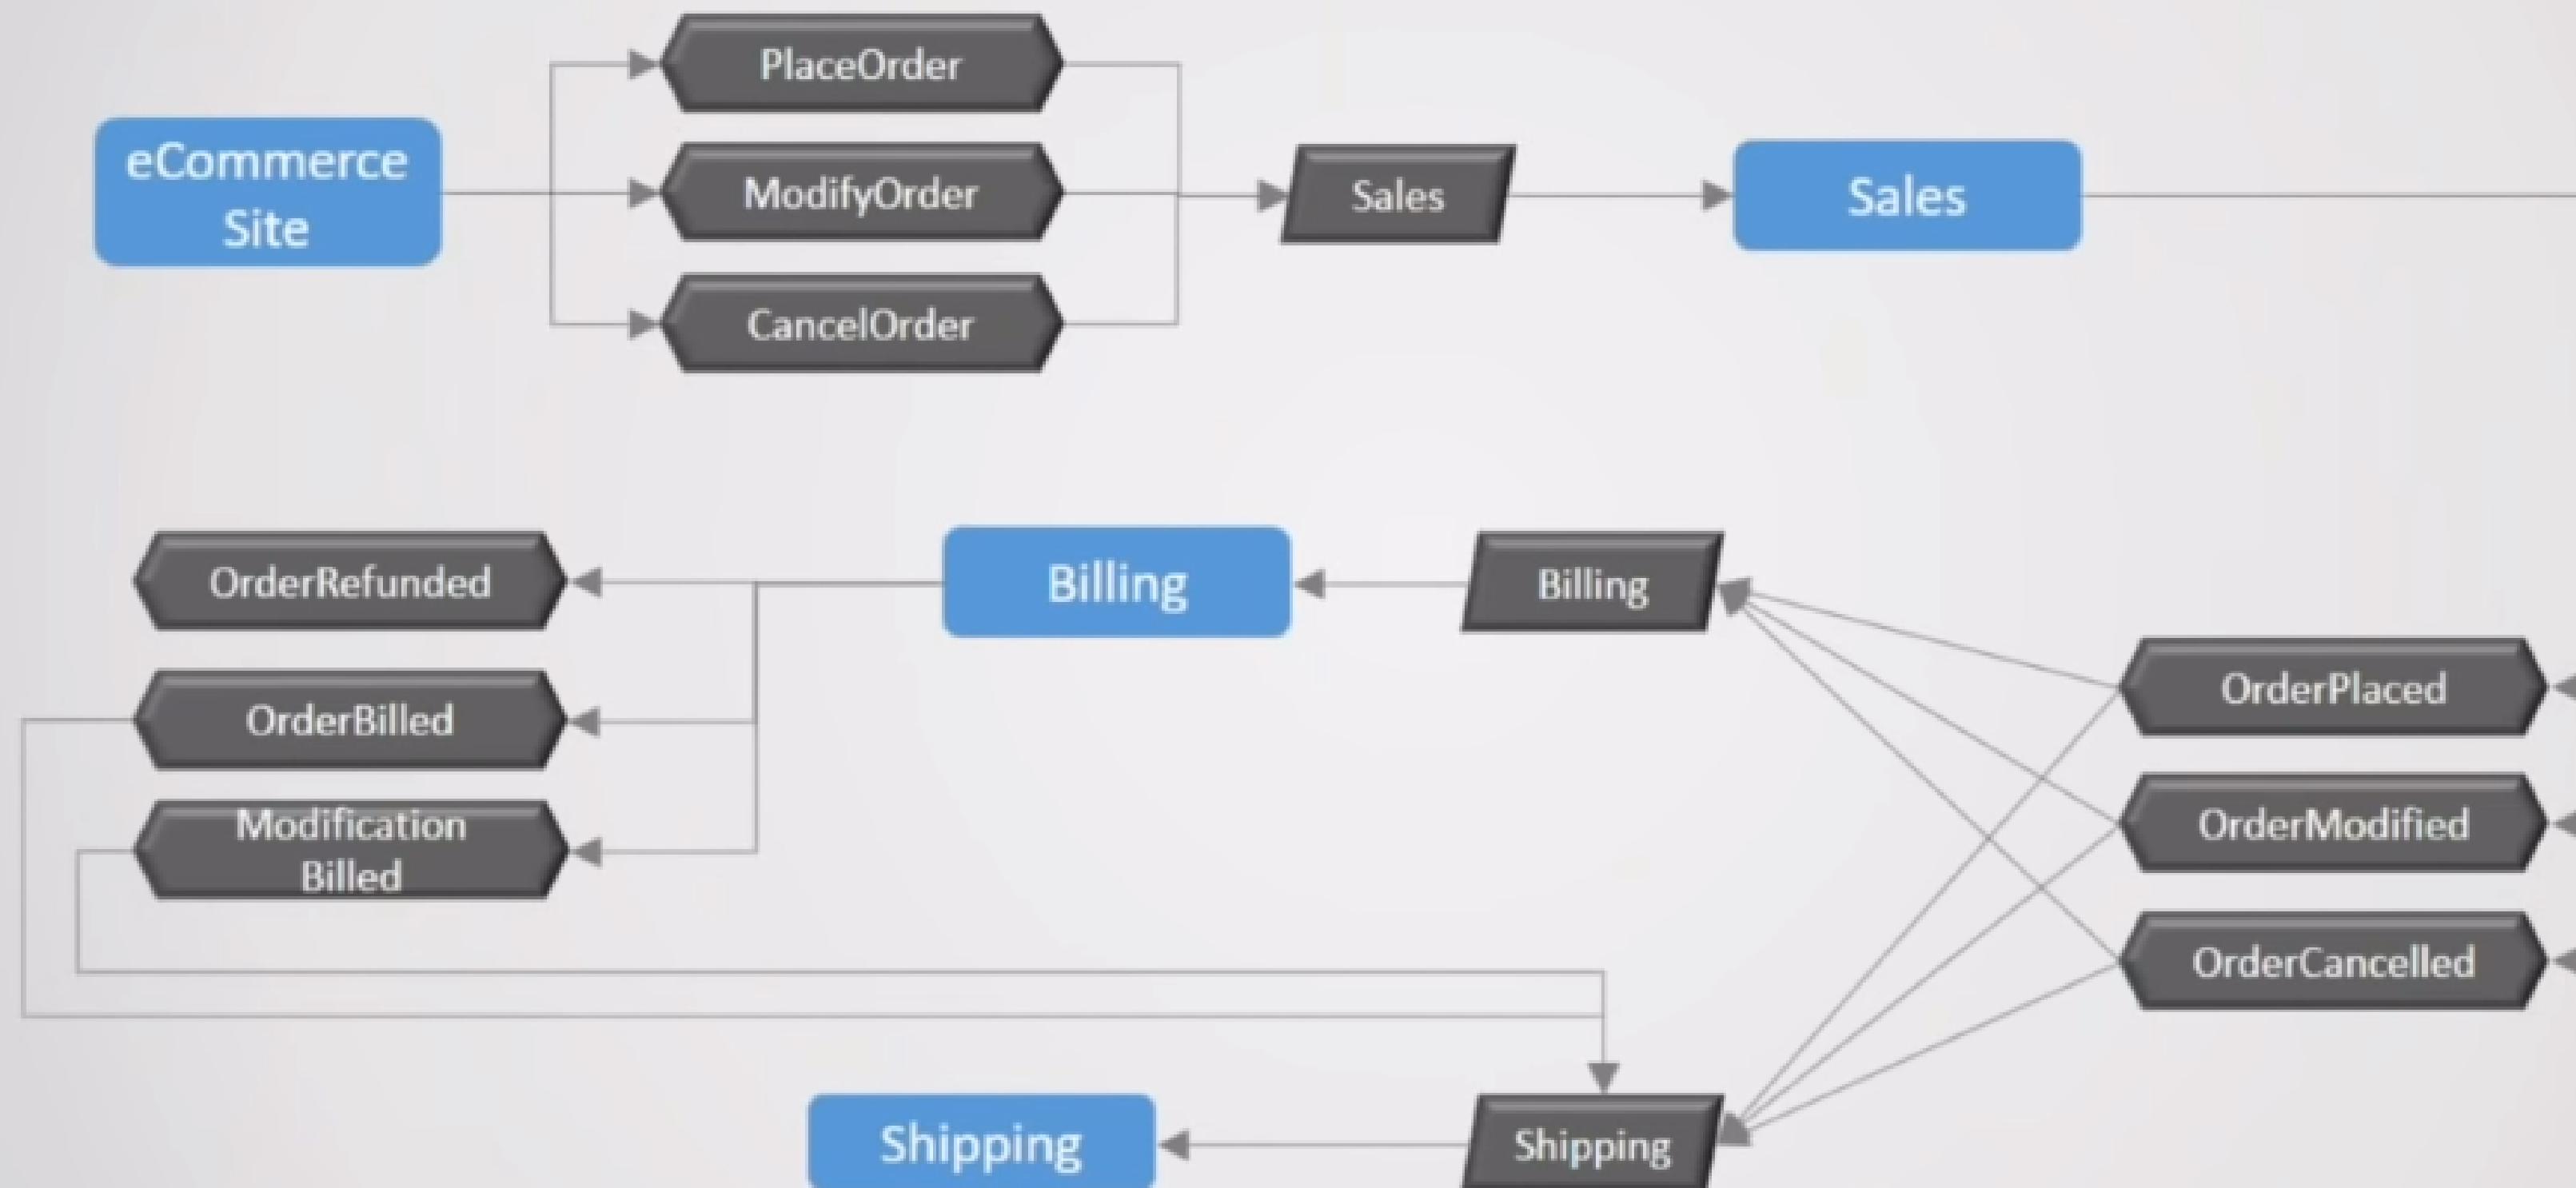


# Example Scenario



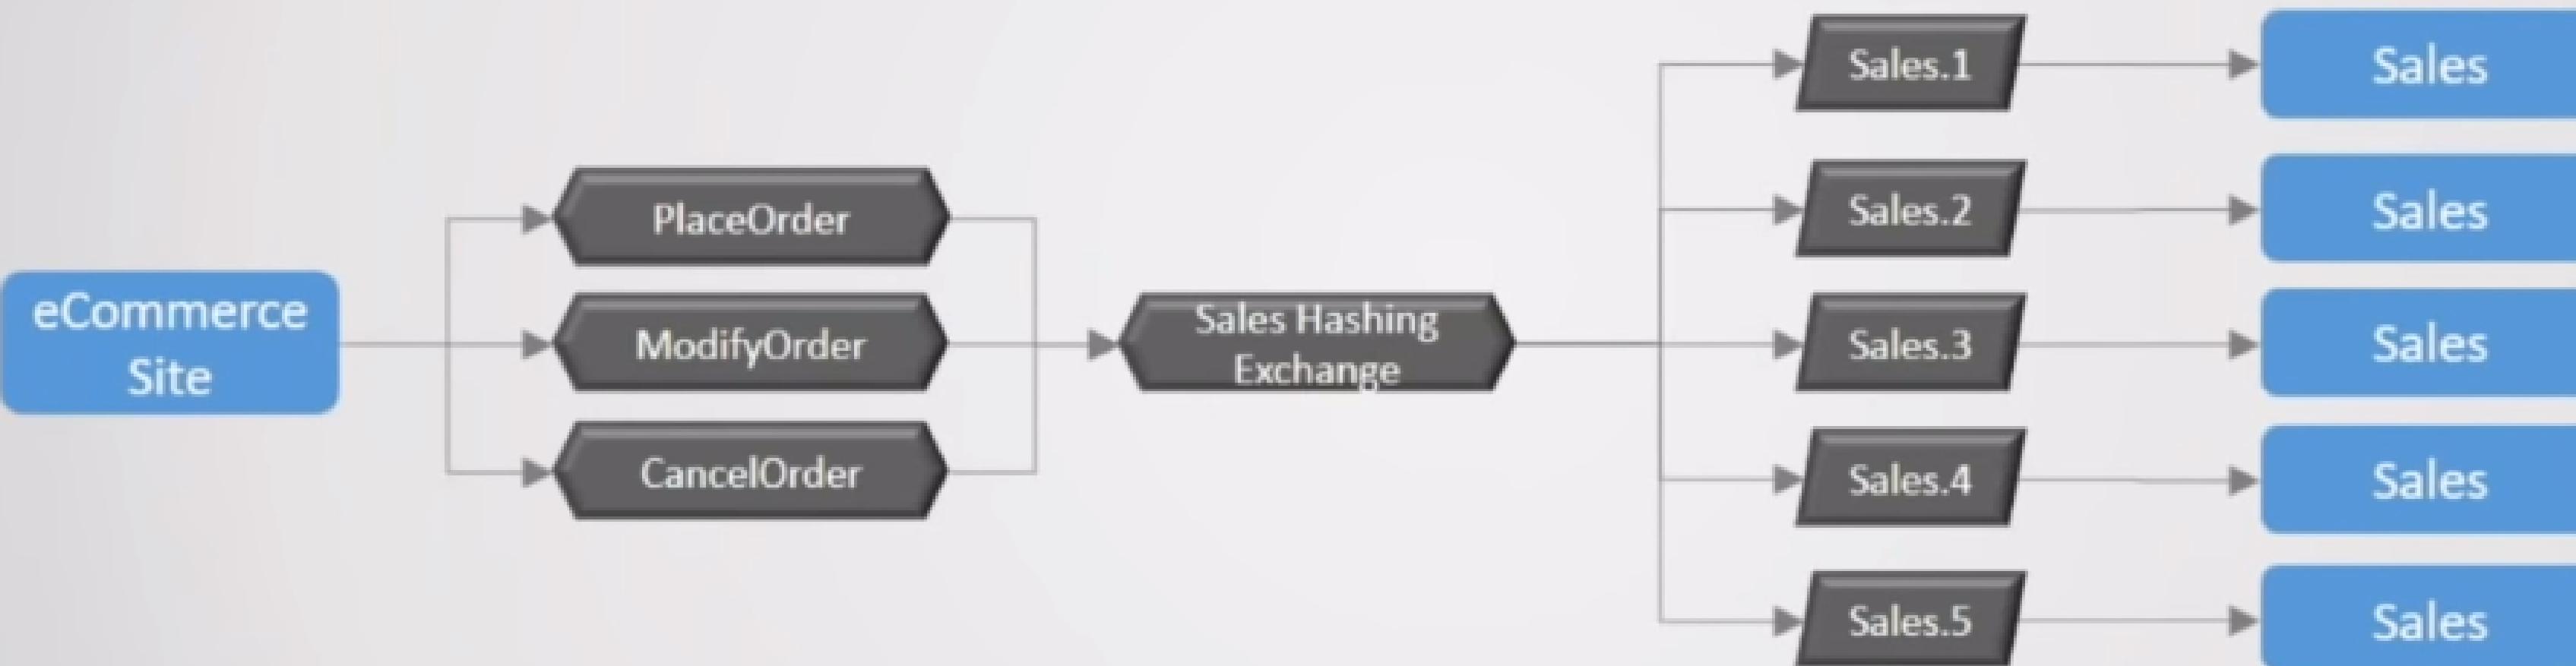
# Example Scenario - RabbitMQ Fanout Exchanges

Fanout Exchange per Event, Single Queue per Application



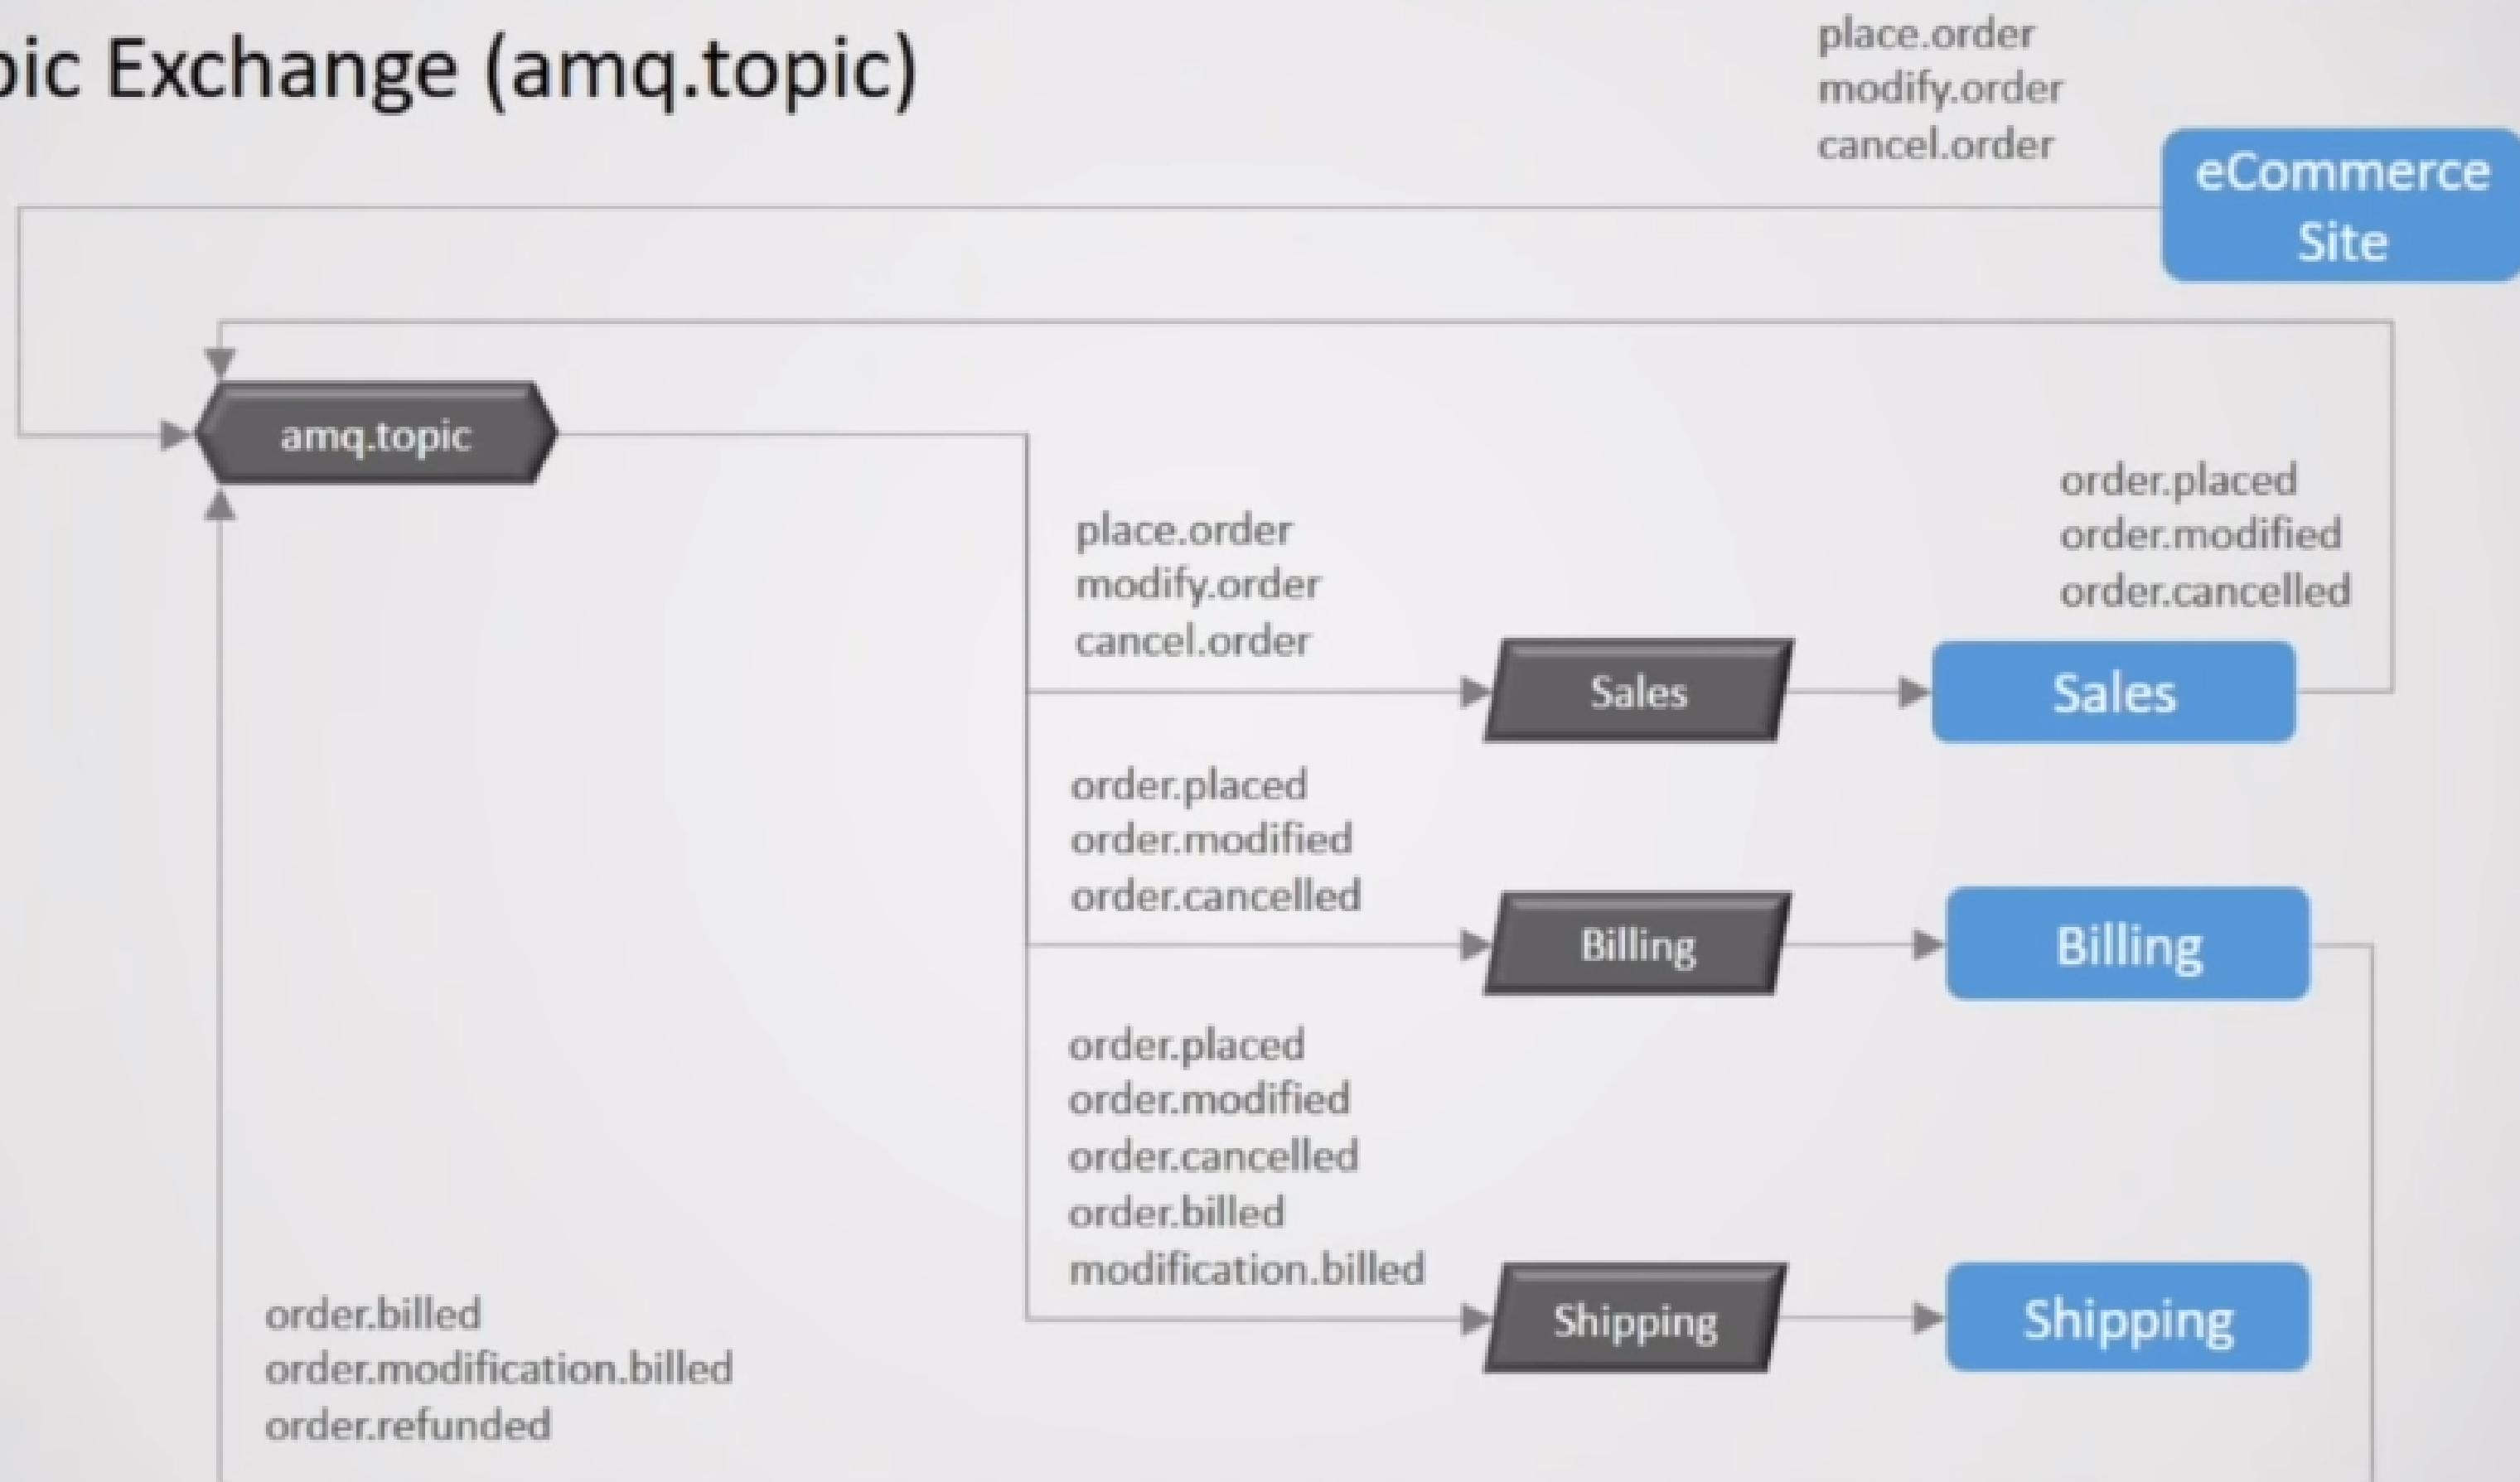
# Example Scenario - RabbitMQ Fanout Exchanges

Scaling our Single Queue, Maintaining Relative Ordering

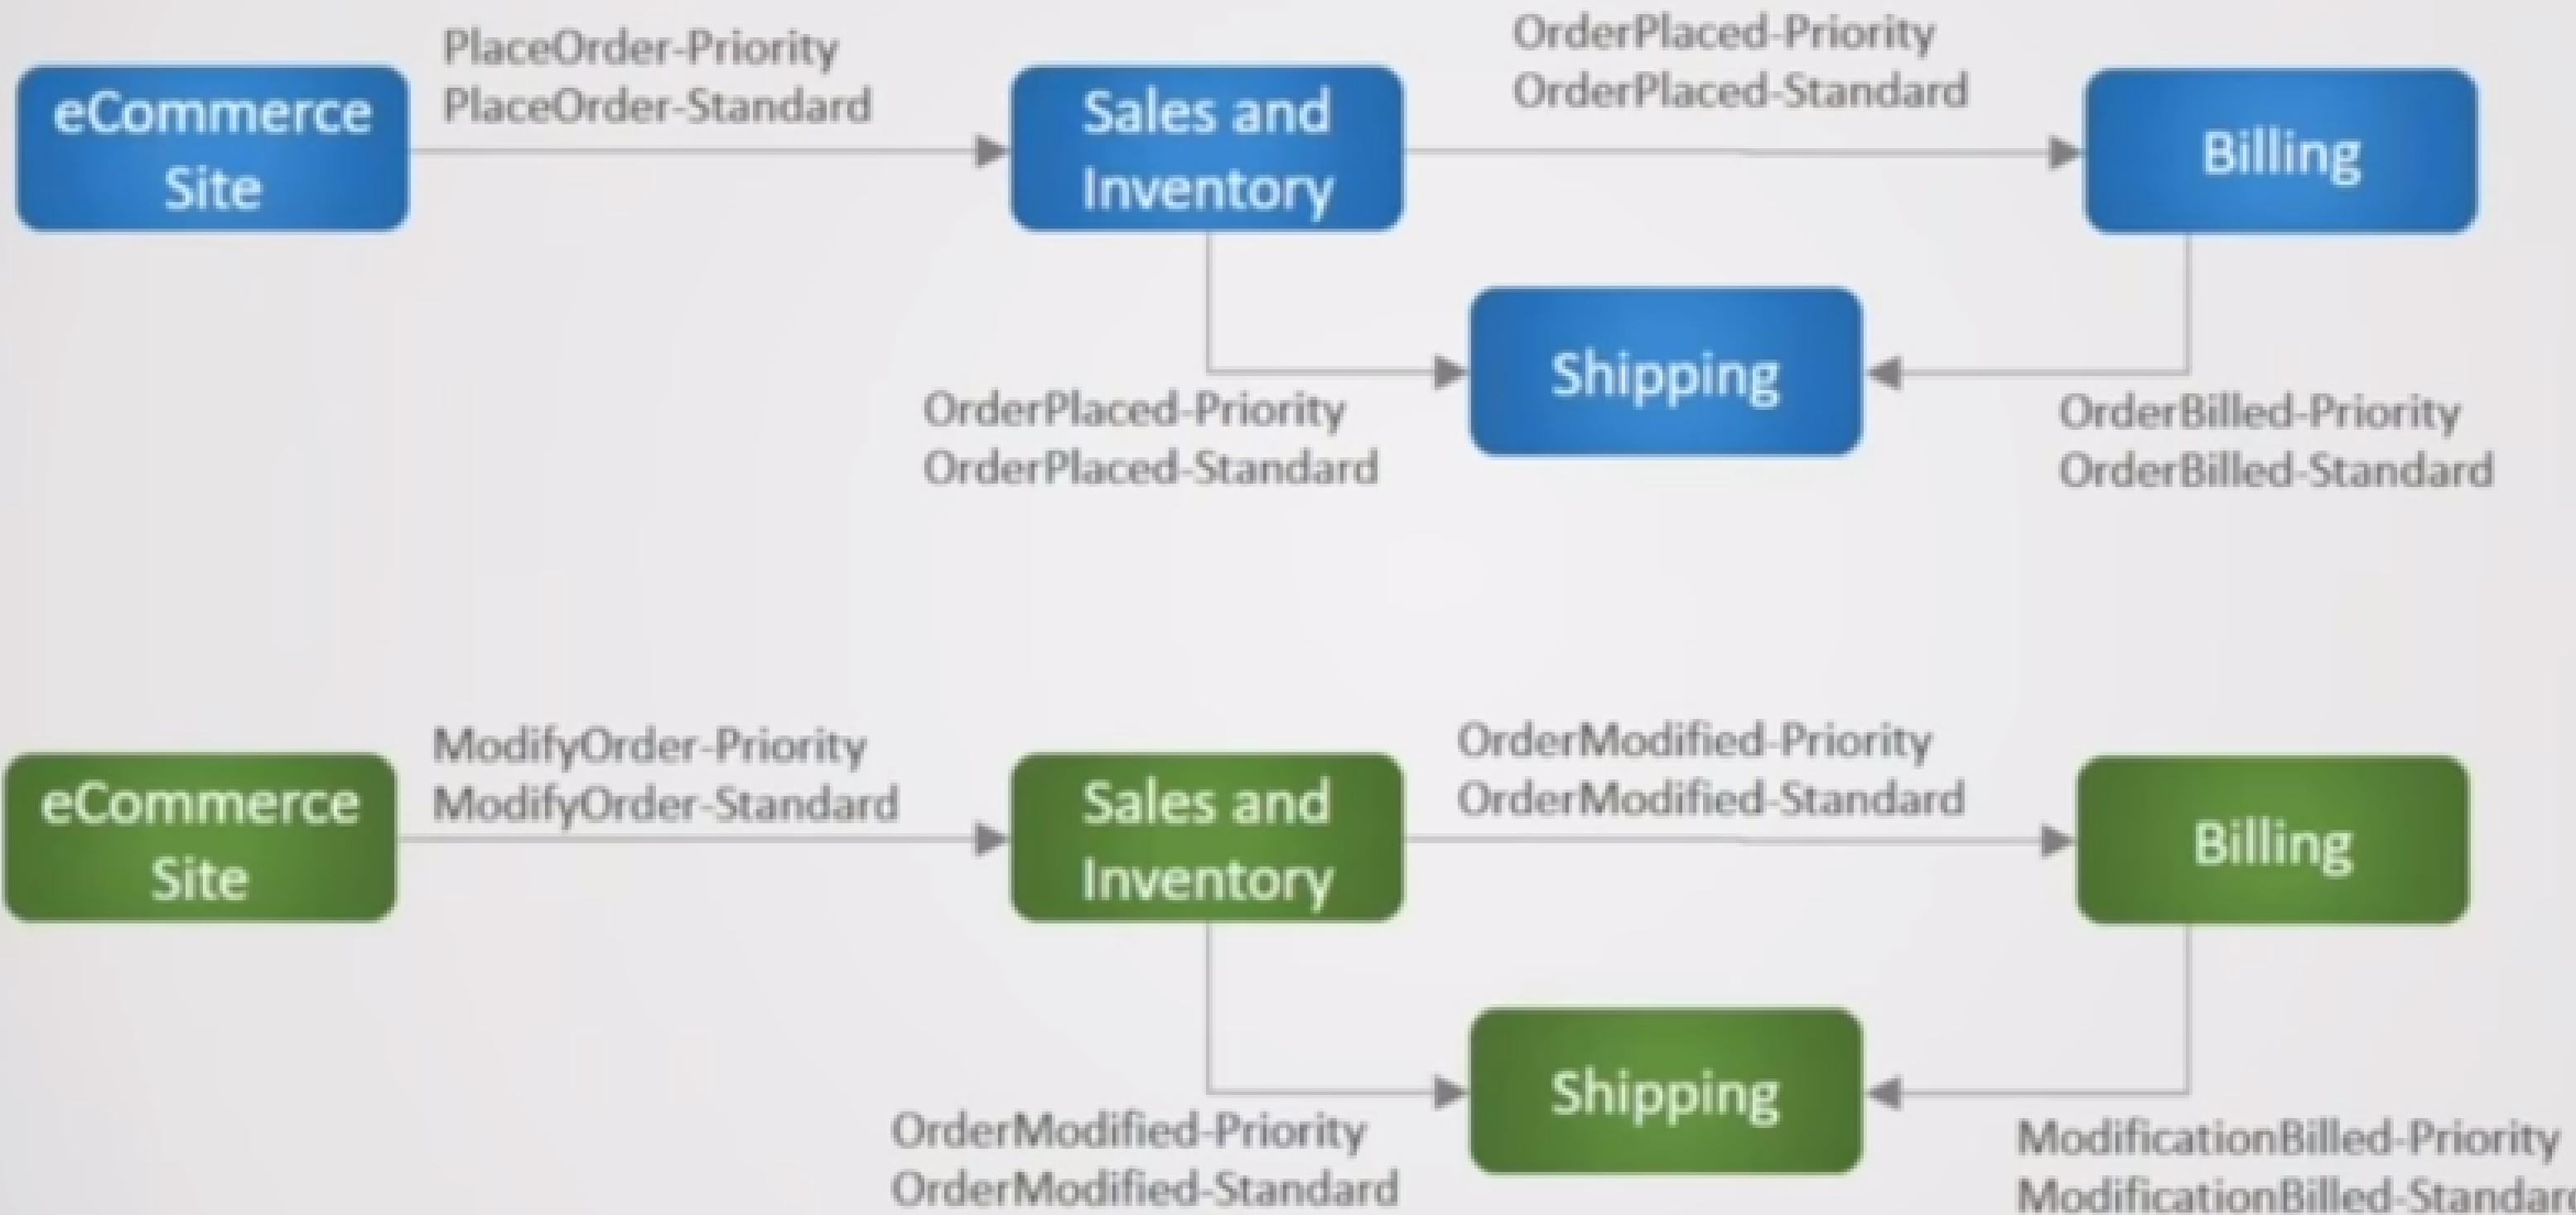


# Example Scenario - RabbitMQ Topic Exchanges

## Single Topic Exchange (amq.topic)

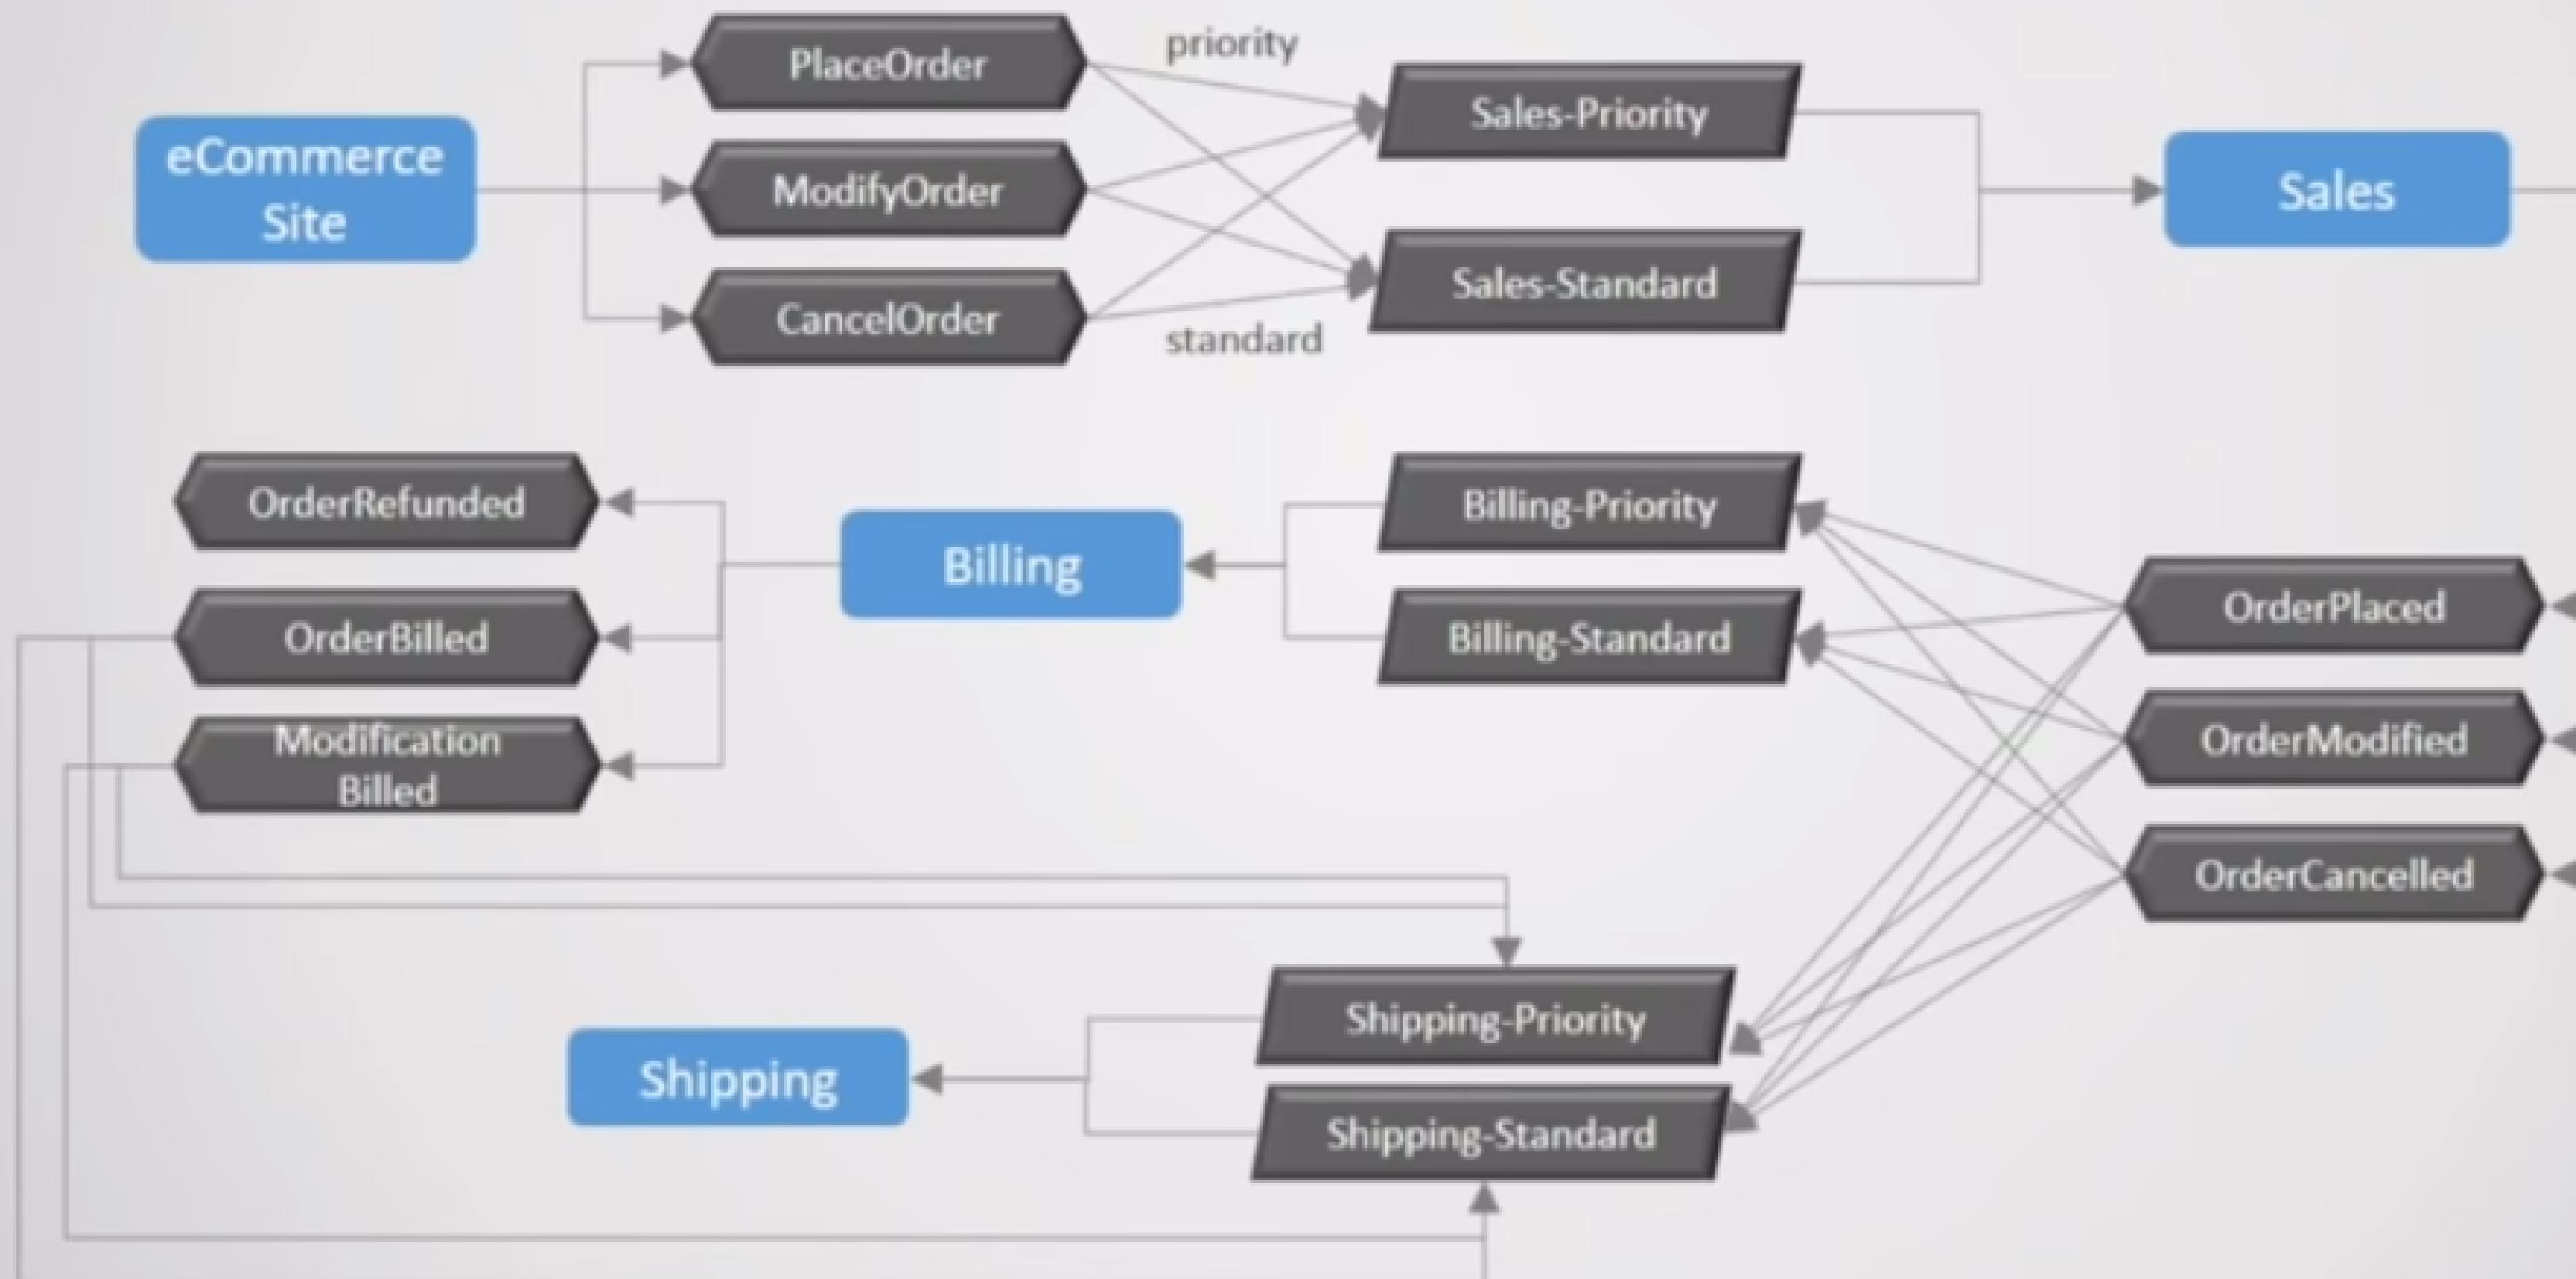


# Modifying the Scenario - Class of Service



# Scenario #2 - RabbitMQ Topic Exchanges

Direct Exchange per Event, Routing Key with Class of Service,  
Queue per Class of Service

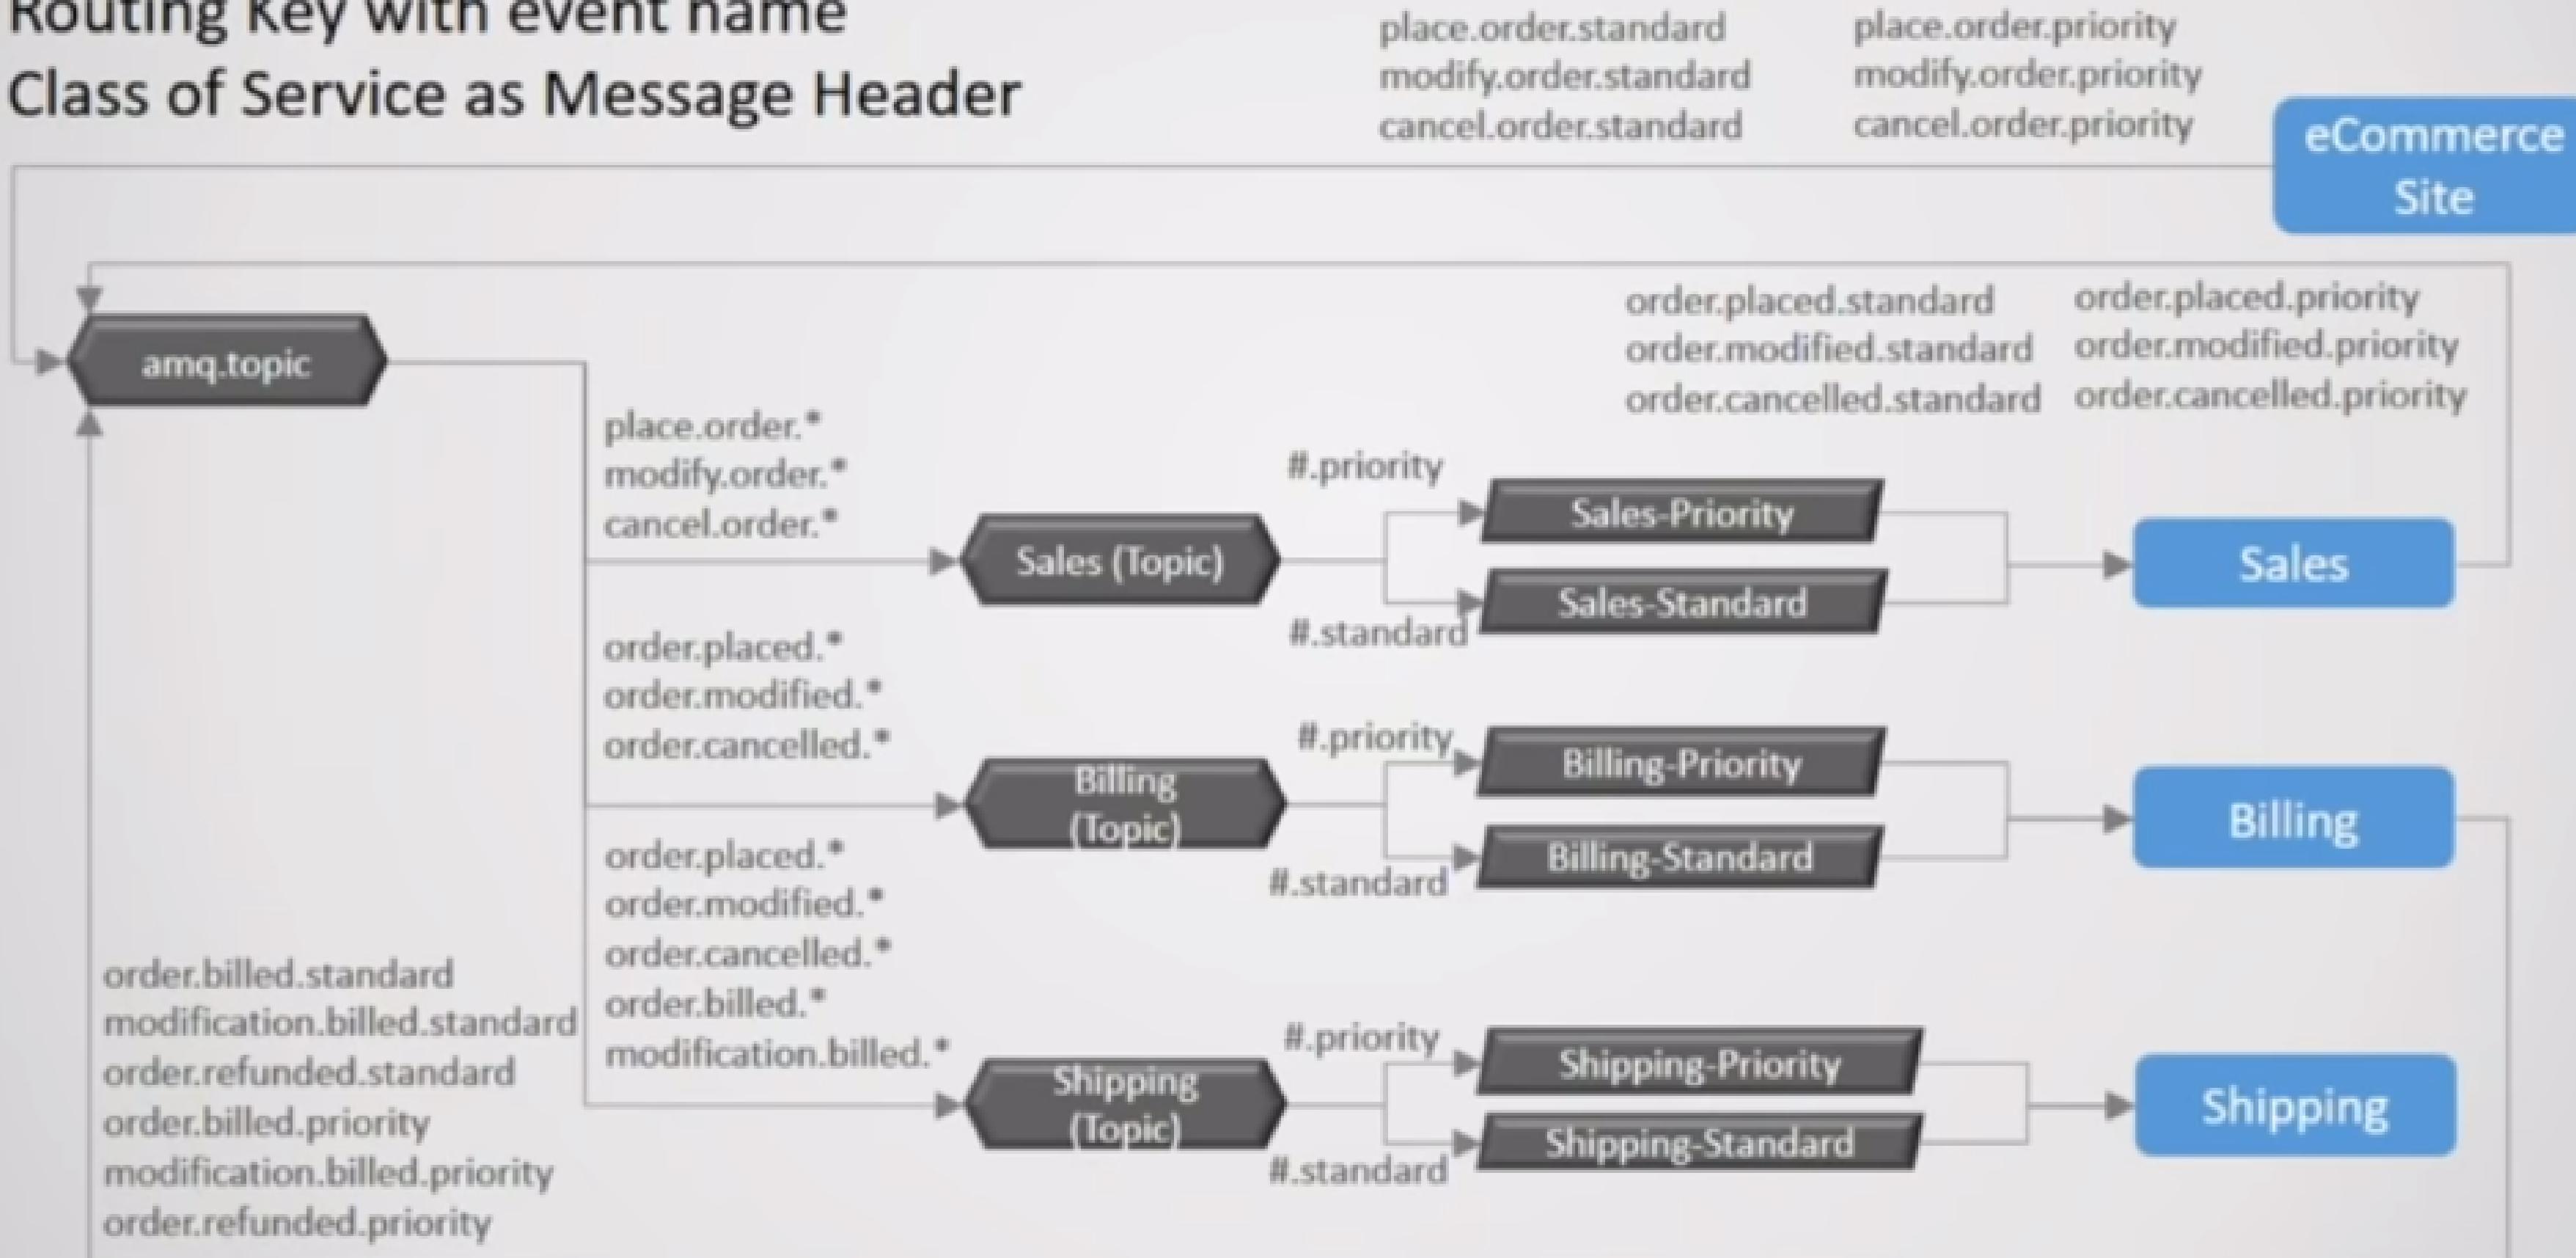


# Scenario #2 - RabbitMQ Two Layered Topic Exchanges

Single Topic Exchange (amq.topic)

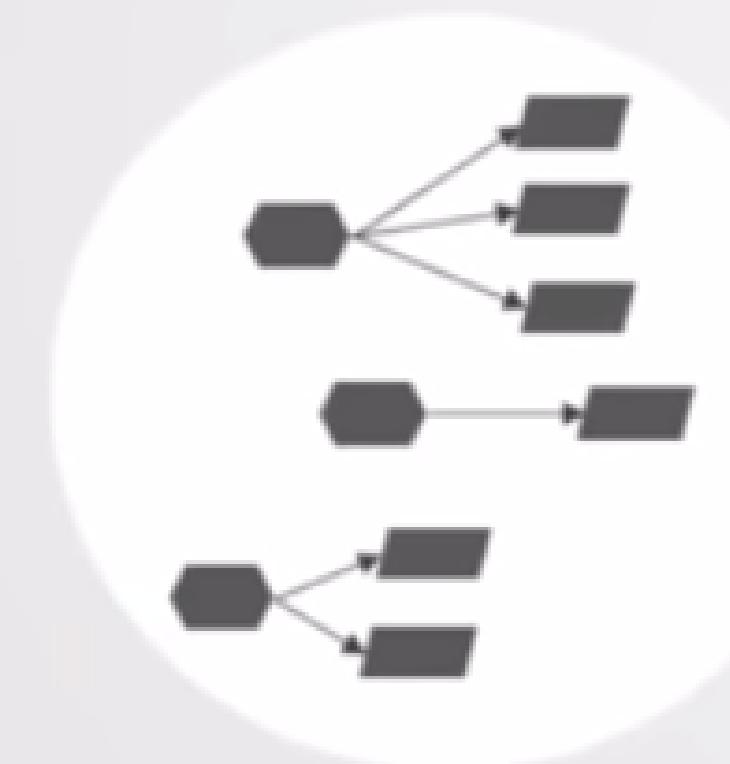
Routing Key with event name

Class of Service as Message Header



# Scenario #2 - RabbitMQ Topic Exchanges

Replicate preferred routing topology without Class of Service,  
in two separate virtual hosts



Priority VHost



Standard VHost