

April , 2021

RabbitMQ

Reliable messaging



- Message Delivery Guarantees
- Message Ordering Guarantees
- Message Durability
- High Availability

Message Acknowledgement Protocols

At-most-once
At-least-once

At-most-once.

This means that a message will never be delivered more than once but messages might be lost.

At-least-once.

This means that we'll never lose a message but a message might end up being delivered to a consumer more than once.

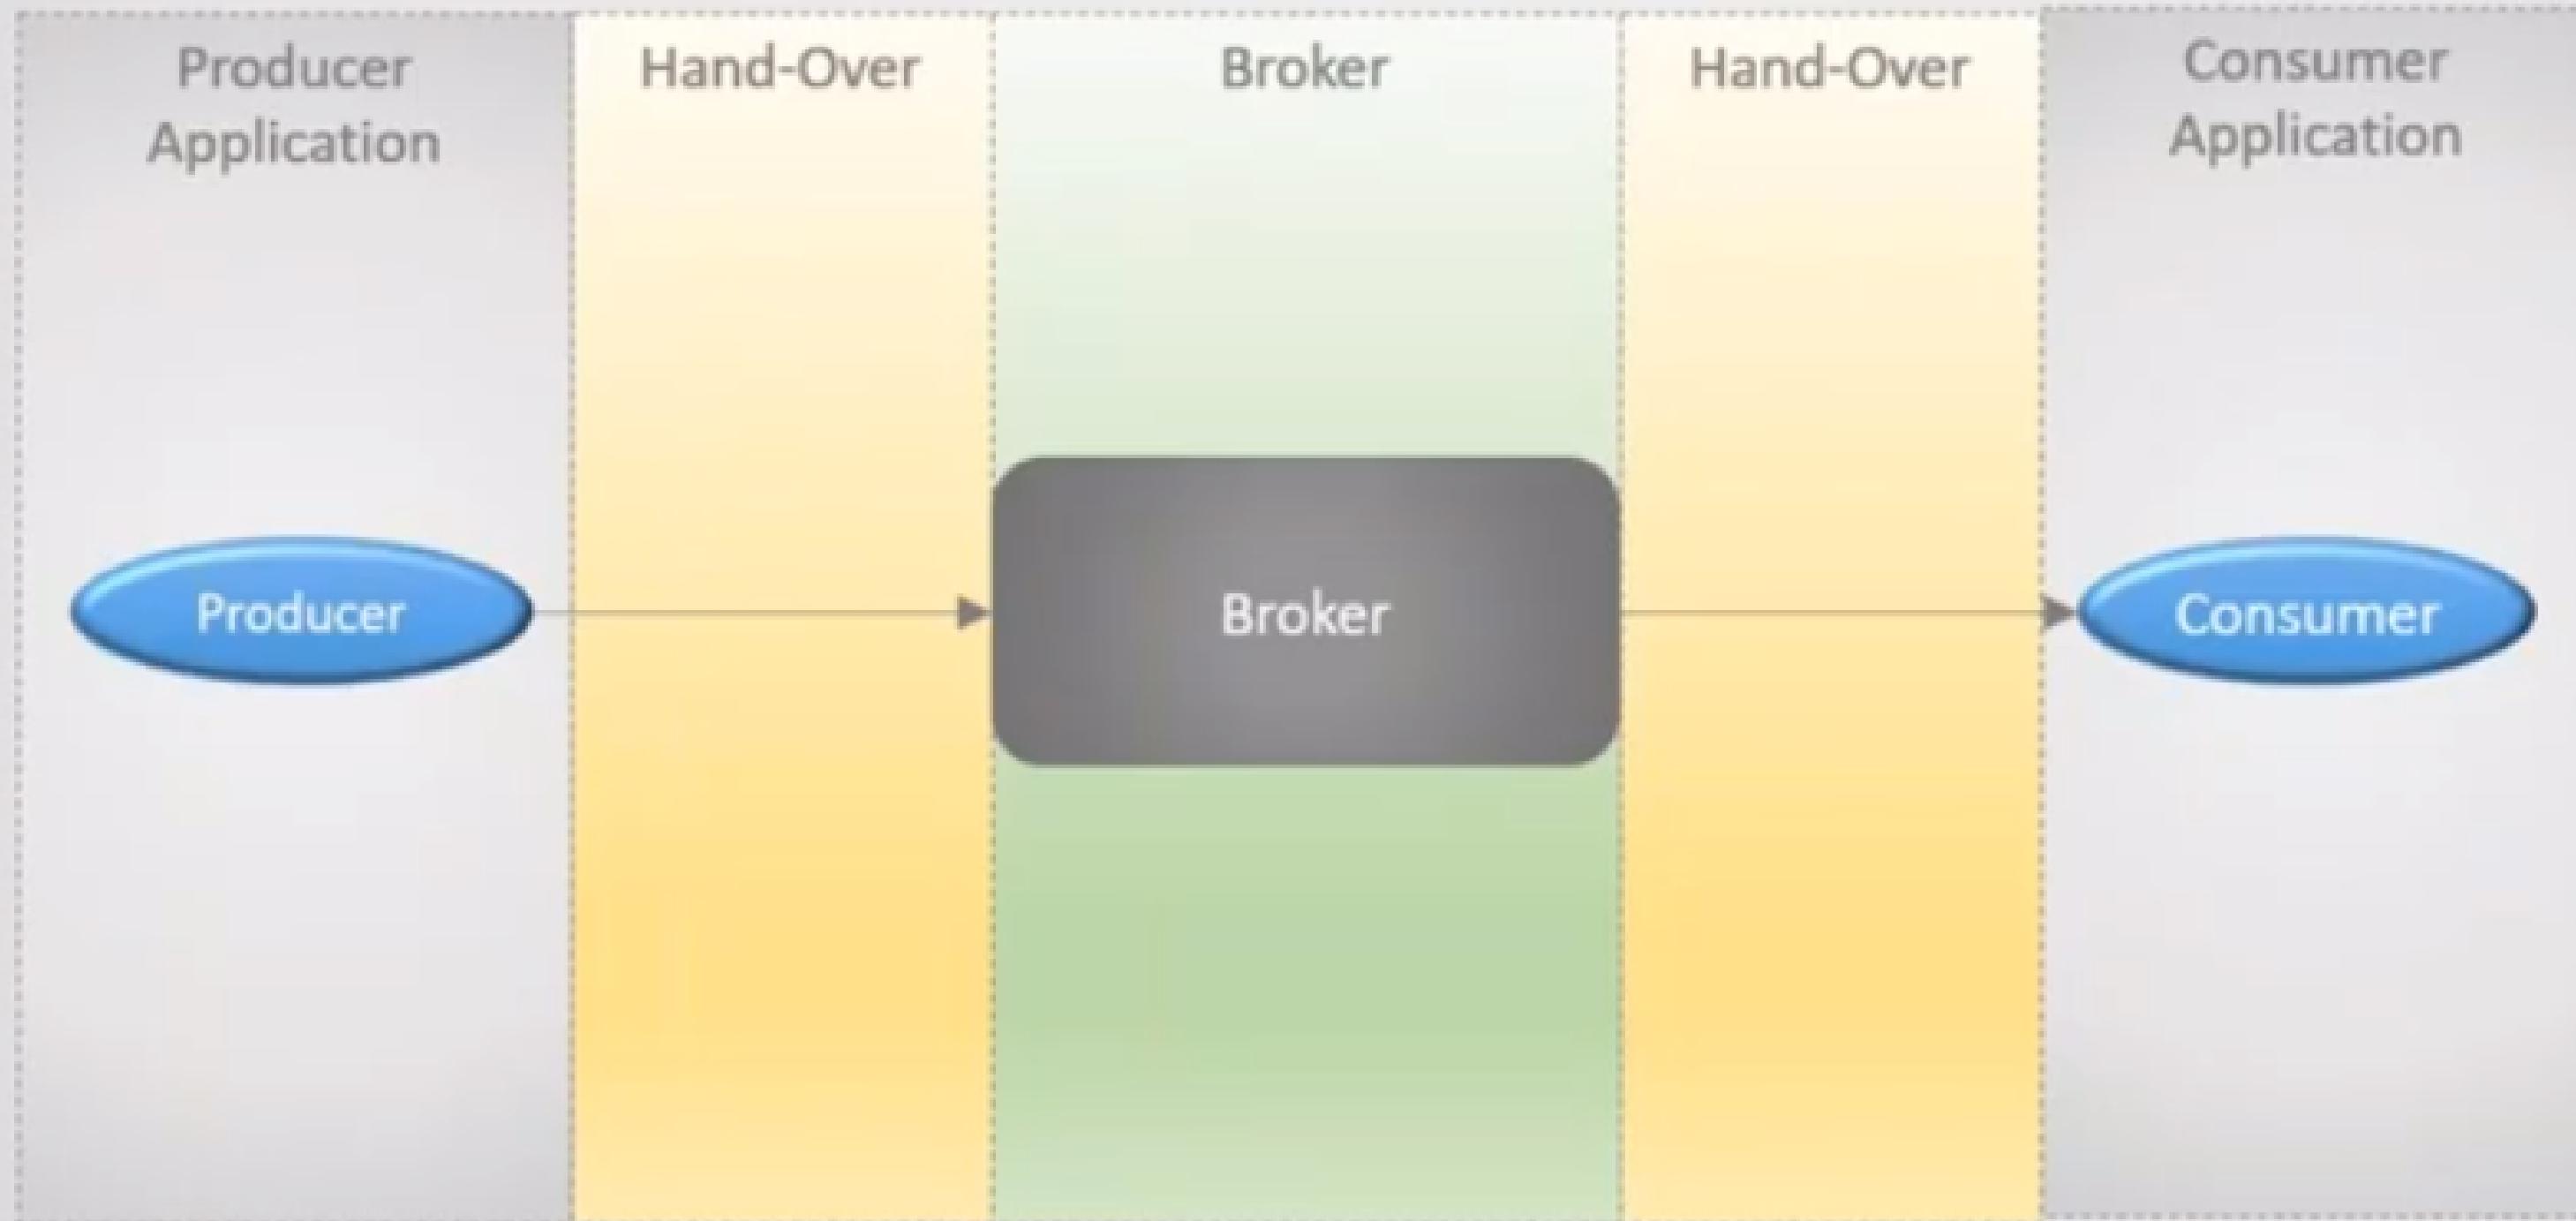
Exactly-once.

The holy grail of messaging. All messages will be delivered exactly one time.

Delivery vs Processing

Delivered twice to be processed once.

Chain of Responsibility



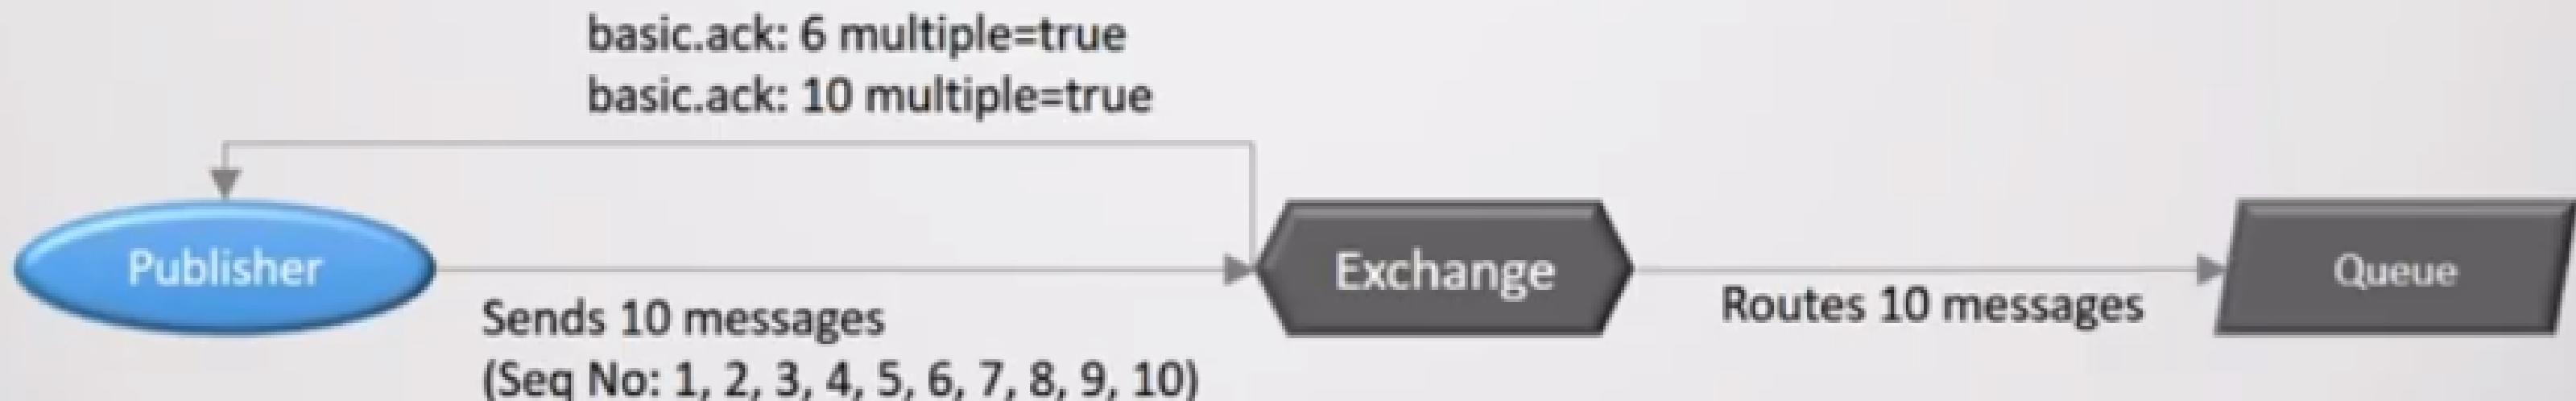
RabbitMQ - Producer Side Acknowledgements

Publisher Confirms

- **basic.ack** (all ok!)
- **basic.nack** (error!)
- **basic.return + basic.ack** (undeliverable!)

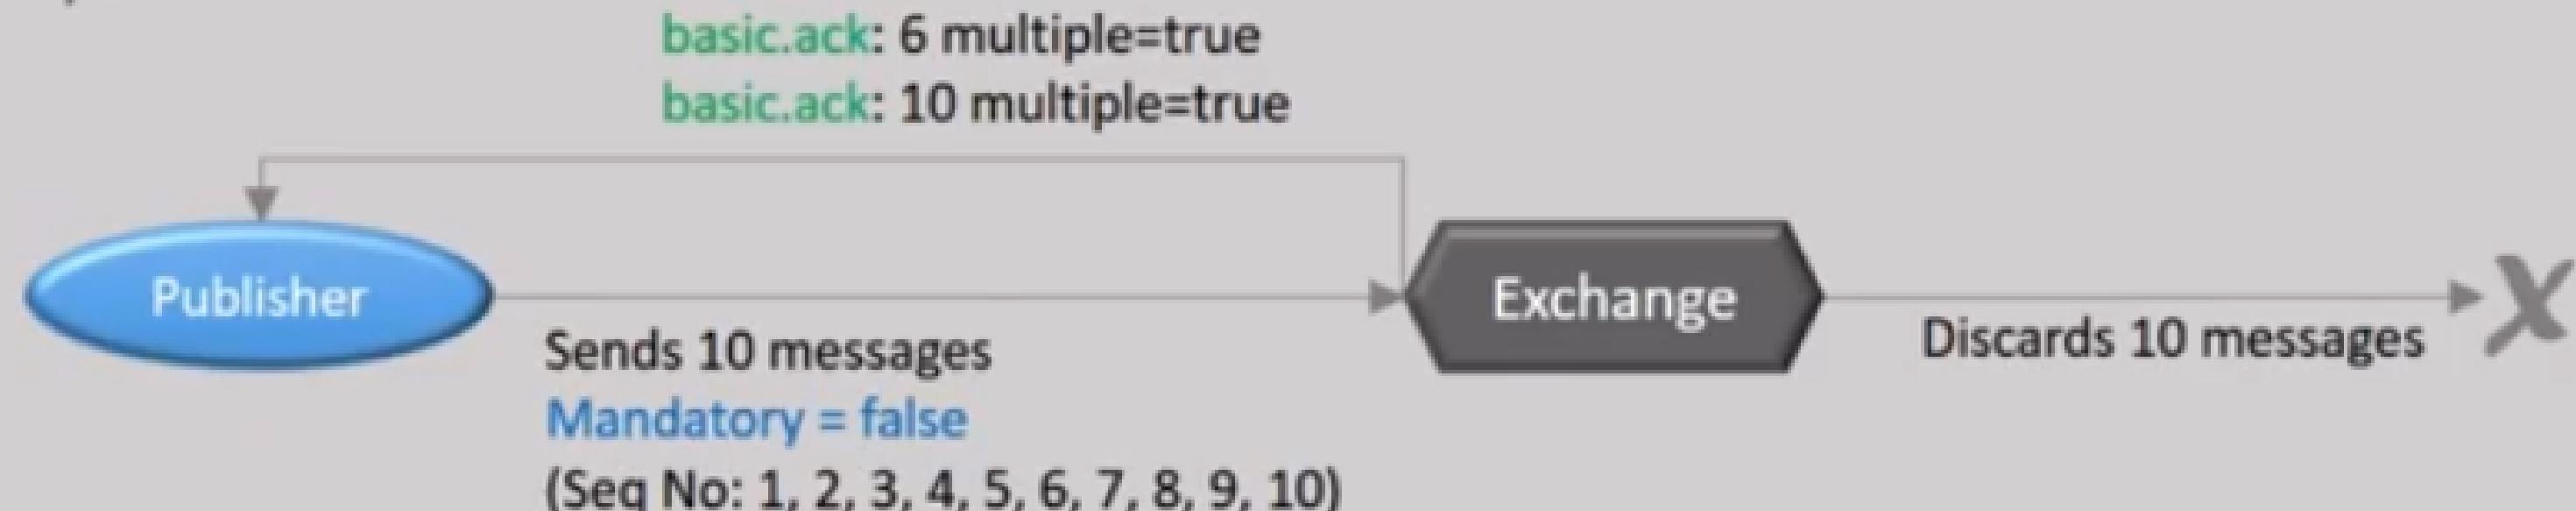
Flags

- **Multiple** (I am acknowledging multiple message deliveries)
- **Mandatory** (give me a `basic.return` if you can't deliver to any queues)

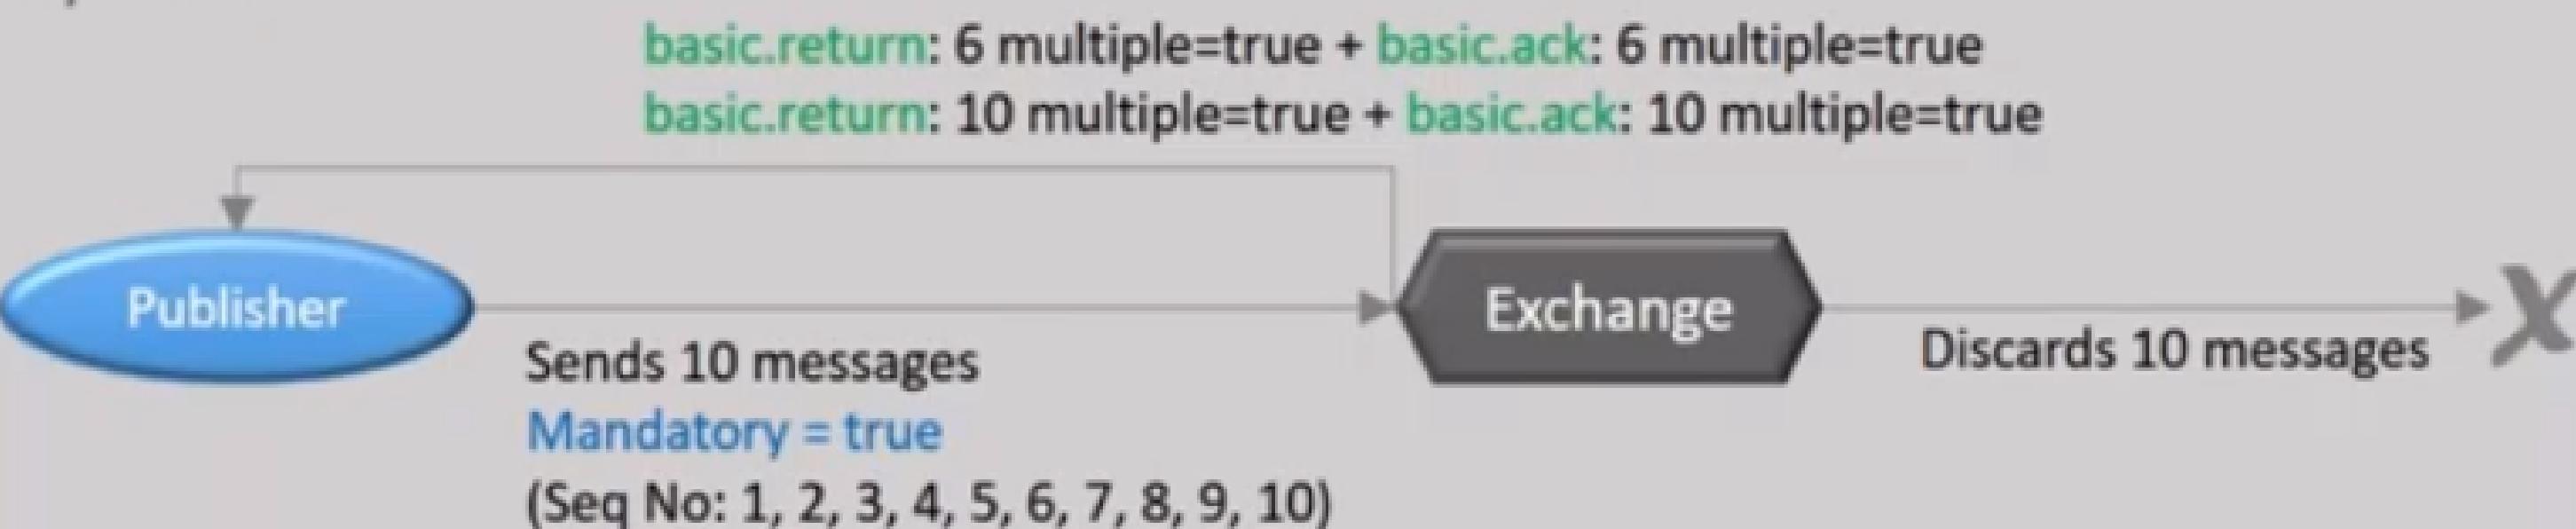


RabbitMQ - Producer Side Acknowledgements

Mandatory=false

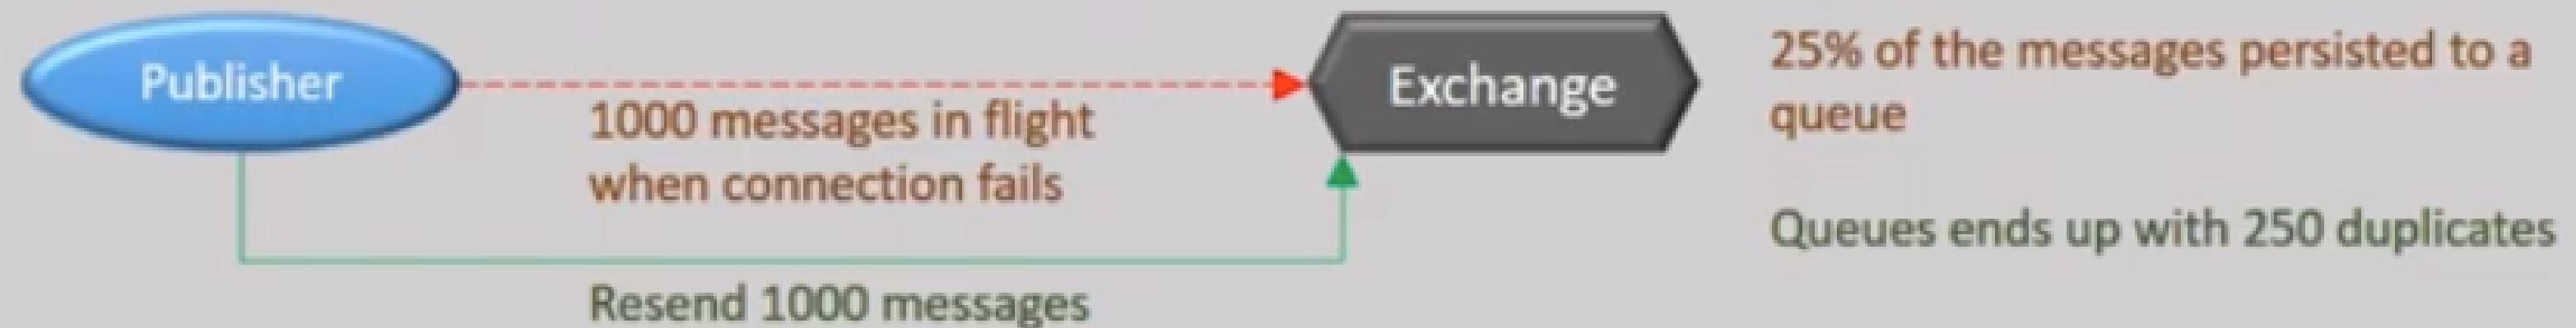


Mandatory=true

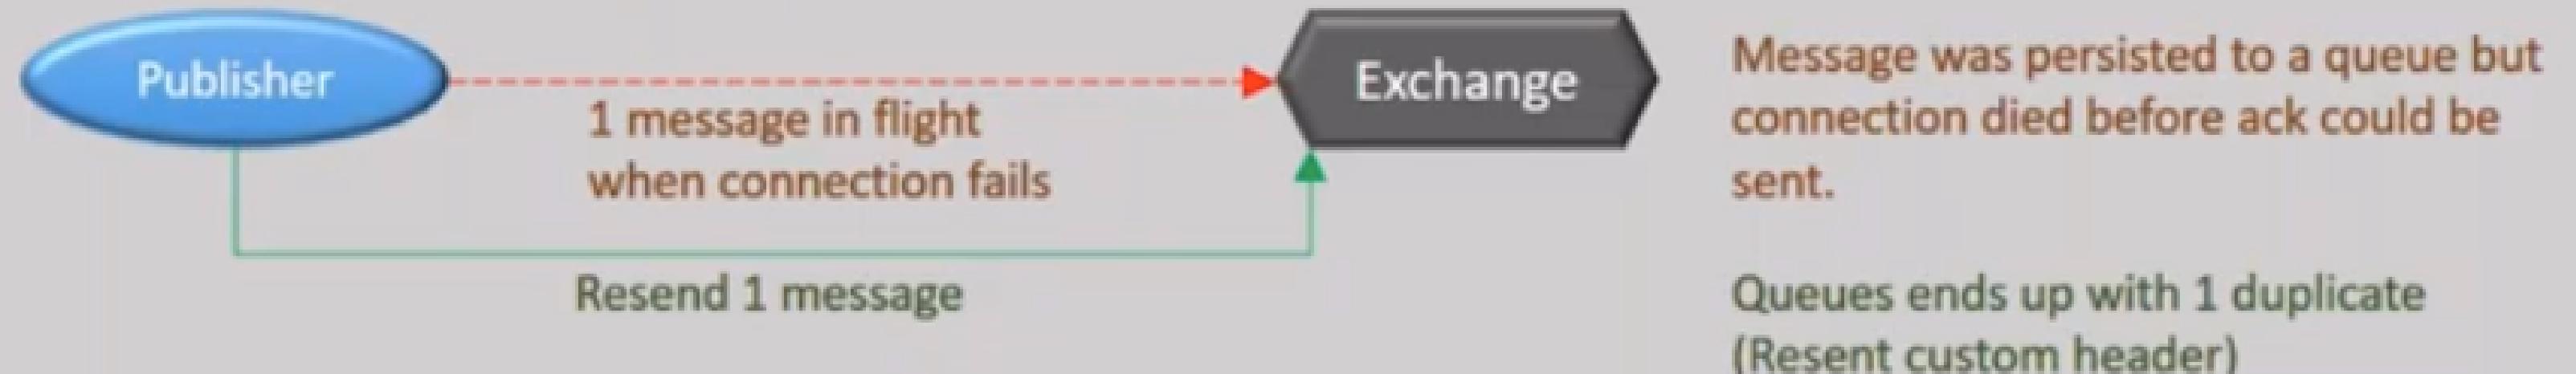


RabbitMQ - Producer Side Duplication

Large # of Messages in Flight = High Throughput, High Message Duplication on Failure



Low # of Messages in Flight = Low Throughput, Low Message Duplication on Failure



RabbitMQ - Consumer Side Acknowledgements

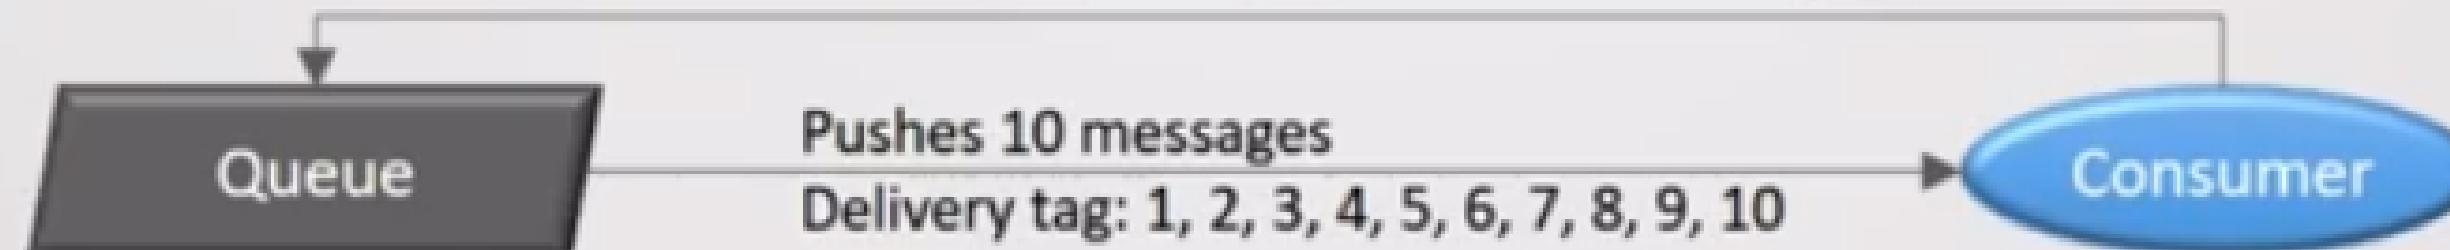
Acknowledgement Mode

- **Auto Ack** (Push me messages as fast as you can!)
- **Manual Ack** (I will explicitly tell you when a message can be removed from the queue)

Consumer Acknowledgements

- **basic.ack** (all ok, remove from the queue!)
- **basic.nack, redeliver=false** (error, but remove anyway)
- **basic.nack, redeliver=true** (error, please redeliver)
- **basic.reject** (same as basic.nack but without multiple flag support)

basic.ack: 1 multiple=false, basic.ack: 2 multiple=false
basic.ack: 3 multiple=false, basic.ack: 4 multiple=false
basic.ack: 5 multiple=false, basic.ack: 6 multiple=false
basic.ack: 7 multiple=false, basic.ack: 8 multiple=false
basic.ack: 9 multiple=false, basic.ack: 9 multiple=false

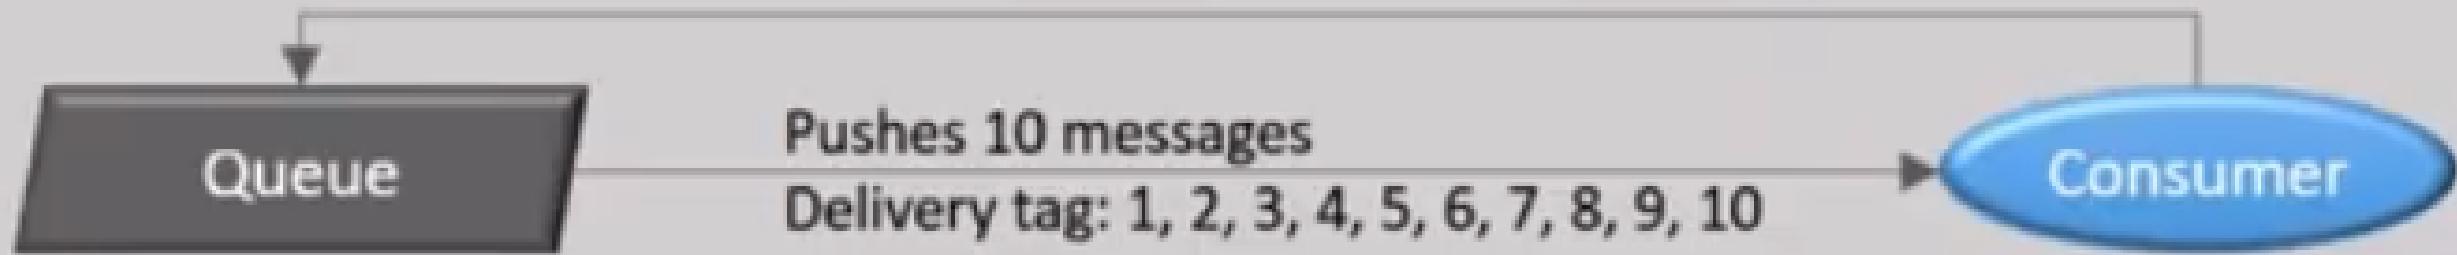


RabbitMQ - Consumer Side Acknowledgements

Flags

- **Multiple** (I am acknowledging multiple messages)
- **Redelivered** (This message is a redelivery)

basic.ack: 6 multiple=true, basic.ack: 10 multiple=true



Redelivered Flag

2. basic.nack: 1 multiple=false redeliver=true

1. Pushes 1 message

3. Delivers message again
with redelivered=true flag

4. basic.ack: 1 multiple=false

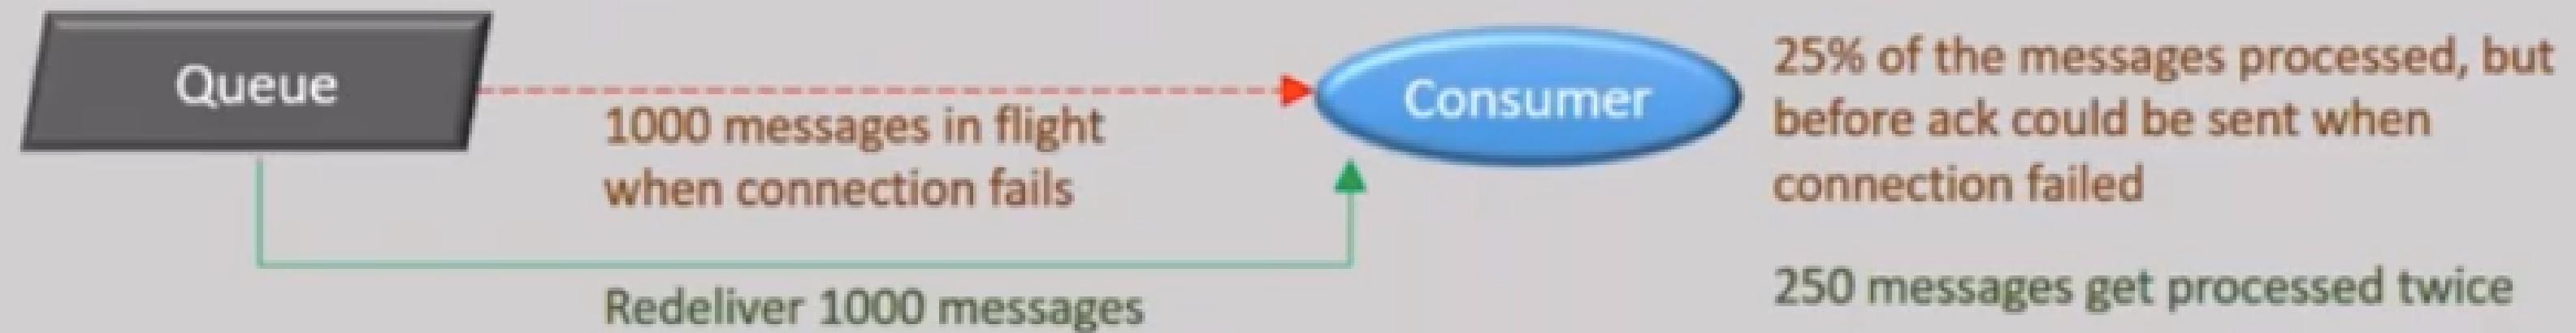
Consumer

Consumer

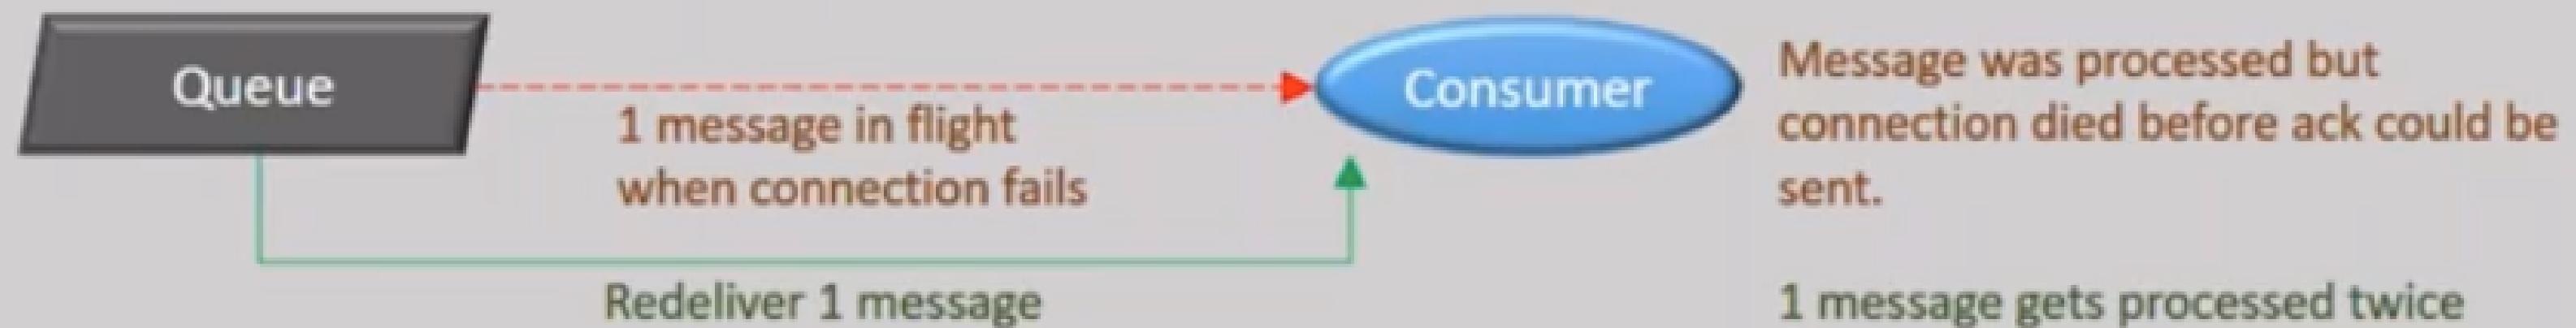
Consumer

RabbitMQ - Consumer Side Duplication

Large # of Messages in Flight = High Throughput, High Message Duplication on Failure



Low # of Messages in Flight = Low Throughput, Low Message Duplication on Failure



RabbitMQ

Broker Durability

Durable Queues
Persistent Messages
Mirrored Queues

RabbitMQ - The Broker

Surviving Broker Restart

- Durable Queue
- Persistent Message

Surviving Broker Loss

- Queue Mirroring
(Clustering)

Non-Durable Queue
Non-Persistent Message

Durable Queue
Non-Persistent Message

Durable Queue
Persistent Message

Mirrored Queue
Persistent Message

Broker Restart

Queue Message

Queue Message

Queue Message

Queue Message

Broker Loss

Queue Message

Queue Message

Queue Message

Queue Message

RabbitMQ - The Broker

Surviving Broker Restart

- Durable Queue
- Persistent Message

Surviving Broker Loss

- Queue Mirroring
(Clustering)

Non-Durable Queue
Non-Persistent Message

Durable Queue
Non-Persistent Message

Durable Queue
Persistent Message

Mirrored Queue
Persistent Message

Broker Restart

Queue

Message

Queue

Message

Queue

Message

Queue

Message

Broker Loss

Queue

Message

Queue

Message

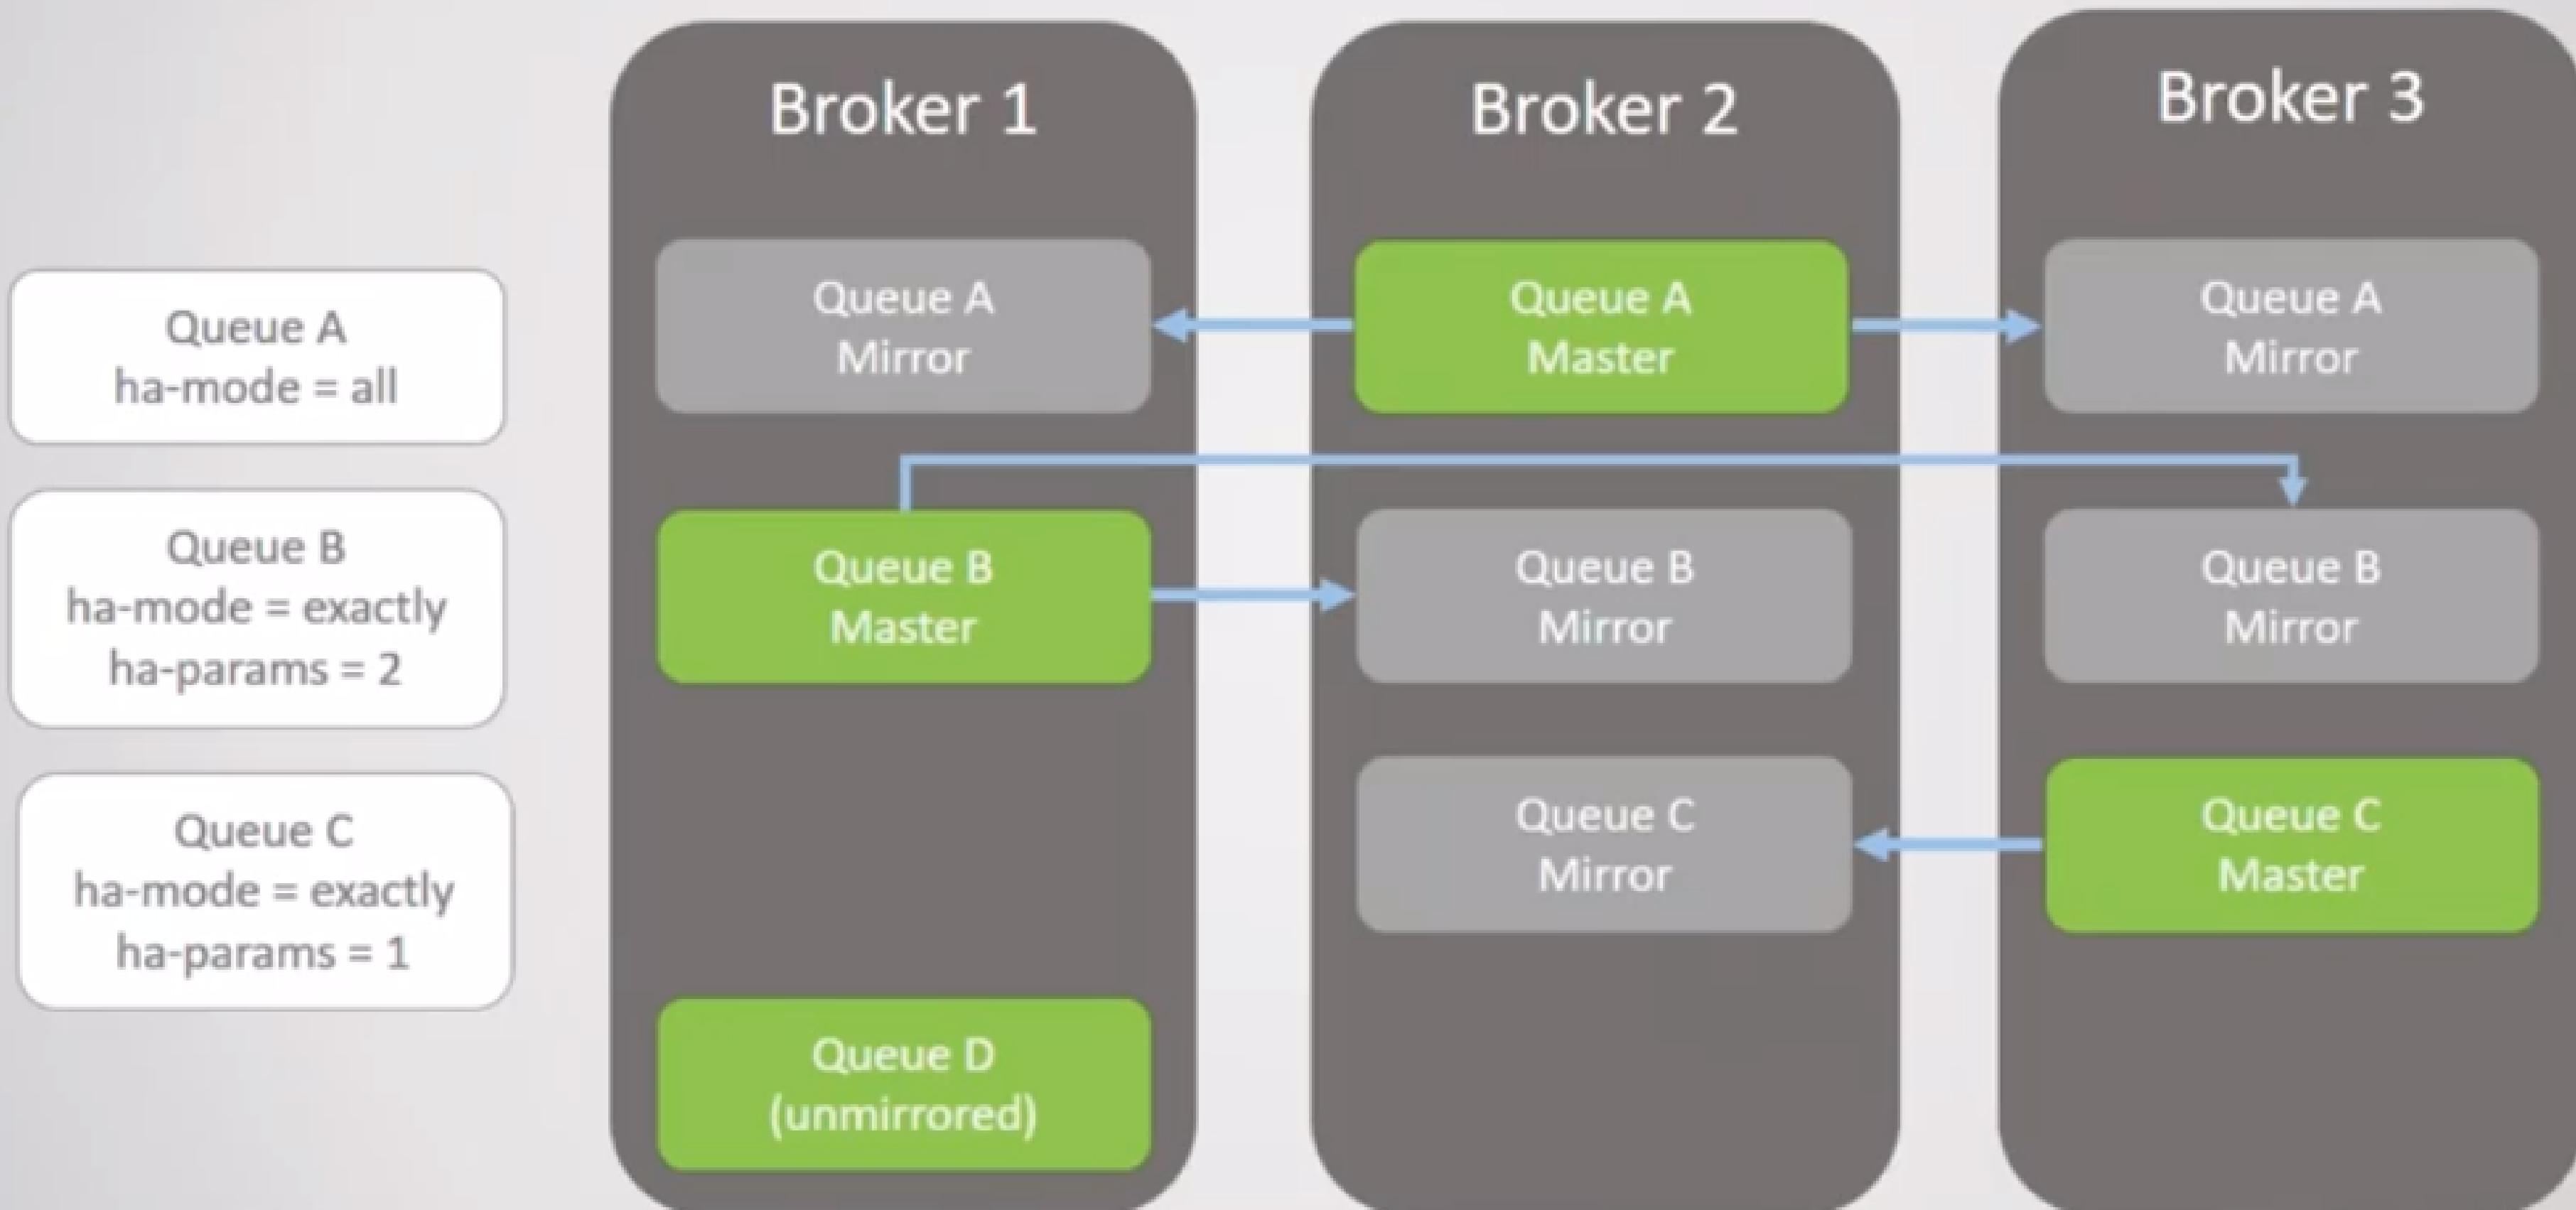
Queue

Message

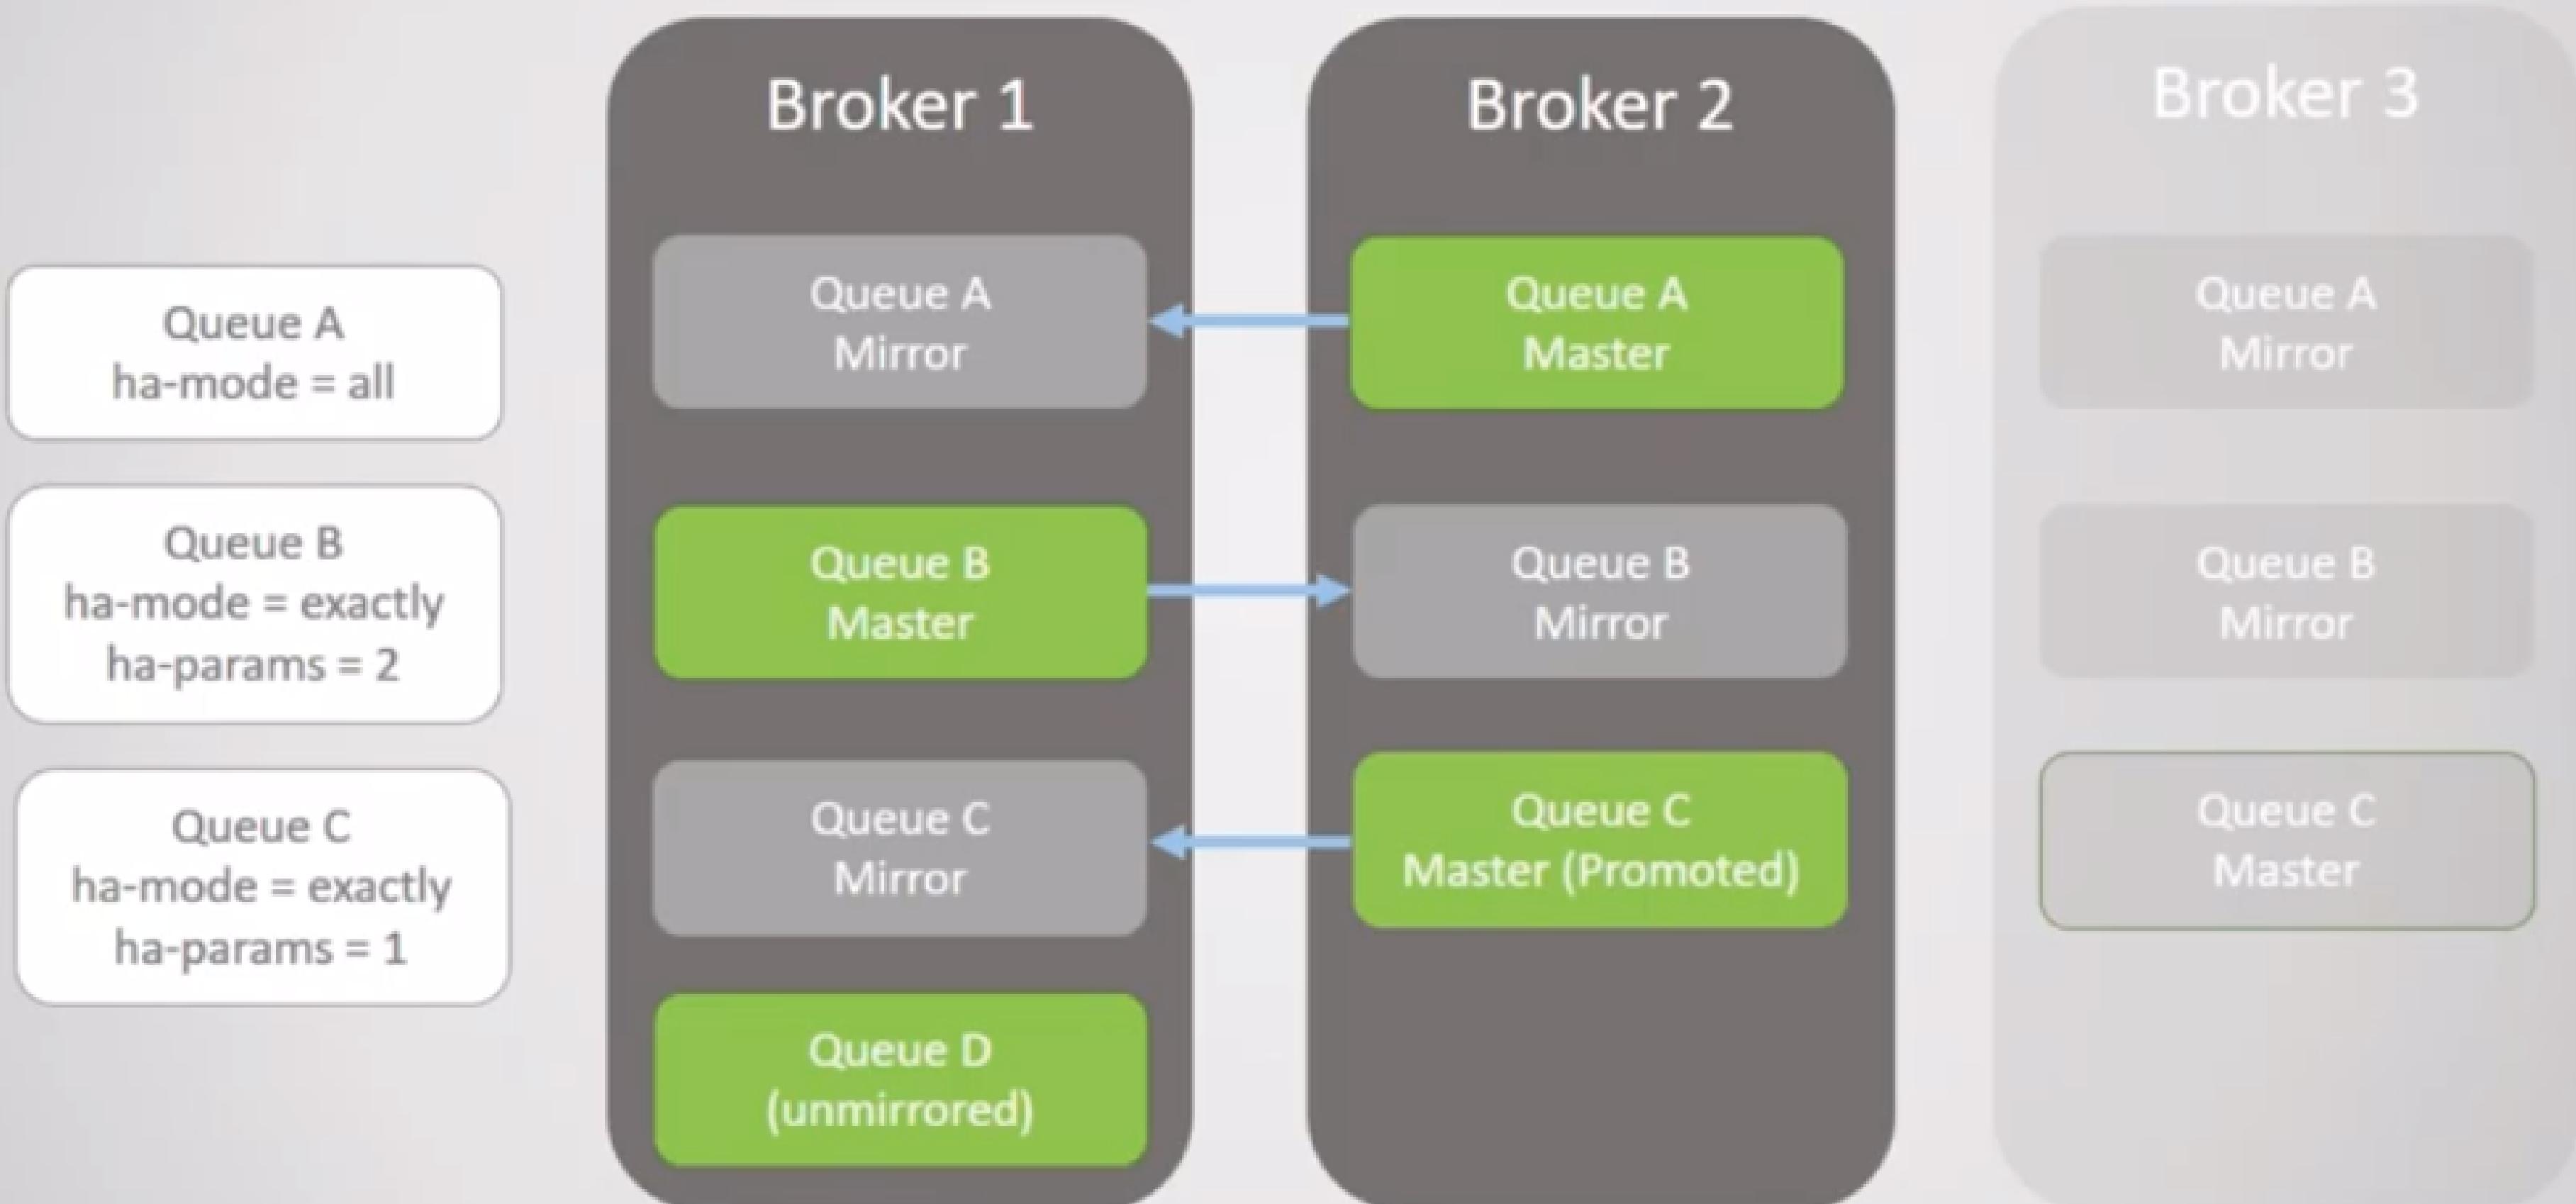
Queue

Message

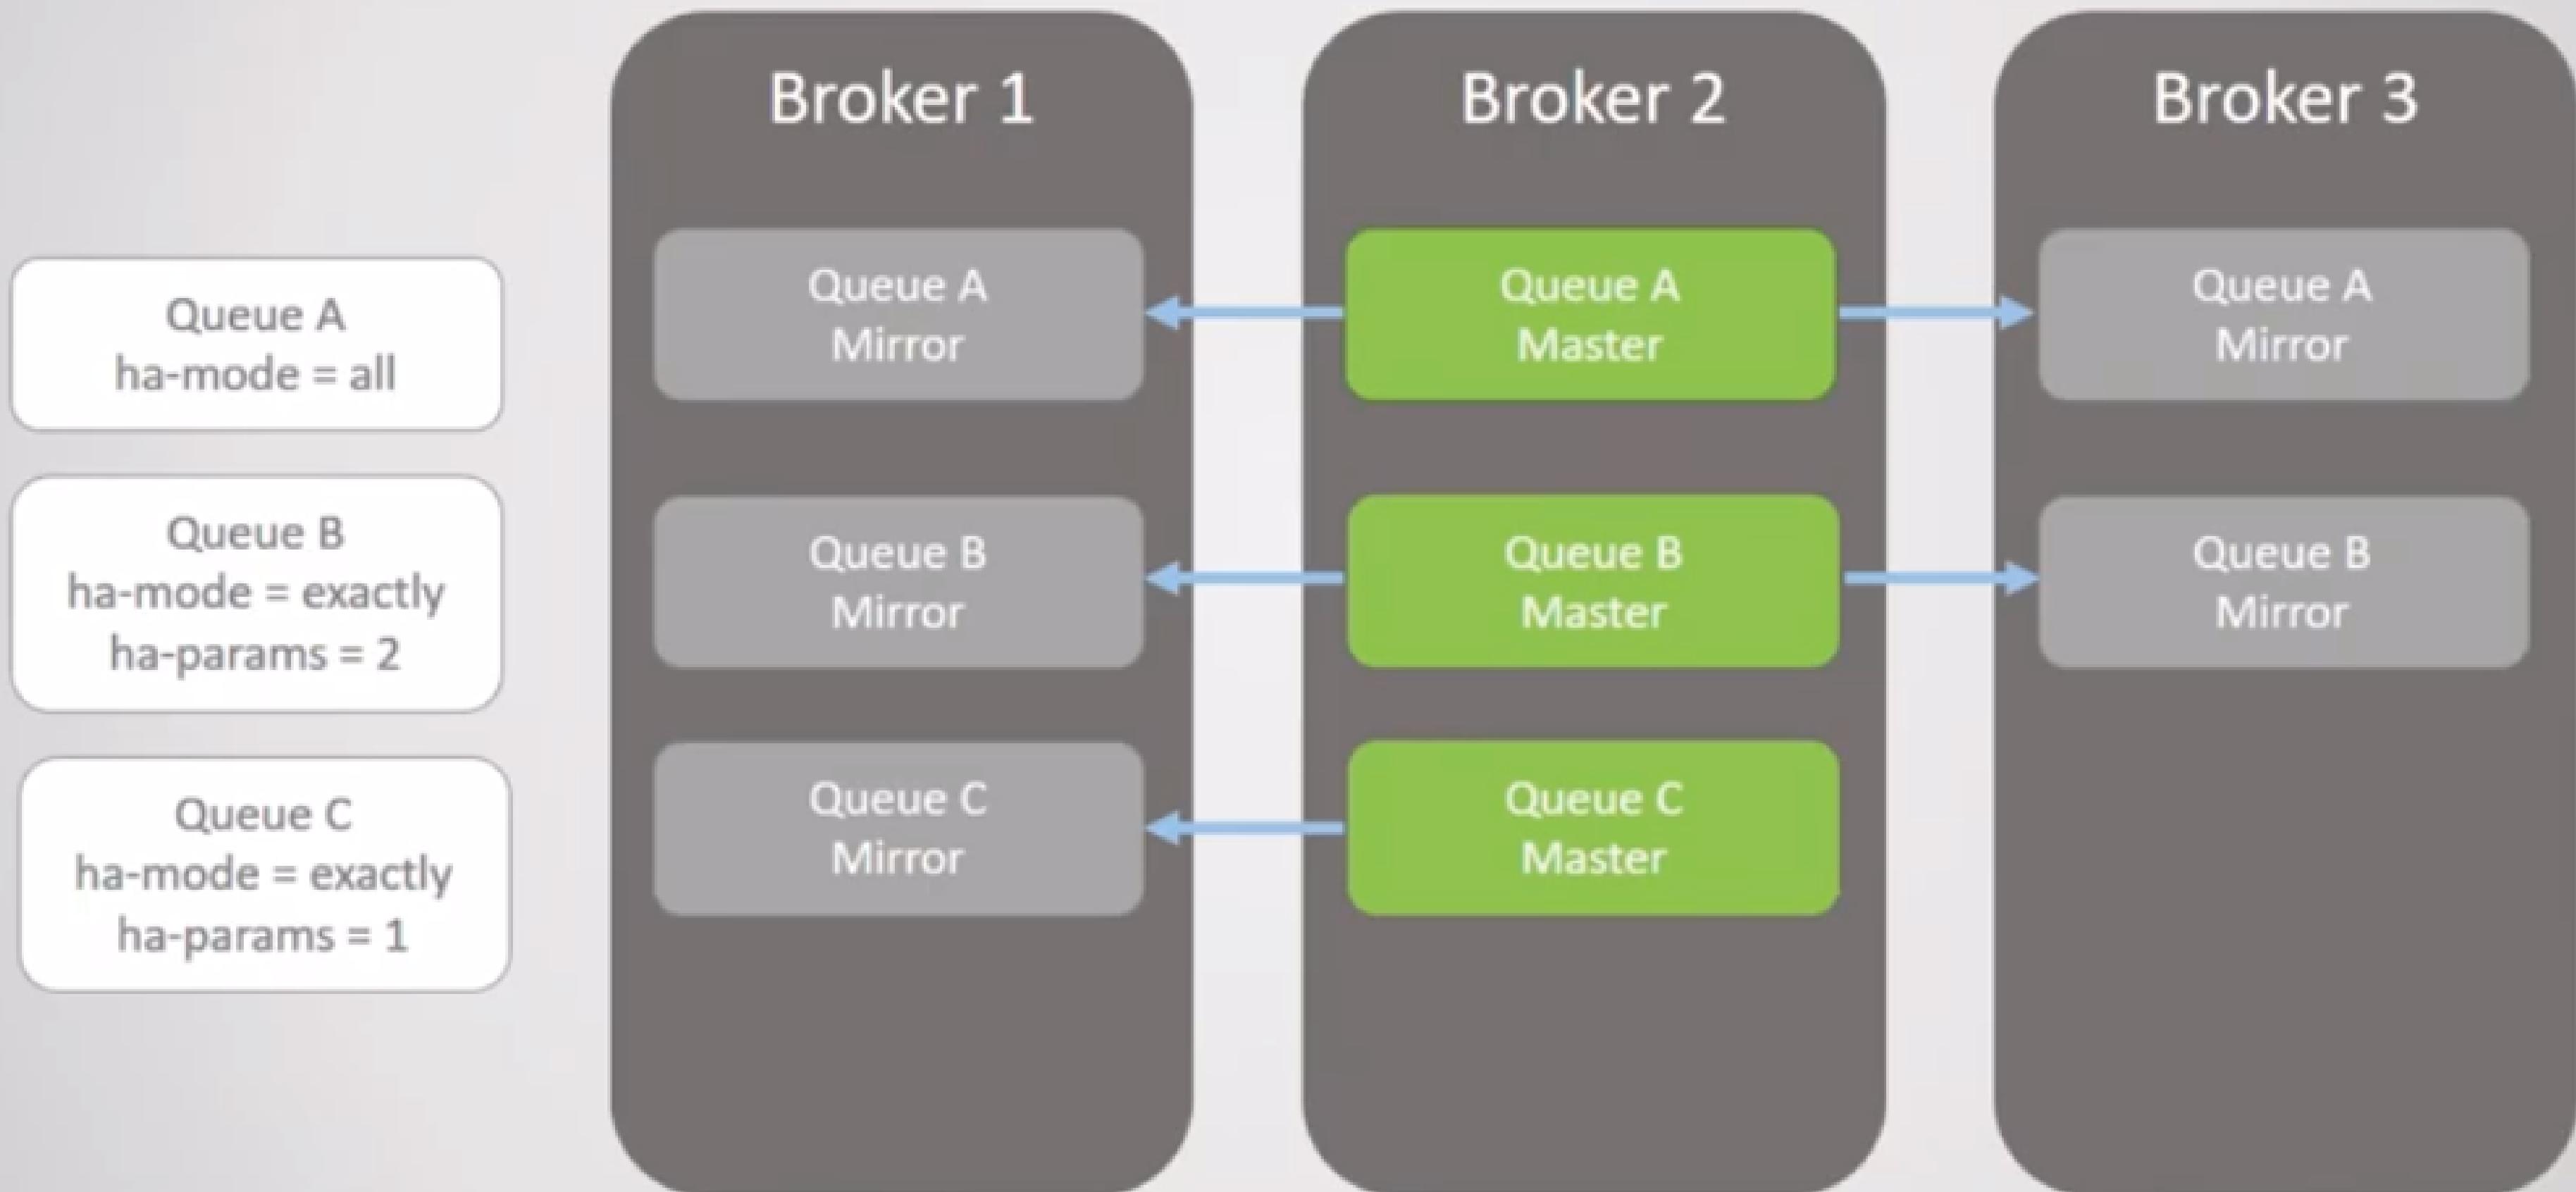
RabbitMQ - The Broker - Queue Mirrors



RabbitMQ - The Broker - Queue Mirrors



RabbitMQ - The Broker - Queue Mirrors

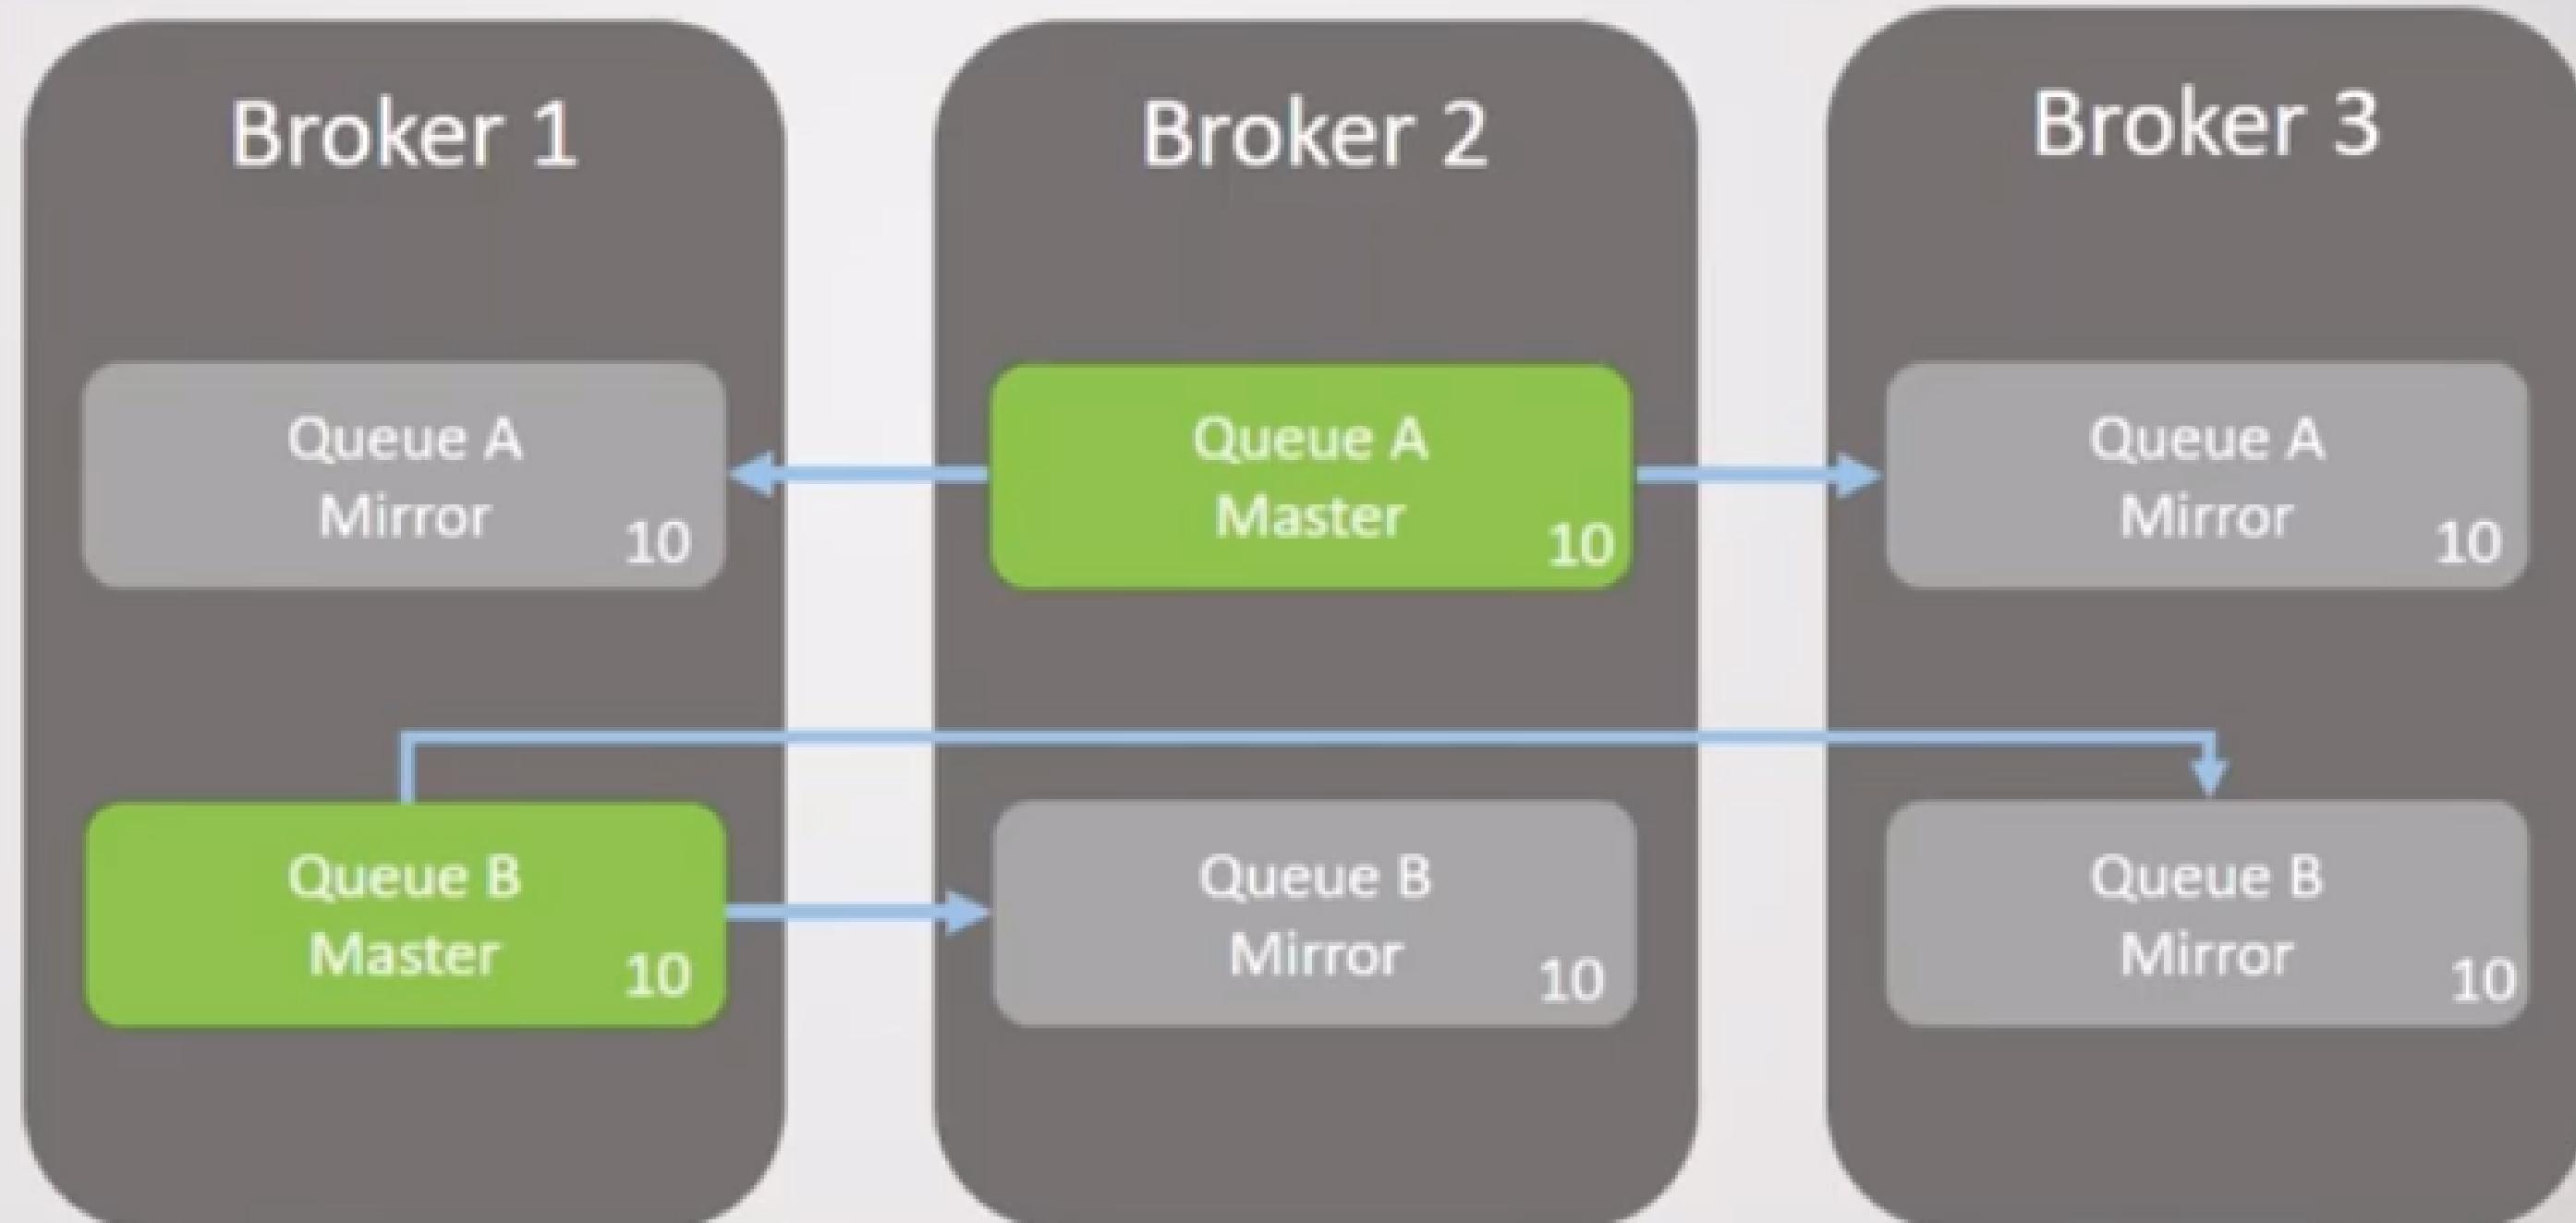


RabbitMQ

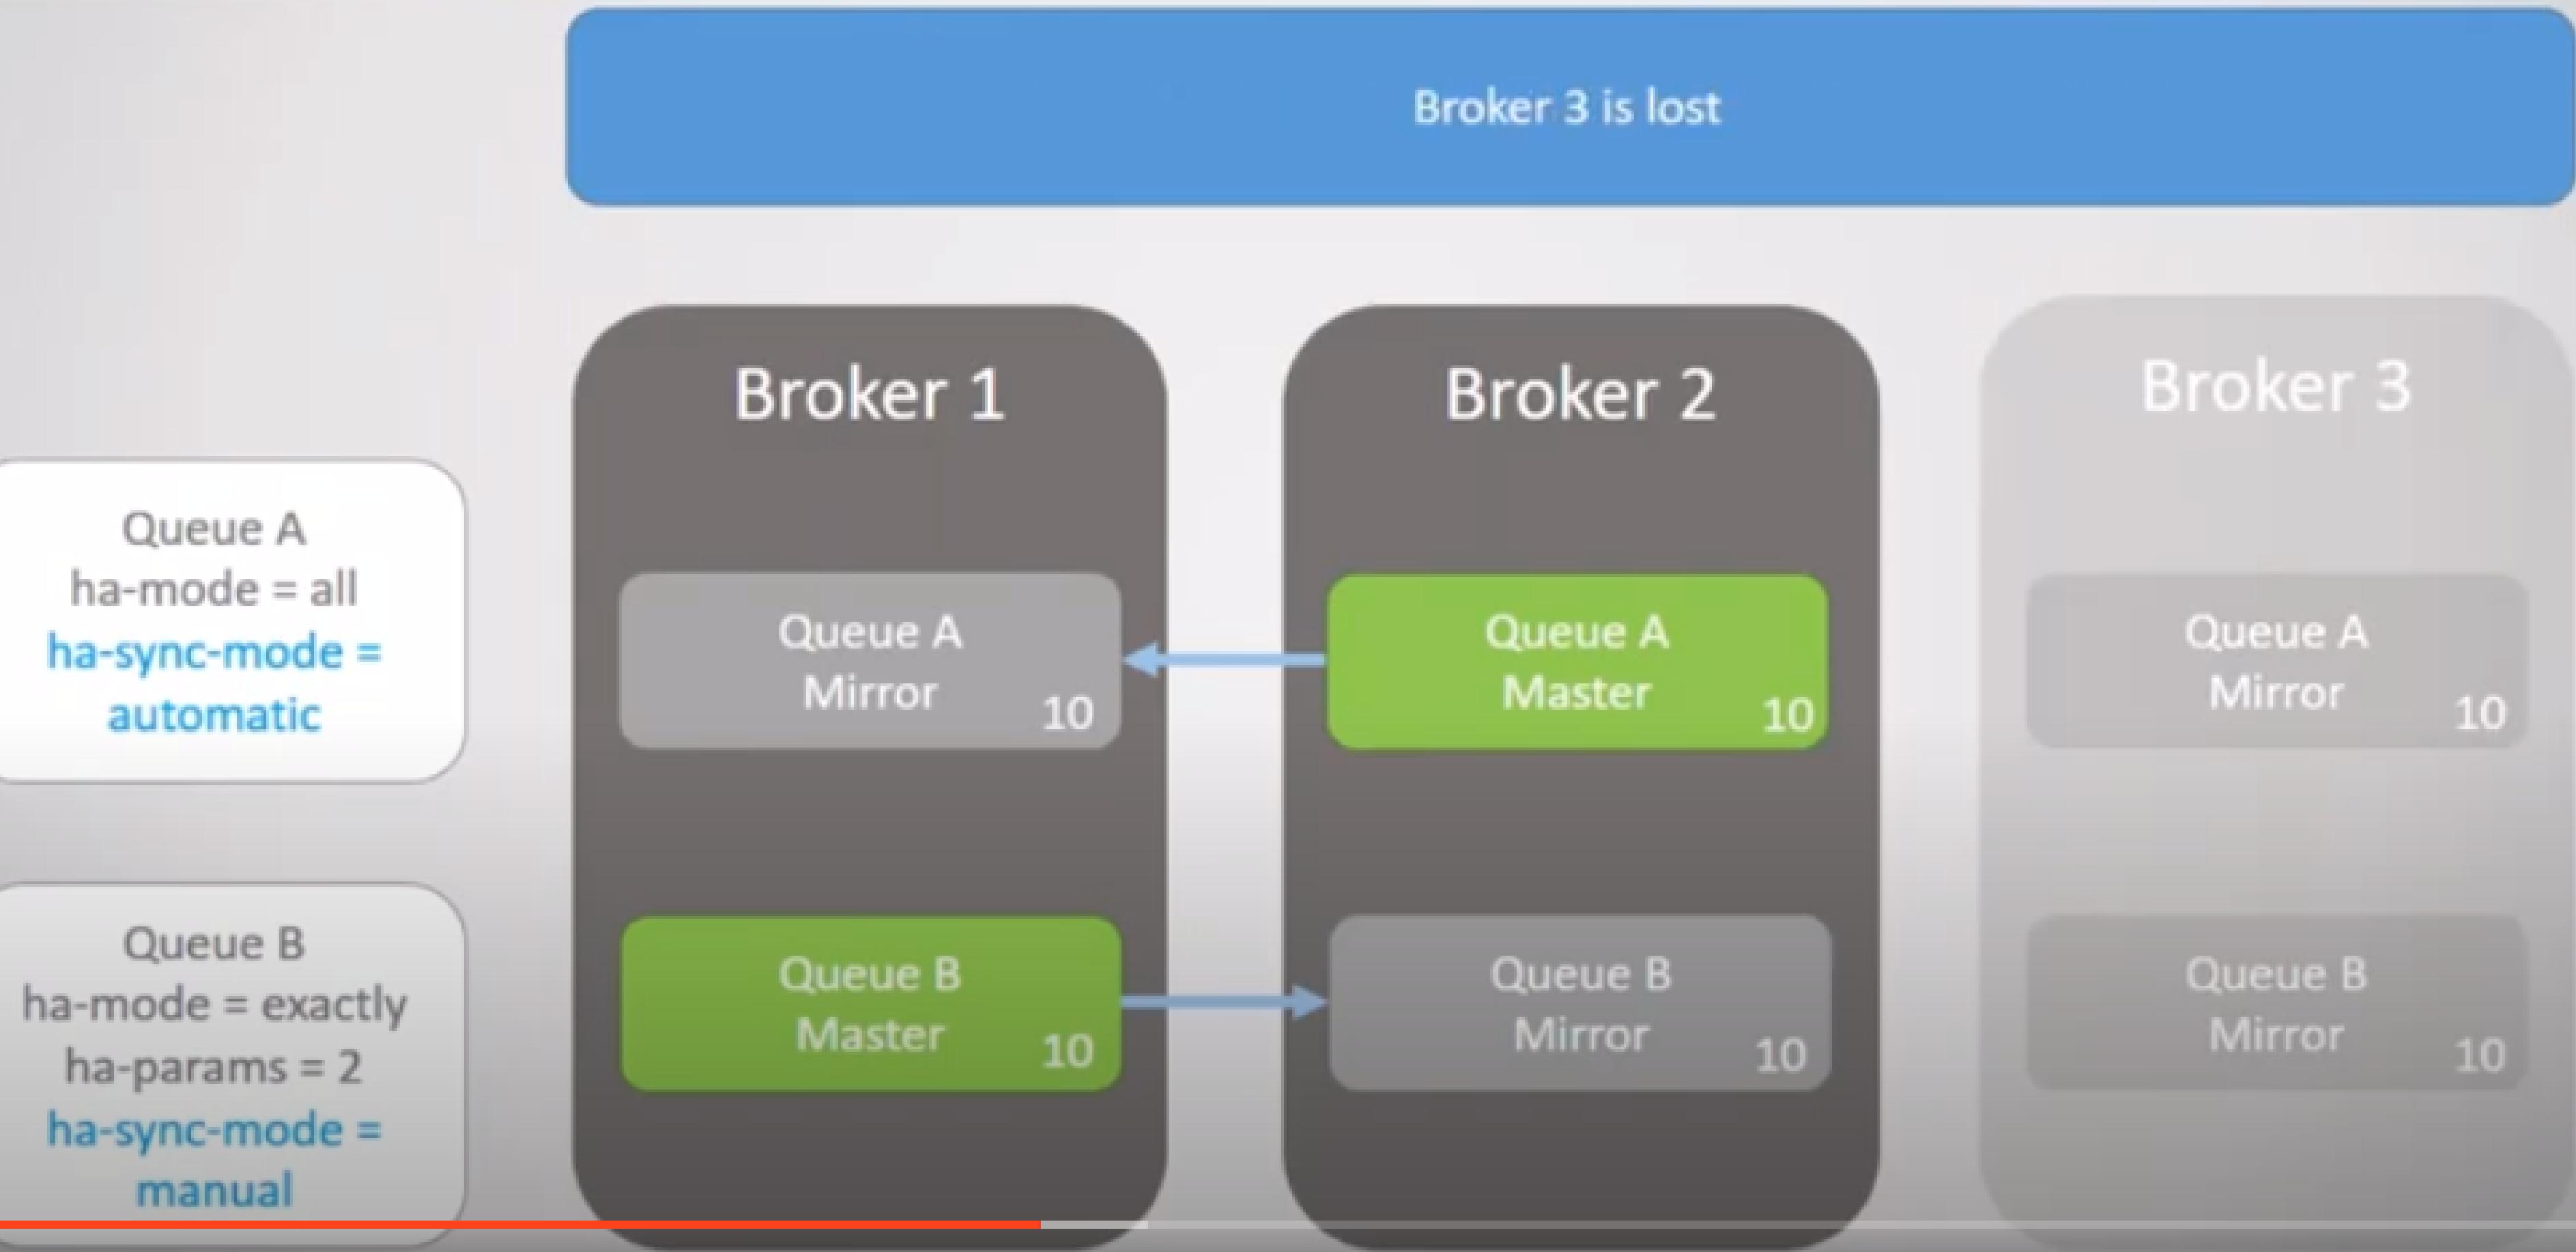
Queue Mirror Synchronization And Queue Failover

RabbitMQ - Queue Mirrors - Synchronization

Three nodes, two mirrored queues each with 10 messages

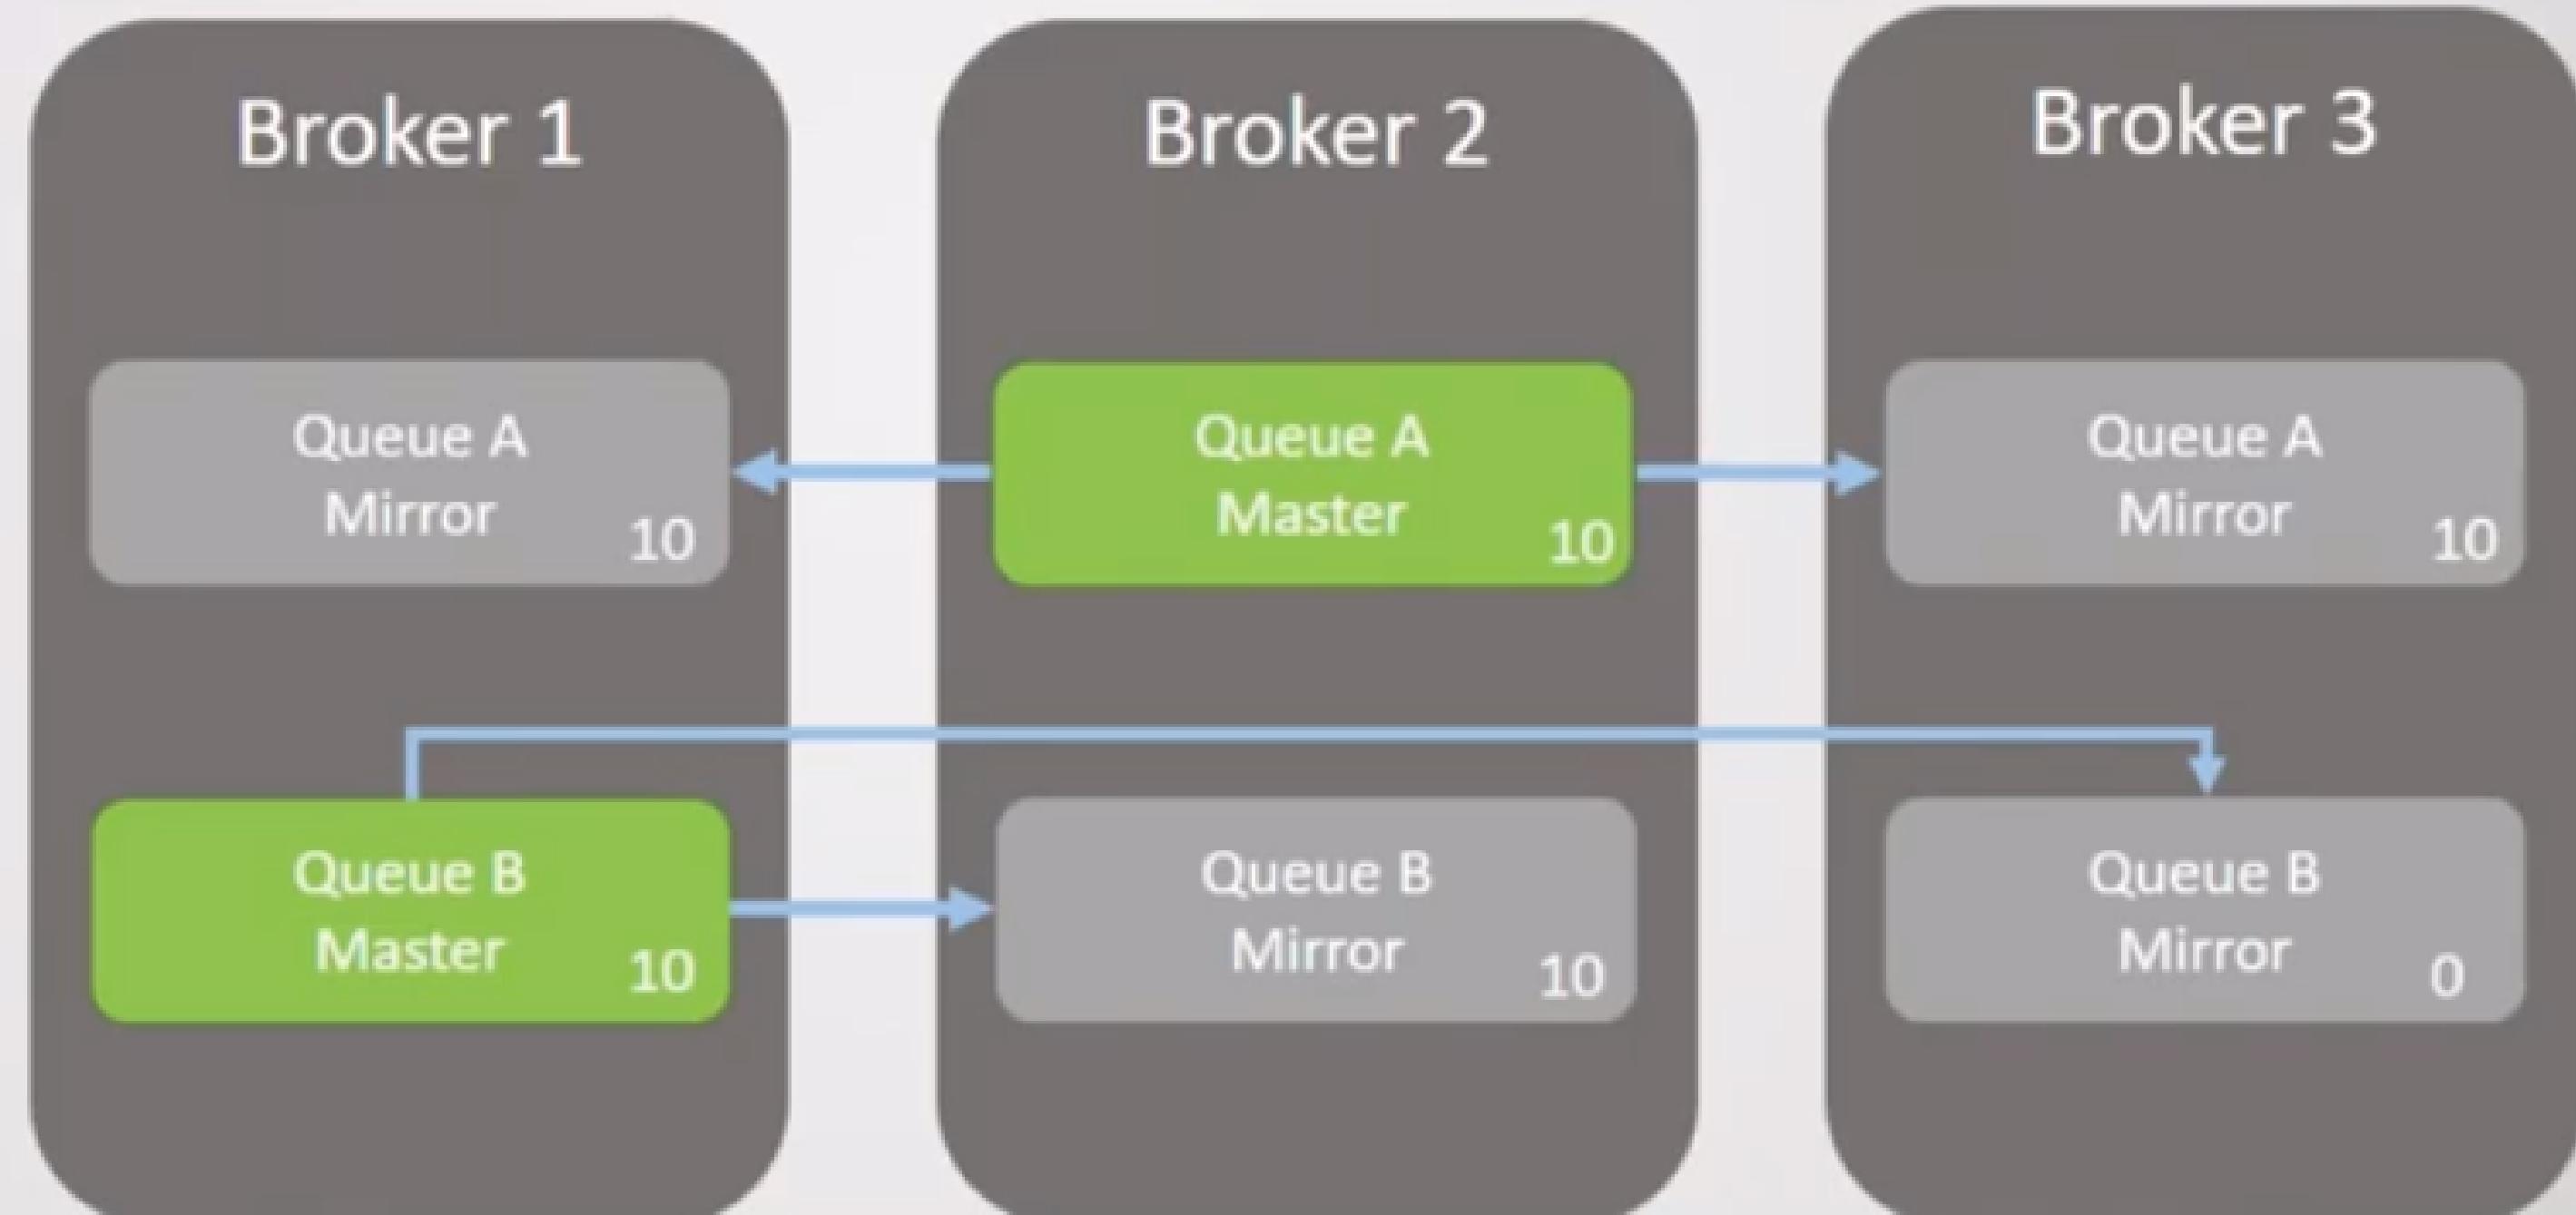


RabbitMQ - Queue Mirrors - Synchronization



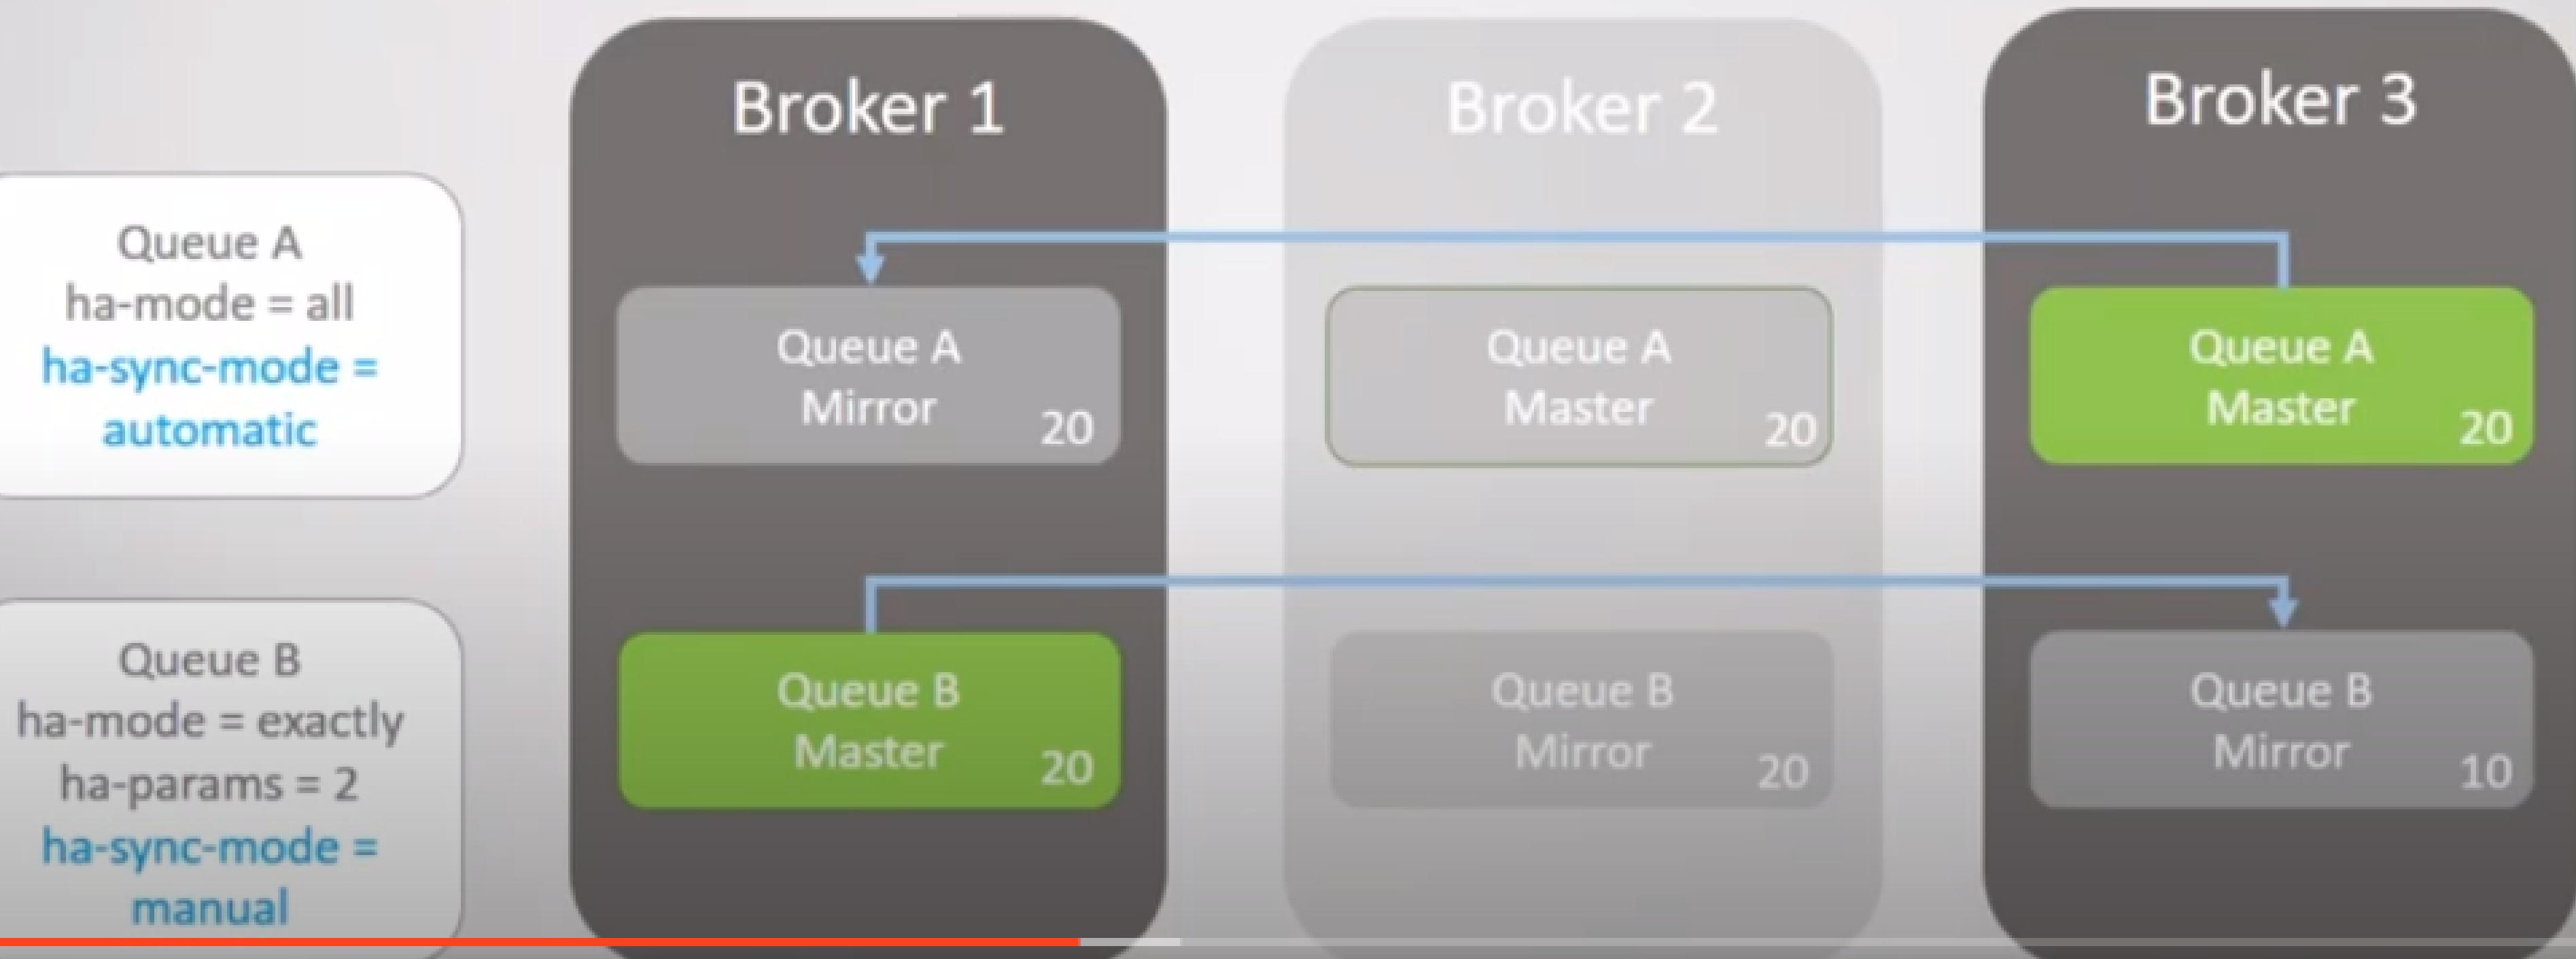
RabbitMQ - Queue Mirrors - Synchronization

Broker 3 comes back.
Mirror A is automatically synchronized. Mirror B remains at 0 messages.



RabbitMQ - Queue Mirrors - Synchronization

Each queue receives 10 more messages.
Broker 2 is lost. Queue A fails over to mirror 3 without data loss.

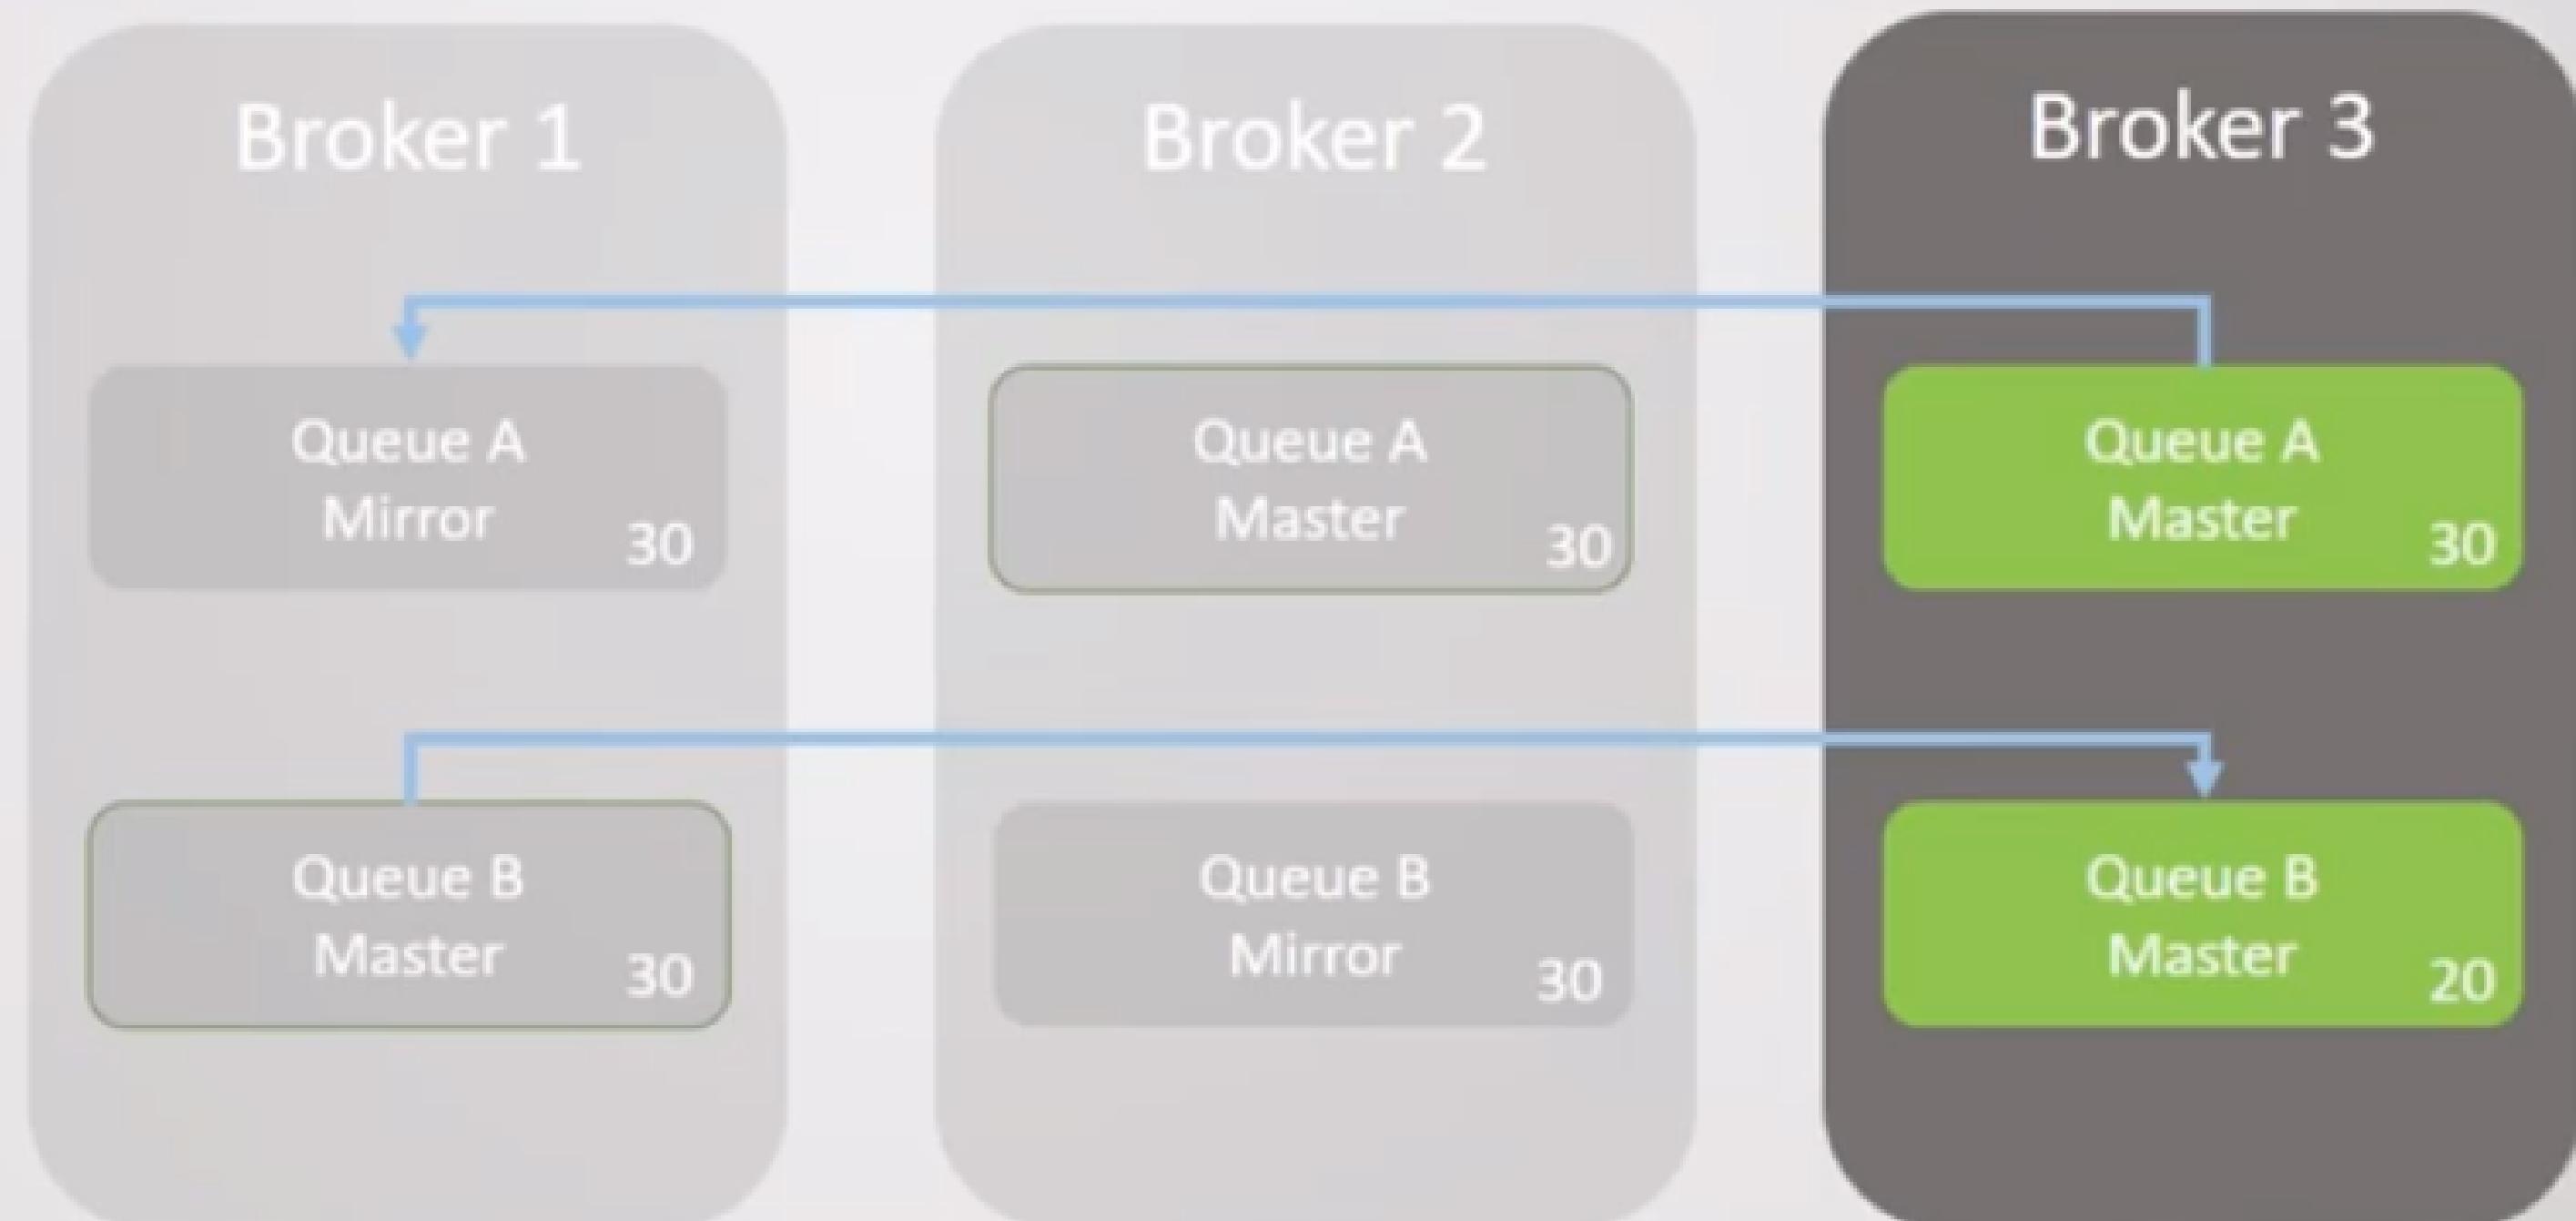


RabbitMQ - Queue Mirrors - Synchronization

Each queue receives 10 more messages.
Broker 1 is lost. Queue B fails over to mirror 3 and loses 10 messages.

Queue A
ha-mode = all
ha-sync-mode = automatic

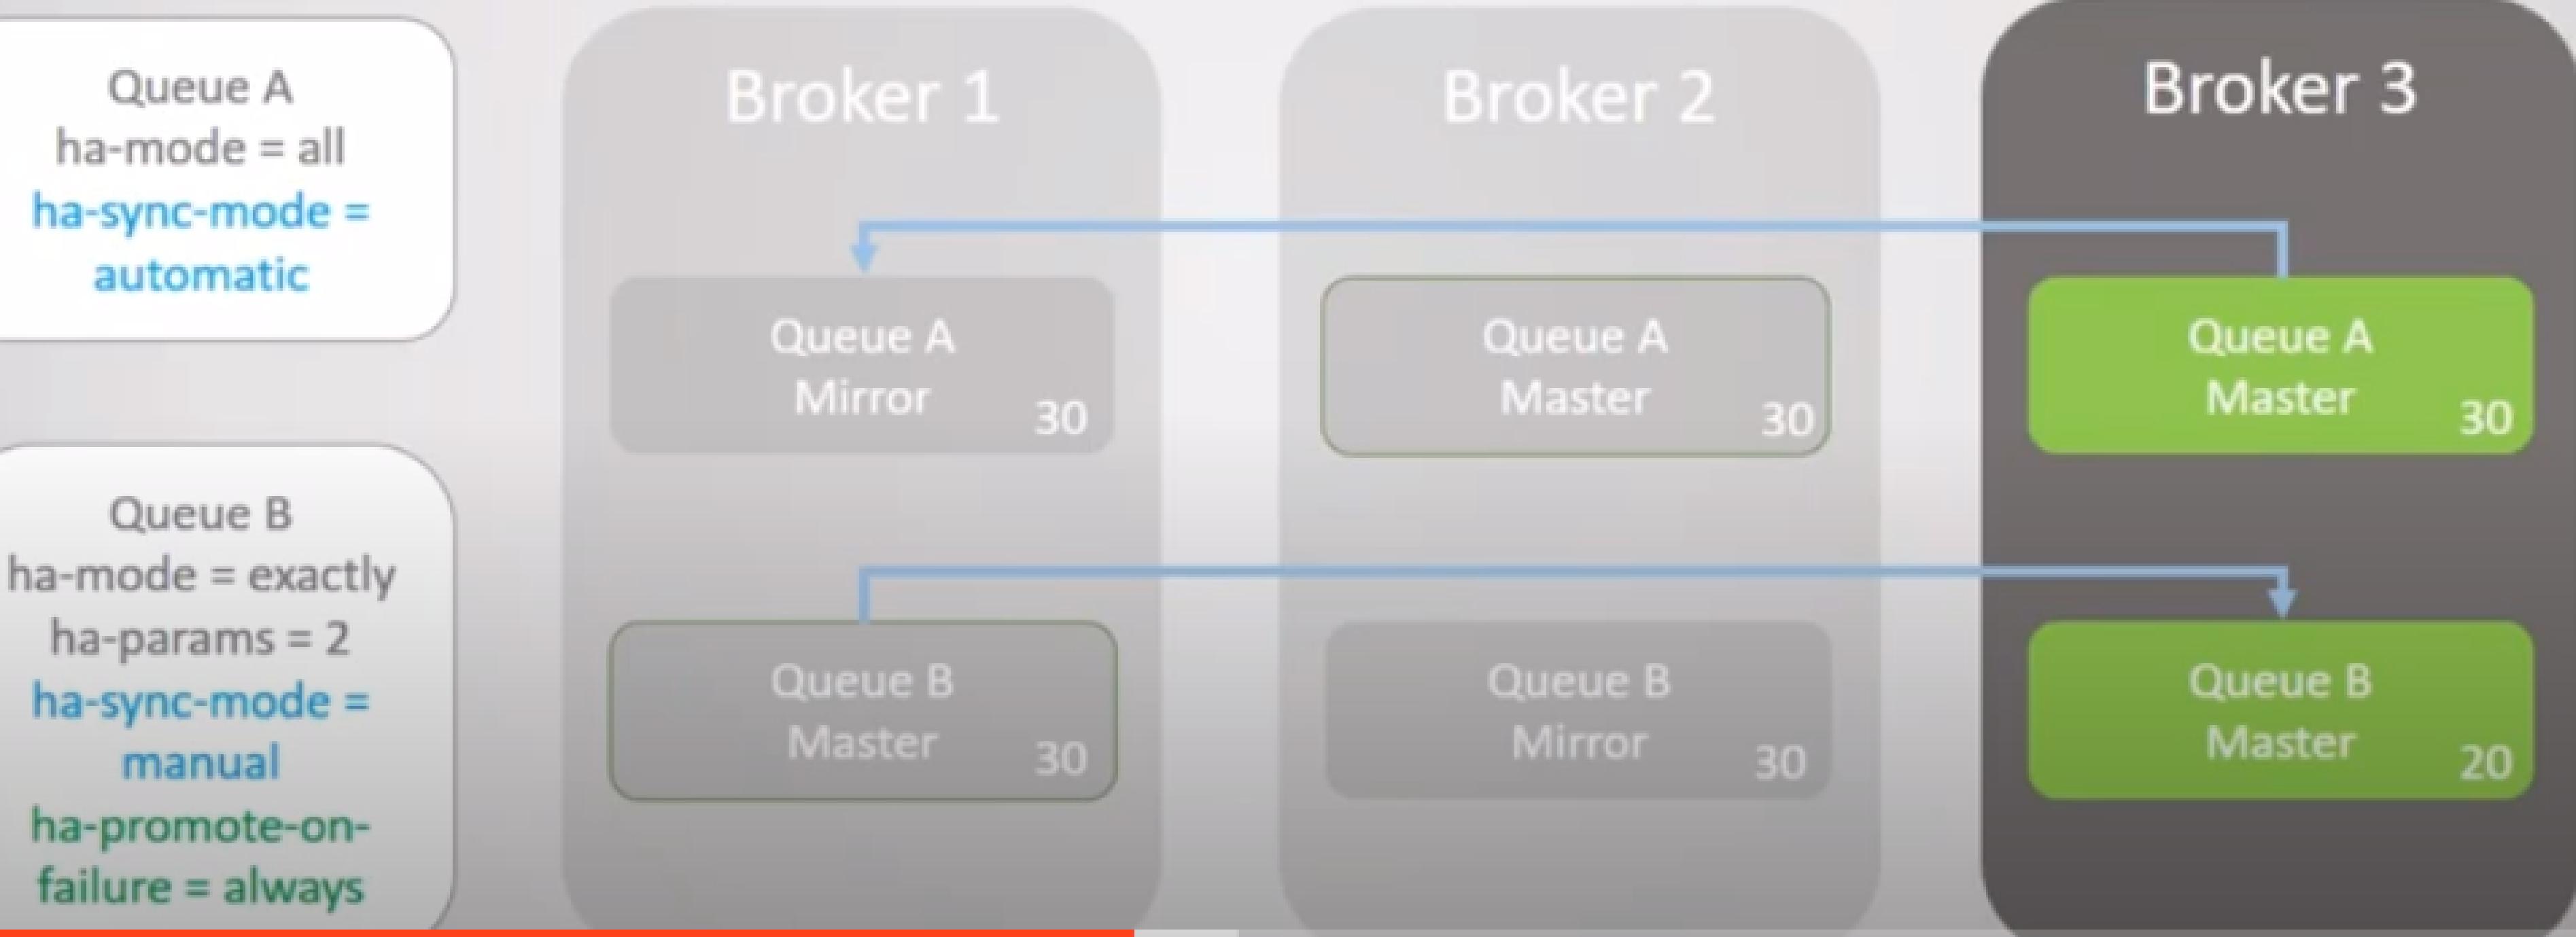
Queue B
ha-mode = exactly
ha-params = 2
ha-sync-mode = manual
ha-promote-on-failure = always



RabbitMQ - Queue Mirrors - Synchronization

Each queue receives 10 more messages.

Broker 1 is lost. Queue B fails over to mirror 3 and loses 10 messages.

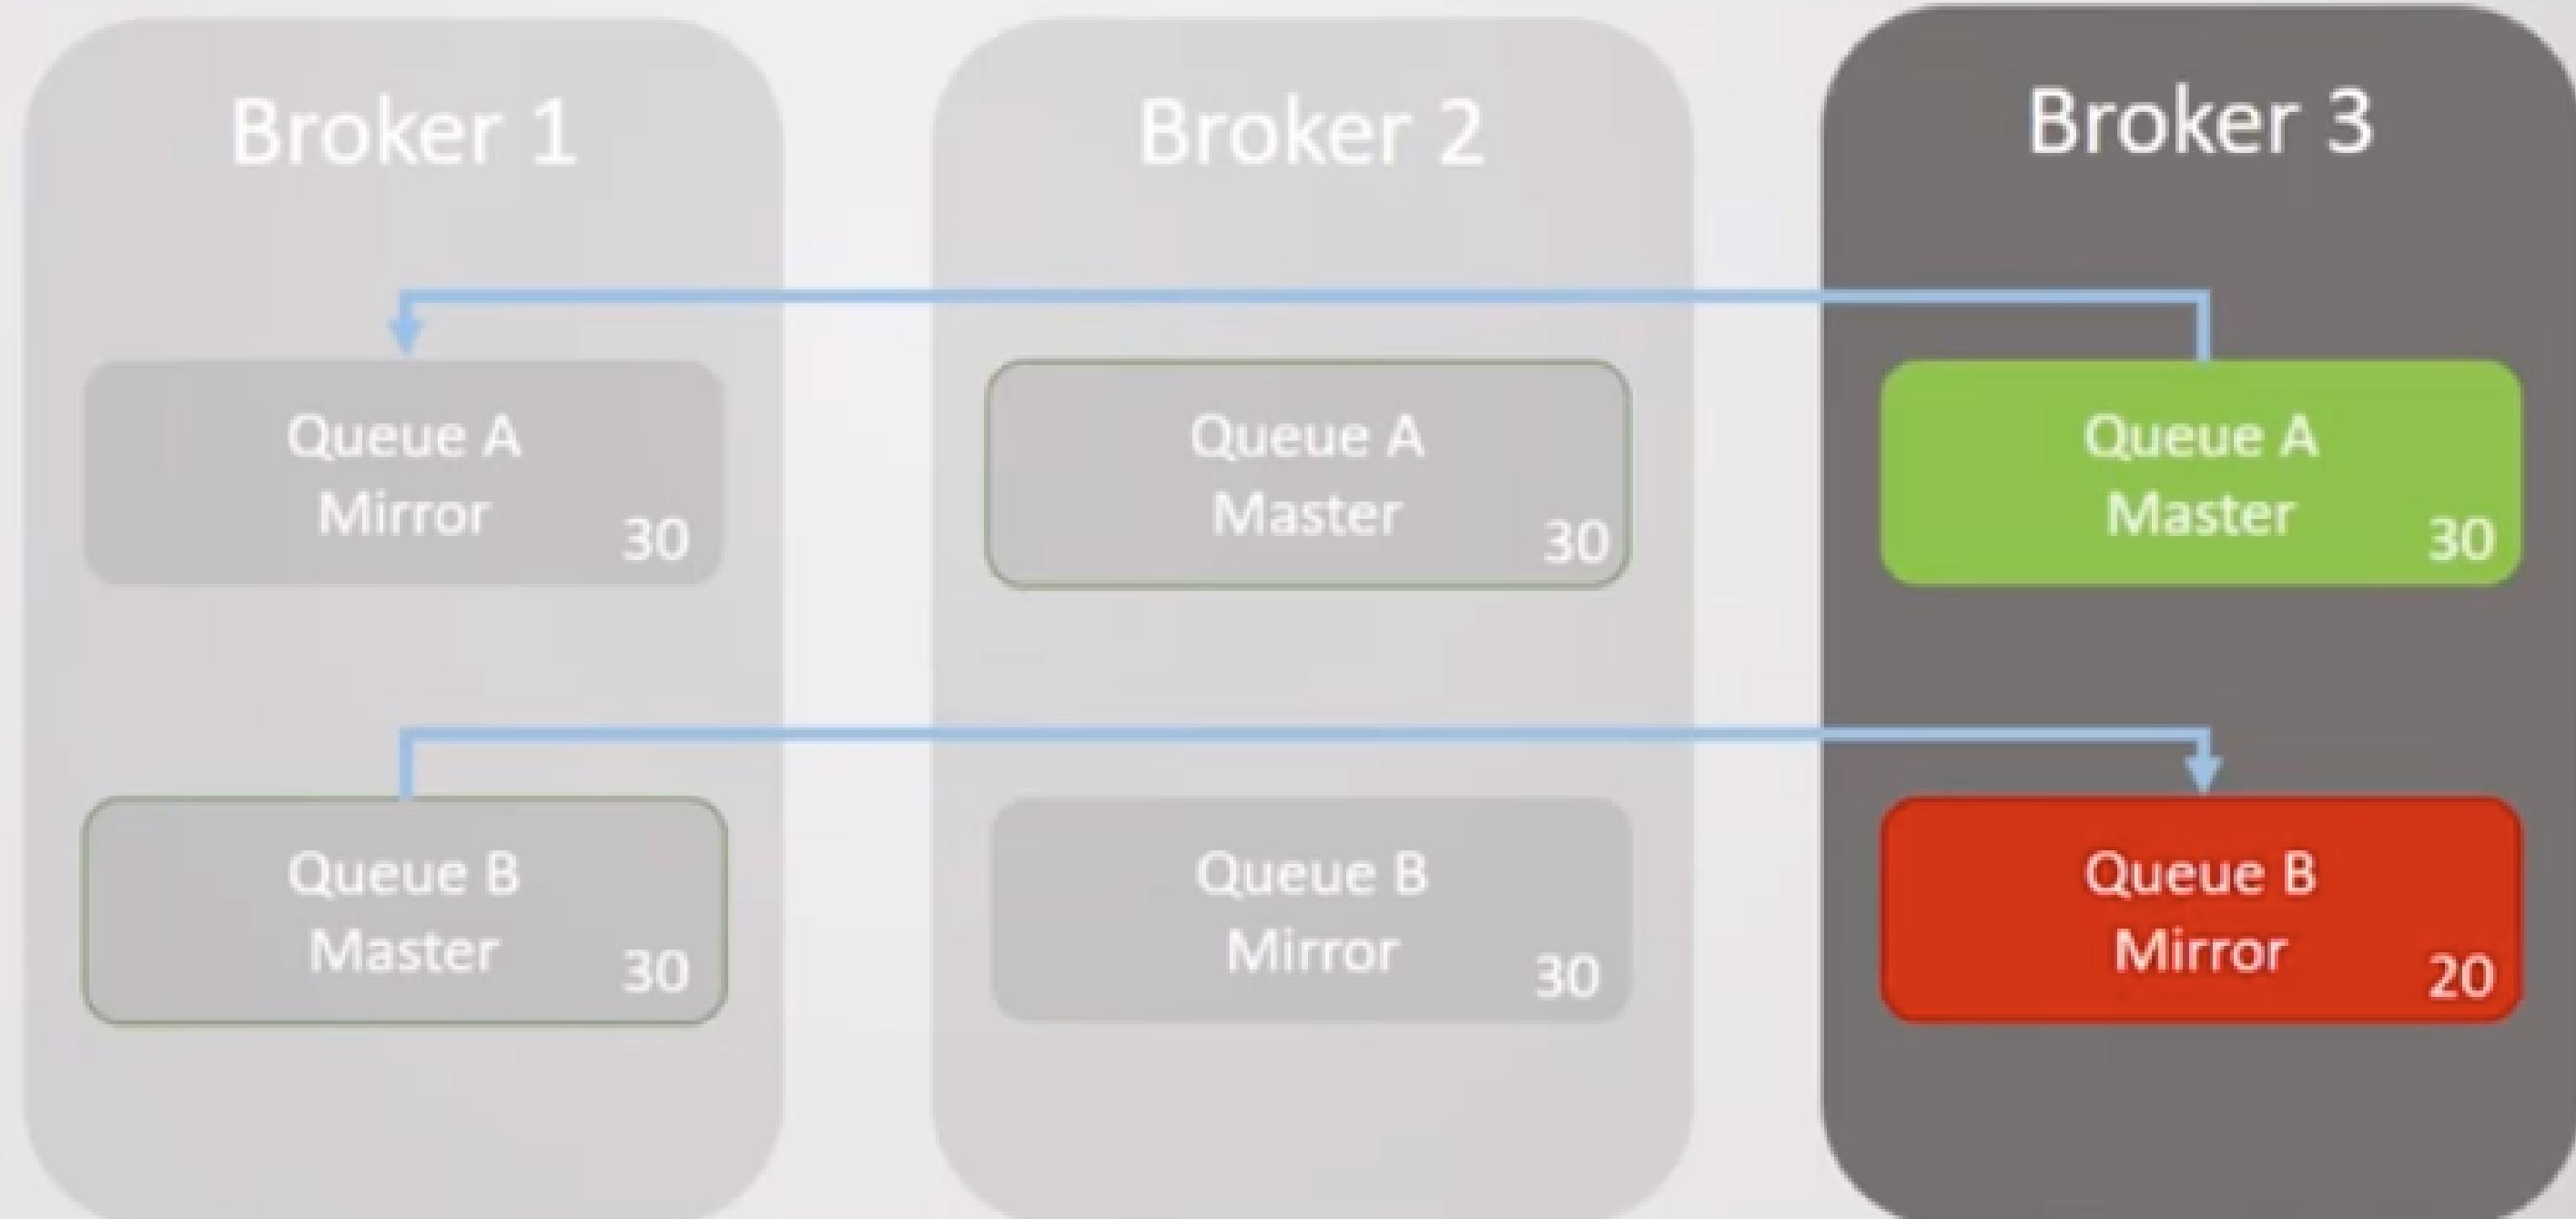


RabbitMQ - Queue Mirrors - Synchronization

Alternate scenario: ha-promote-on-failure = when-synced
Queue B does not fail over as mirror 3 is unsynchronized.

Queue A
ha-mode = all
ha-sync-mode = automatic

Queue B
ha-mode = exactly
ha-params = 2
ha-sync-mode = manual
ha-promote-on-failure = when-synced

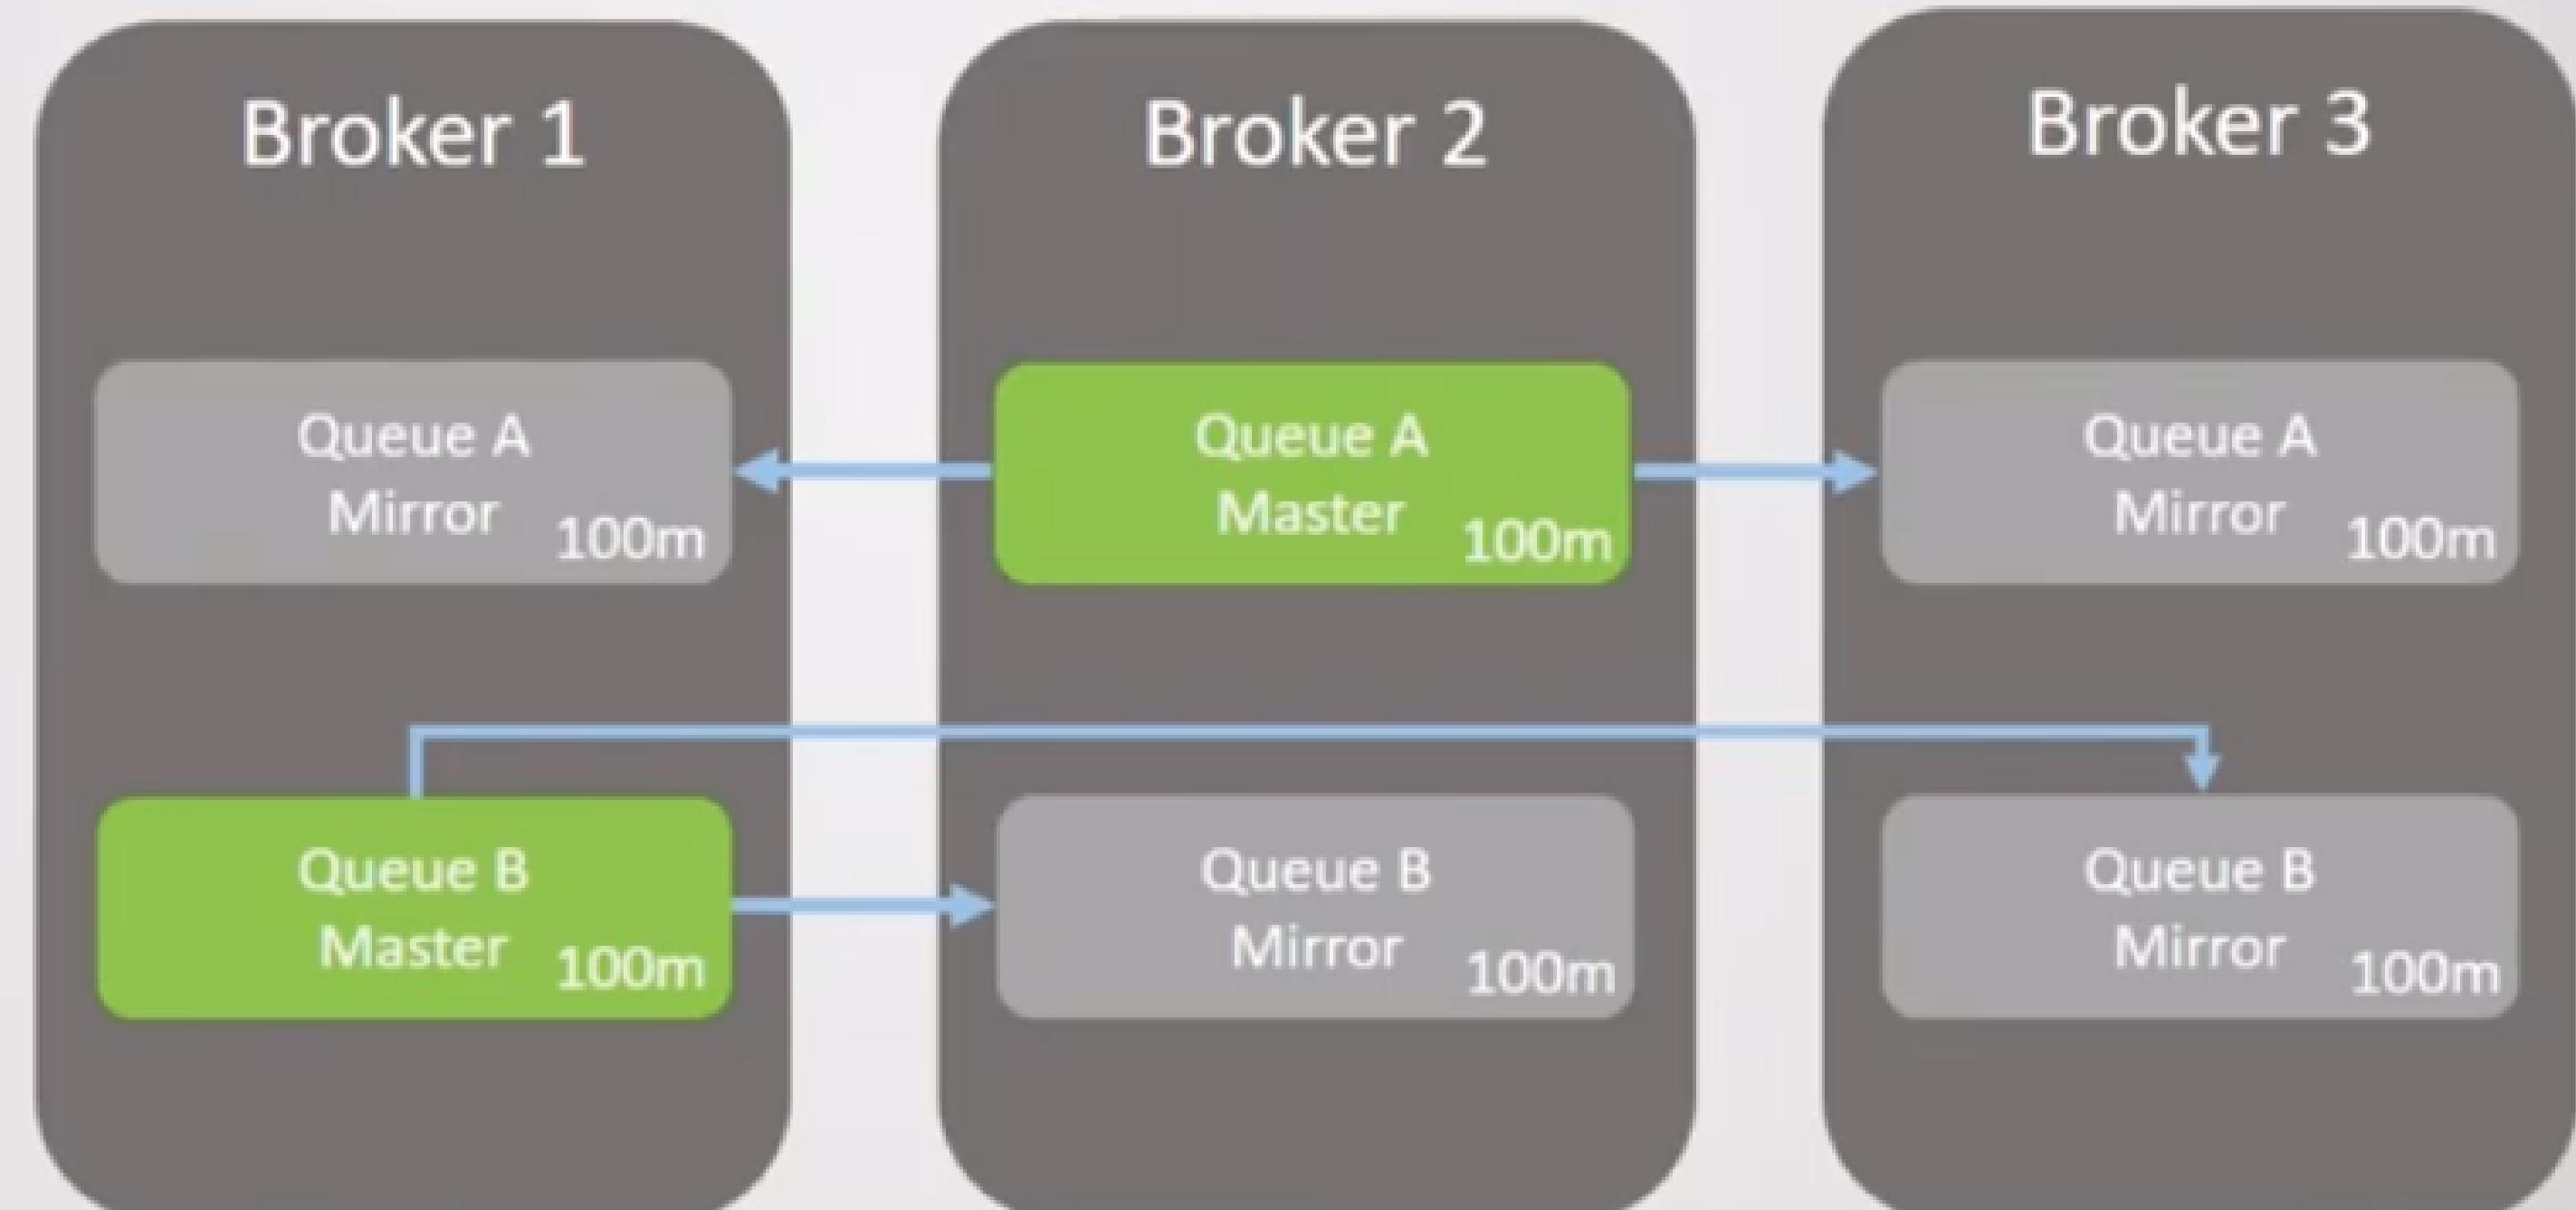


RabbitMQ

Queue Mirror
Synchronization
And
New Mirrors

RabbitMQ - Queue Mirrors - Synchronization

Three nodes, two mirrored queues each with 100 million messages

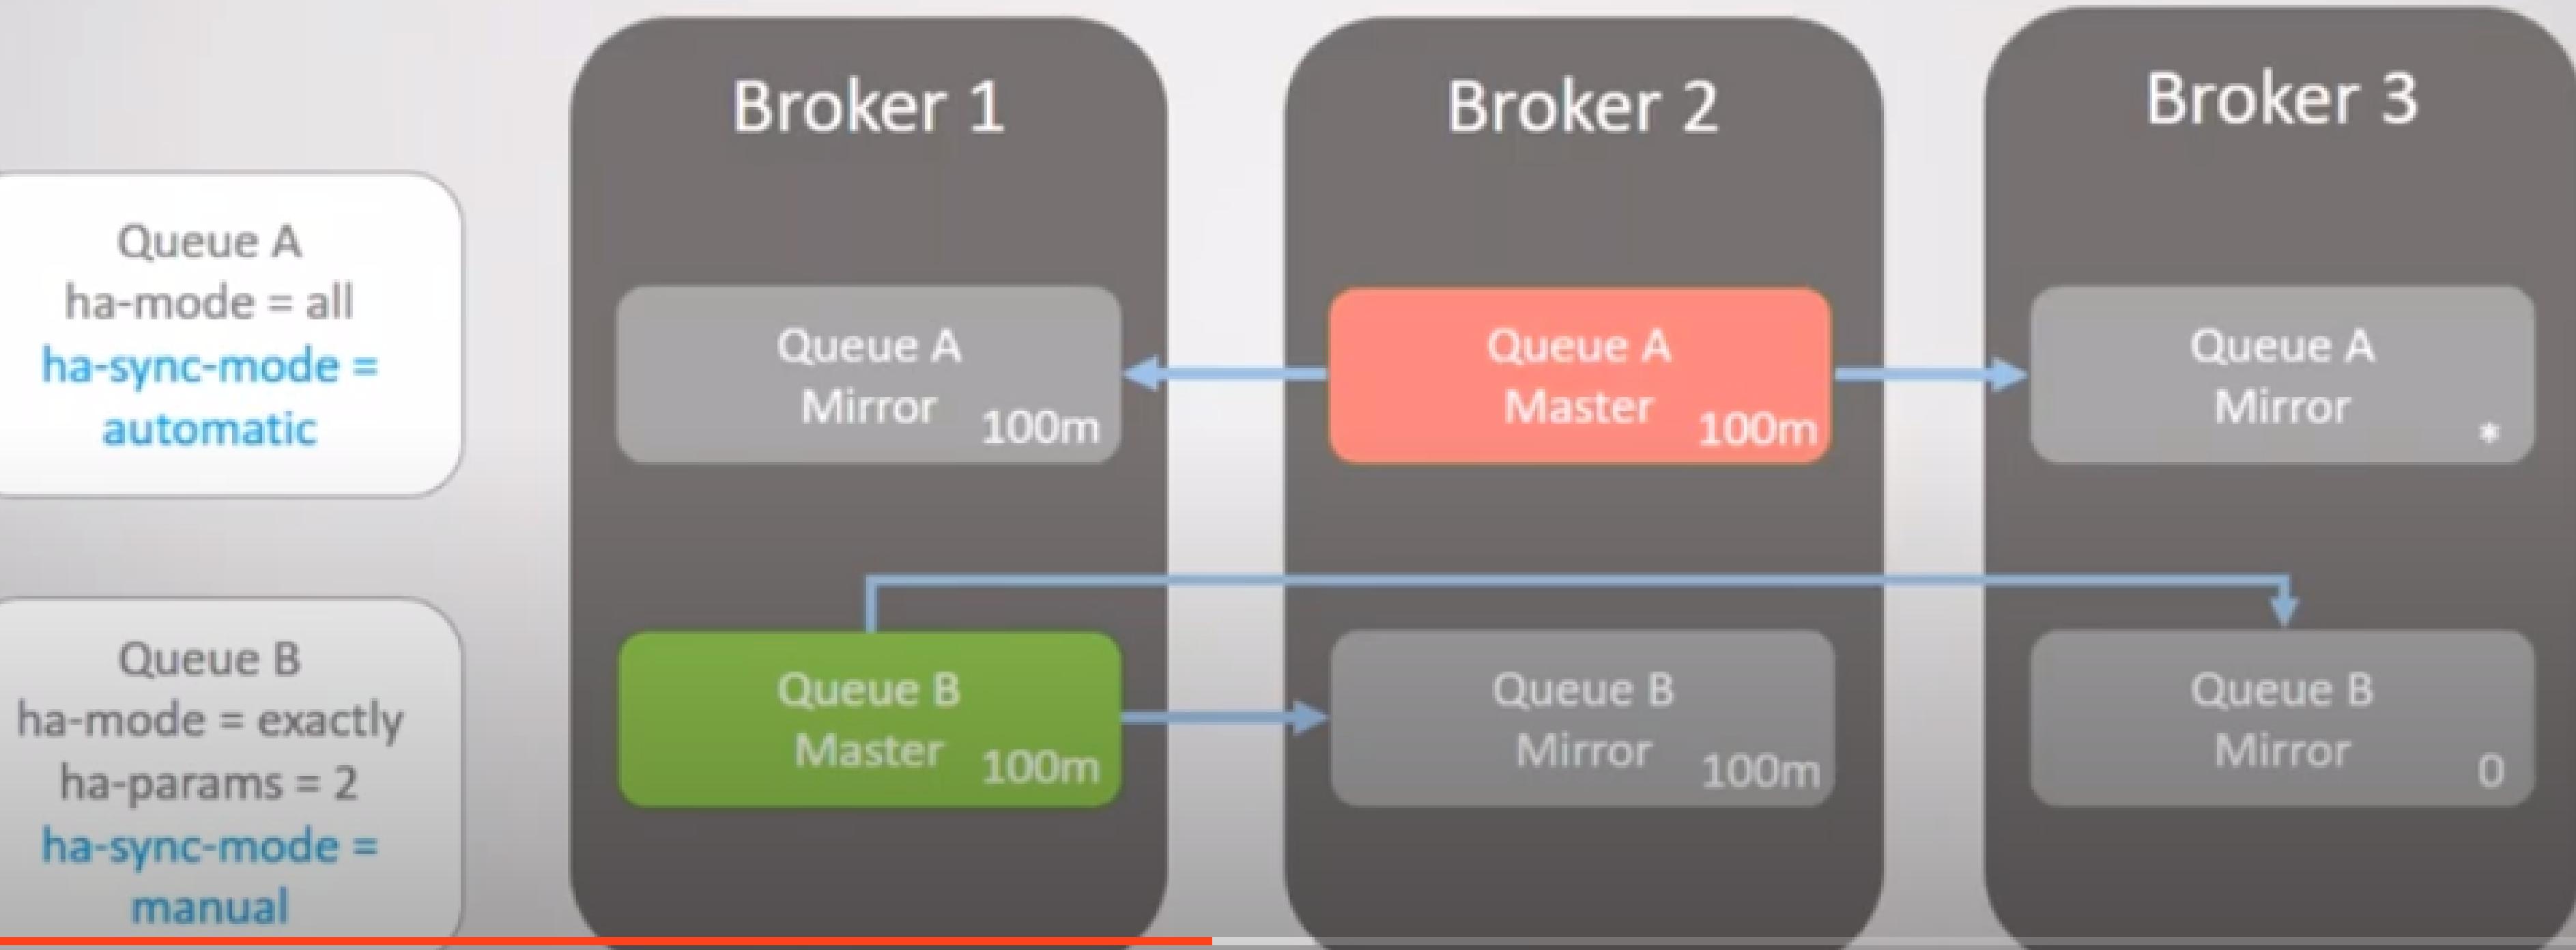


RabbitMQ - Queue Mirrors - Synchronization

Broker 3 comes back.

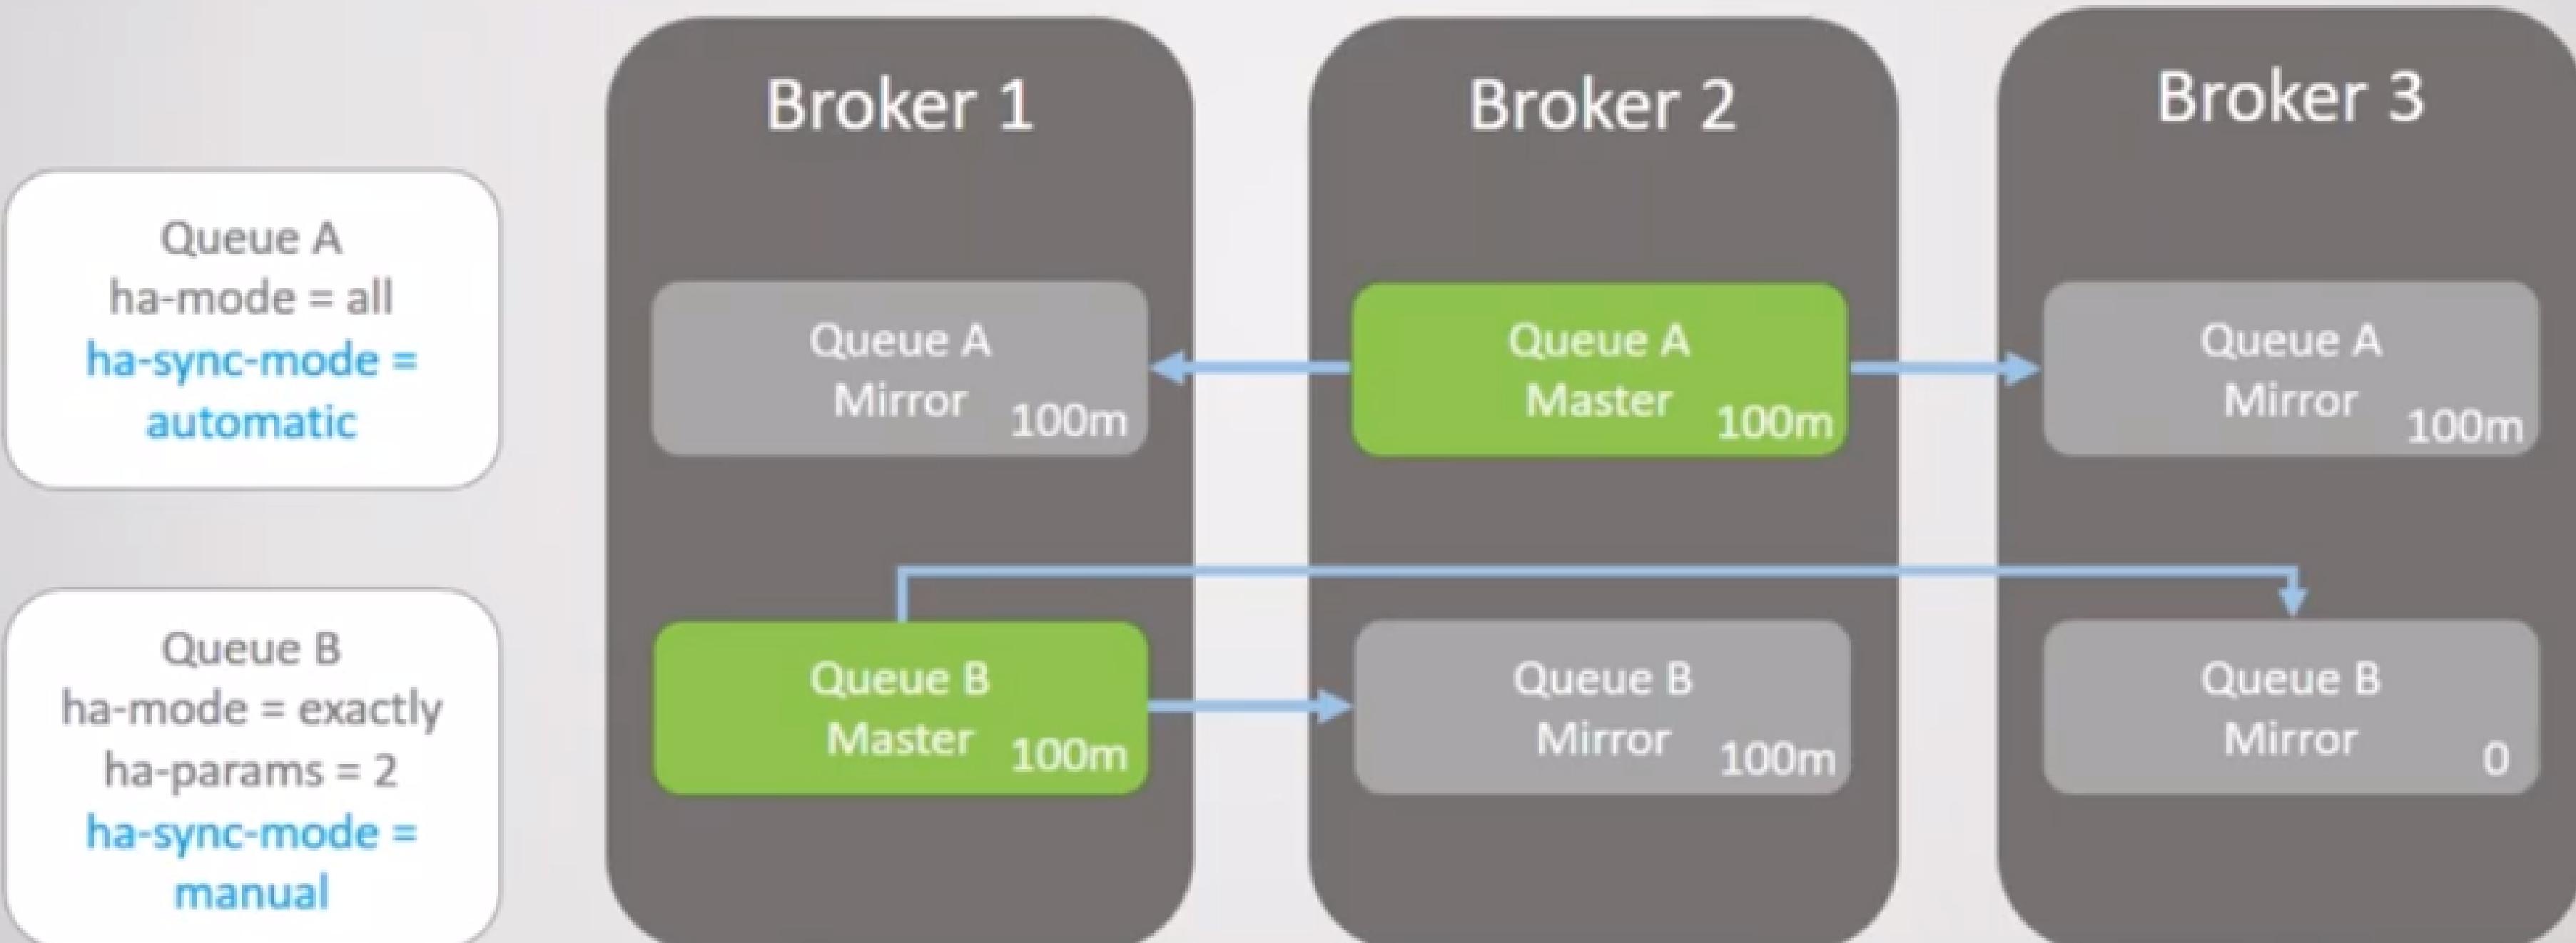
Queue A is unavailable due to synchronization.

Queue B is available but mirror on broker 3 remains at 0 messages.



RabbitMQ - Queue Mirrors - Synchronization

Queue A synchronization completes and the queue becomes available again.



RabbitMQ

Optimizing for High Throughput

- Producers fire and forget
- Consumers use auto-ack mode
- Non-Persistent messages
- Non-Mirrored Queues
- Cluster for throughput

Balancing Data Safety with High Throughput

- Producers wait periodically for acknowledgements
- Consumers group acknowledgements with the Multiple flag
- Persistent messages
- Cluster
- Queue mirroring with 1 mirror*

Optimizing for Data Safety

- Producers wait for acknowledgements after each message or after small number of messages
- Consumers acknowledge each message individually
- Persistent messages
- Cluster for durability
- Queue mirroring with 2+ mirrors.
- `promote-on-failure=when-synced`
- `ha-sync-mode=automatic` for active queues

Synchronous
Replication

Replica Synchronization
Can Cause Unavailability

RabbitMQ

Consumer Side
Redelivered flag

Fire-and-forget
Publisher Confirms

Tunable
Consistency
Vs Availability
Vs Latency
Vs Throughput

Message
Acknowledgements

Configurable
Redundancy